

A Cloud-Based DevOps Toolchain for Efficient Software Development

Ruiyang Ding



A Cloud-Based DevOps Toolchain for Efficient Software Development

Master's Thesis in Computer Science
EIT Digital Master's Programme in Cloud Computing and Services

Distributed Systems group
Faculty of Electrical Engineering, Mathematics, and Computer Science
Delft University of Technology

Ruiyang Ding

23rd August 2020

Author

Ruiyang Ding

Title

A Cloud-Based DevOps Toolchain for Efficient Software Development

MSc presentation

31th August 2020

Graduation Committee

Prof dr. ir. D. H. J. Epema (chair)	Delft University of Technology
Mikko Drocan	Eficode Oy
Dr. J.S. Rellermeyer	Delft University of Technology
Prof. dr. M.M. Specht	Delft University of Technology

Abstract

In the traditional software development life cycle, development and operation are divided into different departments. The conflict between departments and, besides, the lack of automation usually leads to low software development efficiency and slow software delivery. Thus, the concept of DevOps is introduced, which combines different departments and automates the process to make software delivery faster and easier. The DevOps toolchain is one important component for adopting DevOps. On the other hand, the adaptation of cloud technology, especially serverless computing makes it tempting for us to investigate what benefits serverless computing brings to the DevOps toolchain.

In the first research question, we examine the benefits that AWS serverless platforms bring to DevOps toolchain. To answer this research question, we develop a DevOps toolchain hosted in Amazon Web Services (AWS) and leverage the serverless computing service. In addition we examine what does each serverless computing service brings to the DevOps toolchain, examine how does the performance of the DevOps toolchain changes with or without using serverless computing service. Our research shows that serverless computing services such as AWS Fargate could reduce the cost, operation effort, and improves performance by enabling parallel execution. Our experiments show that in contrast to a toolchain hosted in a traditional cloud server vs the toolchain that was developed by us using serverless computing service could reduce the total runtime of parallel execution up to 65%.

In the second research question, we focus on the integrated toolchain build with AWS DevOps tools from AWS serverless platform. We build a demo integrated DevOps toolchain with AWS DevOps tools and compare the integrated toolchain with the non-integrated toolchain that we built. We find that the integrated toolchain significantly reduces the development time by providing an out-of-box solution for the software team. In addition, the better integration with underlying cloud infrastructure provides more functionality such as global monitoring and blue/green deployment. However, we also find that from the experiment that the performance of the integrated toolchain is lower due to the limitation of resources which also come with a high cost.

Preface

Helping customers transform their software development practices to DevOps is one of the main business activities of Eficode, and DevOps toolchain is an essential part of the transformation process. On the other hand, as an advanced partner of multiple cloud providers, Eficode is interested in what cloud technologies could bring to the DevOps toolchain. In this thesis project, I focus on the AWS serverless platform¹ in Amazon Web Services and discuss what change, especially benefit can it bring to a DevOps toolchain.

The process of carrying out this project is not an easy task; setting up the cloud infrastructure requires an enormous number of operational and configuration tasks. The vast but unregulated plugin eco-system of Jenkins also leads to problems like lack of plugin documentation, dependency hell and an unstable plugin such as plugin for using ECS as Jenkins agent. There does not exist that much previous research about the serverless within DevOps toolchain. Moreover, as a student with a software development background, the lack of prior experiences in the related field means much study is needed before I can start the project. All these tedious and unexpected tasks above, plus the tight thesis schedule did make me frustrated in the middle phase of this project. Despite all of these, I managed to achieve the defined goal by answering all the research questions and implement the demo. Through this project, I familiarized with different exciting tools for cloud and DevOps; it opens up a whole new area for me and will help me with my future career in Eficode. I hope the result will give Eficode more insight on the capability that AWS serverless platform could bring for the DevOps toolchain.

Now I'm at the end of this two year's journey, and I feel grateful that I made this life-changing decision to come and study in Europe. I have to say, it was not an easy decision to quit the ongoing research master's study back in Shanghai, China and start all over again in a brand new environment. The two year's study was a journey full of struggle – in both financial and study. However, what it brings is more than what I expected: international experiences by EIT Digital, great friends from different countries, two intern/work experiences, and precious knowledge that combines business and technologies. Most importantly, I could study the topic that I'm interested in and start a career within this field, which is what I was not able to

¹<https://aws.amazon.com/serverless/>

do back in my previous research master's study.

I sincerely thank my supervisor Mikko Drocan from Eficode, Prof.dr.ir. D.H.J. Epema from TU Delft for their guidance and support in writing during this unprecedented time. Thanks to Eficode, for giving me this precious opportunity and sponsoring this thesis project. I also would like to give my gratitude to EIT Digital master school for the scholarship, which made it financially possible for me to finish the two-year's master's study. Lastly, special thanks to Tatu Kairi and Nils Haglund from thesis support team in Eficode for giving the weekly support on my thesis, and Eficode IT for providing the AWS cloud environment.

Ruiyang Ding

Helsinki, Finland
23rd August 2020

Contents

Preface	v
1 Introduction	1
1.1 Problem Statement	2
1.2 Research Approach	4
1.3 Thesis Structure and Main Contributions	5
2 Background and Concepts	7
2.1 Agile software development	7
2.2 Continuous Integration & Continuous Delivery	8
2.2.1 Continuous Integration	8
2.2.2 Continuous Delivery and Continuous Deployment	10
2.3 DevOps	11
2.3.1 Emergence of DevOps	11
2.3.2 Relationship with Agile	12
2.3.3 Elements	12
2.3.4 Toolchain	16
2.4 Serverless Computing	18
2.4.1 History	19
2.4.2 Characteristics	19
2.4.3 Limitations	21
3 Overview of Current Serverless Cloud Services Offering in AWS	23
3.1 AWS Elastic Container Services with Fargate	23
3.1.1 AWS Elastic Container Service(ECS)	23
3.1.2 AWS Fargate	24
3.2 AWS Lambda	25
3.3 AWS CloudWatch	25
3.4 AWS Developer Tools	26
3.4.1 CodeBuild	26
3.4.2 CodeDeploy	26
3.4.3 CodePipeline	26
3.4.4 CodeStar	27

4	Design of DevOps Toolchains	29
4.1	Case Project	29
4.2	Design of Non-integrated DevOps Toolchain	31
4.2.1	Architecture	31
4.2.2	Introduction to Tools Used in the Implementation	32
4.2.3	Infrastructure as Code (IaC)	36
4.2.4	Version Control	37
4.2.5	Continuous Delivery Pipeline	39
4.2.6	Deployment Environment	43
4.2.7	Monitoring	44
4.3	Design of Integrated Serverless DevOps Toolchain	46
4.3.1	Continuous Delivery Pipeline with AWS CodePipeline	47
4.3.2	Source Control with AWS CodeCommit	48
4.3.3	Build and Test with AWS CodeBuild	48
4.3.4	Blue/Green Deployment with AWS CodeDeploy	49
4.3.5	Integration of AWS DevOps Tools using CodeStar	51
4.4	Comparison between Integrated and Non-integrated Toolchain	52
4.4.1	Implementation and Cloud Deployment	52
4.4.2	Extensibility and flexibility	54
4.4.3	Integration Between Tools	55
4.4.4	Visibility	55
4.5	Challenges in Implementation and Design of DevOps toolchains	56
4.5.1	Challenge I: The Enormous and Unregulated Jenkins Plugins System	56
4.5.2	Challenge II: Fargate Does not Supports Container runs in Privileged Mode	56
4.5.3	Challenge III: Slow Starting Time for Agents in AWS Fargate	57
4.5.4	Challenge IV: No Enough Visibility in AWS DevOps tools	57
5	Performance Comparison and Evaluation	59
5.1	Experiment 1: Experiment on Serverless Container Services	59
5.1.1	Test Task and System Description	59
5.1.2	Performance Properties and Evaluation	62
5.1.3	Result and Evaluation	62
5.1.4	Conclusion	65
5.2	Experiment 2: Experiment on Integrated DevOps Toolchain	65
5.2.1	Test Task and System Description	66
5.2.2	Performance Properties and Evaluation Criteria	66
5.2.3	Quantitative Experiment Result and Evaluation	67
5.2.4	Conclusion	70
6	Conclusions and Future Work	73
6.1	Conclusions	73
6.2	Future Work	76

Abbreviations and Acronyms

AWS	Amazon Web Services
CD	Continuous Delivery
CI	Continuous Integration
EC2	Amazon Elastic Compute Cloud
ECR	Amazon Elastic Container Registry
ECS	Amazon Elastic Container Service
ELB	Amazon Elastic Load Balancer
EKS	Amazon Elastic Kubernetes Service
KPI	Key Performance Indicator
OS	Operating System
vCPU	Virtual Processor, Virtual CPU
VM	Virtual Machine
YAML	Yet Another Markup Language

Chapter 1

Introduction

The Agile Manifesto [1] drafted by Kent Beck et al. in 2001, created the Agile software development method. Since then, this software development method has drawn attention to the industry. Agile has become a leading standard for the software development industry, with multiple further enhancements aiming to tackle certain business-specific challenges. The Agile method advocates the shorter development iteration, continuous development of software and continuous delivery of the software to the customer. The goal of Agile is to satisfy the customer with early and continuous delivery of the software [1]. The Agile, which aims at the improvement of the process within the software development team and the communication between the development team and customers [2] improves software development and makes it more iterative and thus faster. However, it does not emphasise the cooperation and communication between the development team and other teams. In real life, the conflict and lack of communication between the development team and operation teams usually becomes the barrier for efficient development of a software project [3].

Hence, in answer to how to solve the gaps and flaws when applying Agile into real-life software development, the concept of DevOps emerged. The term "DevOps" is created by Patrick Debois in 2009 [4], after he saw the presentation "10 deployments per day" by John Allspaw and Paul Hammond. While Agile fills the gap between software development and business requirement from the customer, the DevOps eliminates the gap between the development team and the operation team [5]. By eliminates the barrier we mentioned in the last paragraph, DevOps further enhances and smoothers software delivery. In conclusion, DevOps means a combination of practices and culture which aim to combine separate departments (software development, quality assurance and the operation and others) in the same team, in order fasten the software delivery, maximising delivered without risking high software quality [6][7].

In software engineering, the toolchain is a set of tools which are integrated for performing a specific objective. DevOps toolchain is the integration between tools that specialised in different aspects of the DevOps ecosystem, which support and

coordinate the DevOps practices. The DevOps toolchain helps organisations in creating and maintain an efficient software delivery pipeline, automate the development process [8] which are the keys aspects in DevOps. On the other hand, DevOps relies strongly on tools. There exist specialised tools which help teams adopt different DevOps practices [9].

Traditionally, a DevOps toolchain was to have individual tools which were stand-alone and from different vendors. The tools were usually on-premise. However, such toolchains can also be deployed on cloud virtual machines. In this report, we define this type of toolchain as **non-integrated toolchain**. We also define the toolchain that is delivered as an single platform from an single vendor as **integrated toolchain**. We will introduce more about integrated toolchain in problem statement below.

At the same period that the tools for DevOps emerged and developed, the cloud technologies also developed rapidly. This led to the emigrations of Serverless Computing. The Serverless Computing is a modern cloud computing model in which everything is built and executed in the applications running in the cloud environments without thinking about physical servers [10]. It also allows developers to build the application with less overhead [10] and more flexibility by eliminating infrastructure management tasks [11]. With serverless computing technologies, many new cloud technologies emerged, which gives developers an alternative way from traditional cloud servers or cloud virtual machines. For example, Functional computing allows the application to be divided by functions and designed under the event-driving paradigm without managing the hardware infrastructures. The on-demand nature of the serverless computing could be used to deploy event-based components of a DevOps toolchain, such as post-deploy testing and logging. Managed scalable container services in the cloud enable the user to run the container-based application directly on the cloud, which allows the toolchain scalability. DevOps tools as a service [12] allow the cloud provider to deliver DevOps tools directly on its cloud platform.

Helping the customer in their DevOps transformation is one of the main business activities of Eficode, the company which we are writing our thesis. The transformation is enabled, for example, by defining, developing and maintaining a DevOps toolchain at the customer. As mentioned in the last paragraph, the new changes brought by cloud may further improve the performance and lower the cost of DevOps toolchain development – both money and time. As part of our thesis work at Eficode, we will investigate how serverless computing services enhances the DevOps toolchain.

1.1 Problem Statement

As per the last paragraph, serverless computing could bring enhancements to the DevOps toolchain. Currently, there are several cloud providers that provides cloud services using serverless computing technologies. Among them, Amazon Web

Services (AWS)¹ has the largest market share and is the first cloud provider which provides serverless computing services. According to the report from Gartner [13], the market share of AWS was 47.8% in the year 2018, which makes it the largest cloud provider in the world.

Nowadays, the serverless computing services in AWS has already been expanded to a set of fully managed services called "AWS serverless platform"². This platform includes new AWS cloud products that leverage the serverless computing technologies. These products include, for instance, AWS Lambda³ for function computing and AWS Fargate⁴ for managed container services. AWS also gains the most popularity among the developers that use serverless technologies. The most recent survey report [14] from Cloud Native Computing Foundation (CNCF) shows that 51% of serverless users are using AWS Lambda, while 68% of developers who are not using Kubernetes are using AWS ECS to hosting their containers. As the Advanced AWS partner, AWS is being used as the main cloud providers in the customer projects by Eficode. Furthermore, the company keeps looking for ways to leverage serverless computing services in AWS to produce cost-efficient solutions for the customers.

However, despite the serverless computing is being extensively used, and an enormous number of research papers about the use-cases or benefits of serverless in data analysis [15], for container-based microservices [16], or for IoT applications [17] [18], the benefits of serverless within DevOps has not yet been discussed. There are papers [19] and a book [20] about DevOps toolchain for serverless applications. Nevertheless, there is still a lack of research on how serverless helps DevOps toolchain itself. Thus, our first research question is to fill the gap by answering this question.

The second area we need to investigate in our project is the integrated DevOps toolchain that is powered by serverless DevOps tooling in AWS.

The integrated DevOps toolchain is delivered as a cloud-based single platform that allows development teams to start using DevOps toolchain without the challenge of having to choose, integrate, learn, and maintain a multitude of tools. In other words, the cloud based-integrated DevOps toolchain is to offer DevOps toolchain as a service. In AWS, this is offered by AWS CodePipeline (as the platform) and Several serverless tools that integrated with CodePipeline.

This integrated toolchain is one of the new changes that serverless computing brings, but it also leaves a question to the development team who is trying to build DevOps toolchain in AWS: which kind of toolchain should they select? Should they stick on the previous non-integrated toolchain or embracing the integrated one? The integrated DevOps toolchain provides an out-of-box integrated solution for the whole DevOps lifecycle, which is tempting. However, apart from the advertisement from the vendors of these "DevOps" platforms, we still lack third party

¹<https://aws.amazon.com/>

²<https://aws.amazon.com/serverless/>

³<https://aws.amazon.com/lambda/>

⁴<https://aws.amazon.com/fargate/>

researches about the comparisons between these two.

Based on the above, the research questions could be summarised as below:

RQ1: *How can serverless computing services in Amazon Web Services enhance the DevOps toolchain?*

RQ2: *How does the integrated toolchain build with AWS DevOps Tools compare with the traditional non-integrated toolchain?*

1.2 Research Approach

The first step of our Research is to investigate the current serverless offering in Amazon Web Services (AWS), which is one of the cloud services mainly used in Eficode. We analysis which serverless computing services can be used in our implementation and their possible roles within a DevOps toolchain.

In the next step, we design and implement both traditional non-integrated and integrated toolchain based on the DevOps practices and tools used by Eficode. In the design and implementation of the toolchain, we focus on the following DevOps practices: Version Control, Configuration Management, Continuous Delivery and Monitoring. The goal of the implementation is: First, validate the availability of using AWS serverless computing services in the traditional non-integrated toolchain. Thus, in the process of developing and deploy the toolchain, we could already partly answer the RQ1 by answer how the serverless computing services be used in our DevOps toolchain. Second, the implementation served as the environment for experiments, which could answer both research questions.

Our next step of the study are the experiments. The first experiments is to comparing the metrics measured from the toolchain with and without using certain serverless computing service from AWS. These metrics cover different perspectives, which including cost, performance and ease of use. Bu doing this, we could examine how the serverless cloud computing service could improvement the DevOps toolchain from performance's aspect, which is related to the RQ1.

In the second experiment that related to RQ2, the non-integrated toolchain is used to compare with the integrated DevOps toolchain built by the AWS DevOps tools. We again measure the metrics in these two toolchains. The process is similar to what we do on experiment 1.

Besides, we conduct a study on a comparison between an AWS based traditional toolchain and this out-of-box integrated DevOps toolchain that is also provided by AWS. The reason that we keep the comparison scope within AWS is that by doing this, we make sure that hardware in both toolchains is from AWS, this could eliminate the errors caused by the hardware difference between vendors and focuses on the difference caused by toolchains themselves.

1.3 Thesis Structure and Main Contributions

In Chapter 2, we introduce concepts within the scope of DevOps. We also introduce the concepts in cloud computing which are related to our research. Chapter 3 is focusing on a survey on serverless computing technologies which the DevOps toolchain could make use of. Chapter 4 focuses on the design and the implementation of our DevOps toolchains(both non-integrated and integrated). Chapter 5 focuses on the experiments and evaluations, which show how the serverless computing services introduced in Chapter 3 could benefit DevOps toolchain. We also compare integrated/non-integrated toolchain in Chapter 5. we finally summarise our research and answer the research questions in Chapter 6.

The main contributions of this thesis project are:

- We provide a study on how could the DevOps tools leverage the cloud services to reduced development/deployment difficulties, lower the cost and improving the performance. This part of research could help the software team which is going to employ DevOps understand the practices needed. Besides, the research gives them a clearer scope of the tools needed for implementing the practices.
- We give the overview of two different types of DevOps toolchains. We also implement demo prototypes for each type of toolchain and conduct experiments with these prototypes. The experiment result shows a comparison between different toolchains. It could help the software team understand which toolchain cloud be selected based on the needs.

Chapter 2

Background and Concepts

The DevOps and serverless cloud computing services are the two main fields that our work related to. In this chapter we will introduce both of them.

Section 2.1 introduces the Agile software development method, which is the basis of DevOps. Section 2.2 shows the definition of continuous integration and continuous delivery, as one of the most important practices of the DevOps. Section 2.3 presents the definition, components of DevOps, also, the definition and components of a DevOps toolchain. Section 2.4 introduces serverless computing, including it's concept, characteristics and limitations.

2.1 Agile software development

The term "Agile" represents the fast adaptation and response to the changes in the plan [21]. Agile software development is a method of software development that implements the ideology of "agile". Agile software development advocates the continuous development of software teams. The software development under this methodology will have shorter planning/development time before it delivers to the customers and could better adapt to changes in the environment and requirements.

Iterative Software Development: Agile software development uses an iterative way in the development process. The traditional software development process, like the waterfall method, requires the long and complicated planning process, and a complicated document. Once one phase of the development is done, the teams should not change the output (document and code) of this phase [22]. In contrast, agile software development aims to satisfy the customer with early and continuous delivery of the software [1]. Early means the shorter time before software delivery. Continuous means the development does not end with the delivery. The delivery means to an end of an iteration, together with a demonstration to stakeholders. After delivery, the team continues to the next iteration according to the feedback it gets from stakeholders. In each iteration, the aim of the team is not to add major features to the software, rather is to have a working and deliverable release [23].

In the ideology of agile, the best design of the software product comes from the iterative development [1], rather than the tedious planning.

High-Quality Software: The rapid development does not mean low software development quality. On the contrast, the quality of software design is highly appreciated in agile software development. The automated testing is widely used in Agile. The test cases will be defined and implements from the beginning of the development process. The test goes through the entire development iteration to ensure that the software is of high enough quality and can be released or shown to customers at any time during the iteration [24].

Collaboration: The agile software development processes include collaboration across different groups, i.e. business development team, software development team, test team, and customers. It values face to face communications [25] and feedbacks. The purpose of these communications is first to let everyone in the multifunctional agile team understand the whole project, and second to receive feedback that helps the software in the right development track. The track which aligns with the requirement of the stakeholders [1].

According to the Manifesto for Agile Software Development, compared with traditional software development, the agile software development value these aspects [1]:

- Individuals and interactions over processes and tools.
- Working software over comprehensive documentation.
- Customer collaboration over contract negotiation.
- Responding to change over following a plan.

2.2 Continuous Integration & Continuous Delivery

In the software development, CI/CD refers to continuous integration, continuous delivery and continuous deployment [26]. As we mentioned in 2.1, agile software development requires continuous software quality assurance and iterative development. Currently, CI/CD is one set of the necessary practices for the team to become agile by achieving the requirements above. Figure 2.1 shows the relationship between these 3 practices.

2.2.1 Continuous Integration

Continuous integration is the base practice of all practices within CI/CD, and continuous delivery/deployment is based on the continuous integration [26]. The continuous integration means the team integrate each team member's work into the main codebase frequently(multiple times per day). "Integrate" means merge the




Figure 2.1: The relationship between continuous integration, continuous delivery and continuous deployment [26]

code into the main codebase [27]. The continuous integration rely on following practices: *Source Code Management*, *Build Automation*, *Visibility* and *Test Automation*. The definition of these practices are:

- *Test Automation*: Test automation means using separate software to execute the software automated without human intervention. It could help the team to test fast and test early [28]. There are three approaches to test automation: first, the automated unit tests which design by developers in the early stage of the project. Unit testing aims to verify if each component of the software project works properly; Second, the functional testing which tests the business logic behind the user interface. The function testing is a type of black-box test which the tester only care about if the output is as expected under certain input; Third, the graphical user interface (GUI) testing test if the user interface meets the functionality requirement. The context of GUI testing is to simulate the user's operation on the user interface.
- *Build Automation*: Automate the process of creating software build. This means to automate the dependency configuration, source code compiling, packaging and testing. It is viewed as the first step to continuous integration [29]. There are two types of build automation tools, the build-automation utility and build-automation servers [30]. Build-automation utility means the tool to generate build artifacts by compiling the source codes. The common tools that belong to this type include Cmake, Gradle and MSBuild. Build-automation server is the tool which executing build-utility tools; it allows the build to be triggered from the outside or be scheduled on a time basis. Build-automation server is usually web-based. Continuous integration serv-

ers, such as Jenkins and Circle CI, is considered as build-automation server.

- *Source Code Management:* In continuous integration, the team, maintains a single source repository and use version control system. In practice, this means one branch in the version control system act as the "mainline", while everyone works off this mainline [27]. However, everyone needs to merge the code to the mainline every day. For making sure that the mainline code still works after the merge, the mainline that merge the new code needs to be built and tested. we will further introduce this practice in Section 4.2.4.

With the help of these practices, the workflow [27] in continuous integration for each developer in the software team is as follows: In the development of each feature, the developer first pulls the code from the main codebase. During the development, new test cases should also be added to the automated test. Automated testing runs on the code after the developer finishes the feature development. This is for maintaining the code quality and minimise the number of bugs from the beginning. The actual practice for implementing this step is to have build automation tools compiled the code locally in the development machine.

After the step above, the developer already has the executable and the high quality (passed the automated test) code in the development machine before submitting the change to the repository. This represents the principle of quality and automation in agile software development. In the next step, the developer commits changes to the repository, which is the main codebase, and the system check the conflict and do the test/build again, to make sure that there are not any bugs missed in the test on the development machine. If the code passes this build and test, it will be merged to the main codebase, and the integration is done.

2.2.2 Continuous Delivery and Continuous Deployment

Continuous delivery is a practice that the software development team build software that can be released at any time during the life cycle [31]. This practice ensures that the software always high-quality and in a deployable state [32]. Continuous delivery provides a clear way for software development teams to become agile [33][34]. In the last section, we introduce the concept of continuous integration. Continuous delivery is based on continuous integration, but it further automates the software deployment process. In the software deployment pipeline, the team divides the build into several stages, first build the product, and then push the product into a production-like environment for further testing. This ensures that the software can be deployed at any time. However, in continuous delivery, deploying software to the production environment is done manually. The benefit [32][31] of continuous delivery includes:

- High quality of code: The automate and continuous testing ensure the high quality of code.

- Low risk: The software team could release the software at any time. The release process is easy, and it is also harder to make a mistake.
- Short time before going to the market: The iteration of software development is much shorter. The automated testing, deployment and environment configuration short the development life cycle. The always ready-to-deploy status shorten the time from development to market.

The continuous deployment is based on continuous delivery. The only difference is continuous deployment automates the deployment process. In continuous delivery, the software is deployable but not deployed without manual approval. In the continuous deployment, each change that passed automated build and testing will be deployed directly. Continuous deployment is a relatively new concept, and most companies have not yet put this practice into production [35]. Although continuous delivery is a necessary practice for companies to become DevOps, it has been widely used.

2.3 DevOps

The fundamental goal of DevOps is to minimise the service overhead so that it can respond to change with minimal effort and deliver the maximum amount of value during its lifetime.

– Markus Suonto, Senior DevOps Consultant, Eficode

DevOps is a set of practices that aims to combine different, traditionally separated disciplines (e.g. software development, operations, QA, and others) in cross-functional teams with the help of automation of work to speed up software delivery without risking high-quality [36].

2.3.1 Emergence of DevOps

In the pre-DevOps era, development and operations were two different teams with different goals. The interface between them is based on the ticket system, and the operation team performs ticket management. As we mentioned at Section 2.1, the goal of Agile is to shorten the deliver life cycle and quickly delivery software to the customers. Therefore, when practising agile development methods in this situation, developers try to deliver code, and they will develop earlier. However, the operation team usually will delay the process for quality control or other reasons. In practice, this causes the delay between the code change and the software delivery to the customers [37]. The lack of communication and conflict between developers and the operation team slow down the software delivery process and also make it harder for the teams to be real Agile. Therefore the concept "DevOps" is being proposed at 2008, for eliminating of the boundary between developers (Dev) and operation team (Ops).




Figure 2.2: DevOps Practices and Workflow [38]

2.3.2 Relationship with Agile

DevOps is the extension and evolution [39][37] of Agile. DevOps and Agile both driven by the collaboration ideology and the adoption of DevOps needs Agile as the key factor [39]. DevOps has a different focus on agile. DevOps focus on the whole delivery and customer satisfaction while agile is focused on the development with the requirement and customer [40]. Figure 2.2 shows the workflow and practices of a team working under DevOps.

2.3.3 Elements

In this section, we will introduce the necessary elements that an organisation need to includes while introducing DevOps as a common practice. Eficode proposed it's DevOps capability model which act as the baseline for the DevOps transformation. This model consists following elements enablement, organisation and culture, environments and release, builds and continuous integration, quality assurance, visibility and reporting, technologies and architecture. Combined with our research that based on other materials related to DevOps, we summarise the DevOps capability model at a more general level, into four aspects.

Culture

According to Walls (2013), this is being done by promoting the culture with 4 characteristics open communication, incentive and responsibility alignment, respect and trust [41].

Open communication means open discussion and debate. Communication can help minimise the gap between developers and operations teams. Traditionally, communication within the team is carried out through a very formal and stand-

ardised ticketing system. However, in teams adopting DevOps, communication is not limited to the ticket system. Instead, the team will maintain free communication throughout the life cycle of the product, and they will discuss requirements, timetables, and anything else. In addition, information sharing is also very important[42]. Metrics and project status are available to everyone in the team, so each member can clearly understand the scope of the team's work.

The **incentive and responsibility alignment** mean the entire team (consisting of Dev and Ops) has the same goals and assumes the same responsibilities. The transition from "Dev" and "Ops" to DevOps requires people who used to be in charge of only development and operation start taking the responsibility in both sides [42]. Such transition means if the product is failed, individuals or part of the team will be not solely blamed. This "no blame" culture could help each engineer be willing to take the development responsibility for the whole system [43].

Respect means all employees should respect and recognise the contribution of other teams members. A DevOps team is not a single team without any division of jobs, there is still an operation part within a team [44]. Therefore each part of the team needs to **trust** the other parts are doing the best to benefit the whole team. On the other hand, the person in the operation team will take development responsibility, and the developers will also put their hands-on operation and management[45]. To make people with different roles work in a team, trust and respect each other is critically important.

Organisation

In the organisational level, the DevOps emphasises the collaboration between different parts of an organisation. This is closely related to the "culture" part of this section. Within a team, each member should be a generalist who could understand all aspects of a project. There will not be a dedicated QA, operation or security team within a team. Instead, these are jobs that belong to everyone [43][4]. The structure and rule of the organisation should provide all members with opportunities to learn all skills needed for building the whole system.

The DevOps Handbook [4] published in 2016, proposed ways to organise an organisation for DevOps transformation. One of the principles in organising the teams is to keep the team boundary and comply with the "two pizza rule" proposed by Amazon in 2002. The rule is to keep the team size small for having a more productive team meeting. A small team could help to reduce the inter-team communication, keep the team scope bounded and small [4]. Furthermore, the smaller team also means less bureaucracy in team management. There are four benefits to have a small team:

- The smaller team allows each team member to understand the whole project easily.
- The smaller team could reduce the amount of communication needed. It could also limit the growth rate that the product could have.

- The smaller team could decentralise power. In DevOps, each team lead could define the metrics which become the overall criteria of the whole team's performance.
- In a smaller team, failure does not mean a disaster for the company. This fact allows the team to fail. Thus each employee could train their headship skill in the team without too much pressure.

Having a loosely-coupled architecture is another important organisational aspect for DevOps. The first benefit is the better safety. In the organisation with a tightly-coupled architecture, because each component is closely coupled with each other, even a small change could result in large failure [4]. The second benefit of the loosely-coupled architecture is productive. In a traditional organisation, the whole organisation shares the same development life cycle. The result of each team will be merged, tested together and deployed together, which is time costly when configuring the test environment and dependencies. A loose organisation enables each team to finish the development life cycle (from planning to deployment) independently. Each team could update their products independently, which gives the team more flexibility to align the product with the change in the customer requirement. The update of each team's product should not affect other teams product.

Automation

In the DevOps, automation means automation within the whole development and operation process. The organisations which employ DevOps aim for a high degree of automation[46]. With automation, people could be free from the repetitive work and reduce human error. It could help build the DevOps culture of collaboration, and it is seen as the cornerstone of the DevOps [47]. The main practices regarding automation are the automated testing, continuous delivery and automated operation. The automation operation includes several practices such as automated monitoring and alarming, automation infrastructure provision and environment configuration.

The continuous delivery pipeline is the core of automation within the scope of DevOps. As per discussed at 2.2.2, the continuous delivery will ultimately automate all steps between the developer to commit the code to the product in the production. In addition, the continuous pipeline brings together all automated steps within DevOps life cycle.

Infrastructure as code is a practice which helps to achieve automated operation part, specificity, environment configuration and infrastructure provision. The Infrastructure as Code (IaC) means everything at the software infrastructure level is defined as code [48]. Because it is code, the developer could use the automation methodology used in the software development to manages and deploy these codes. According to Christof et. (2016), under IaC, infrastructure can be shared, tested, and version-controlled [7]. This could help emphasise the automation within the

operation scope. In addition, IaC the team could be free from the tedious environment configuration and shorten the product development lifecycle. Automating server configuration with IaC helps the developers and operation staff know the server configuration equally [47], which help build the culture of shared responsibility and trust.

Monitoring and Measurement

Monitoring is to continuously collect the matrices from the running system. Monitoring provides the team with good visibility on the whole system. The team could get an update on the system status, and find the problems in the system in time. To conducting the monitoring, the monitoring system needs to do the measurement, which is to collect data properly from the system. The measurement is defined as reducing the uncertainty through observation, which producing quantitative result [49]. The organisation should properly use the result (metrics).

In the DevOps way of development, the testing is the key to maintain the quality of the software continuously. However, when the product enters the production, we cannot test the software any more. So, we need monitoring to keep track of the status of the product [50]. According to State of DevOps report from DORA and Google Cloud, the good monitoring structure and the wisely usage of the data from monitoring for making the business decision could improve the software delivery performance [51]. Thus, monitoring is an important component of DevOps.

With monitoring, the software team could keep tracking the status, and maintain the quality of deployed production. Monitoring enables the management teams to track the KPIs during the production. The monitoring has also enabled the team to collect the data from customers' usage behaviour. This helps the agile development team to improve in the next iteration of the product [42].

The development of monitoring should be in parallel with the main product, and the monitoring system can be already be used against the "staging deployment" (see Figure 2.1) at the early stage of the iteration. By the practice of parallel development, the development team can improve the monitoring system continuously together with the main software system. In addition, the parallel development helps the team to find the gap in the monitoring earlier [50].

As we mentioned in the "Culture" section, the collaboration is an important part of the DevOps culture. Collaboration needs the communication and information sharing between the development(Dev) and operation(Ops) team. The monitoring could be one of the channels between the Dev and Ops since it can expose the information of the whole system, which helps team members to understand the system as a whole. This helps the team achieving the point we mentioned at 2.3.3 (Culture) that the project status and matrices should be available to every team members.

2.3.4 Toolchain

A DevOps toolchain is a set of tools that are integrated to aid the software development, deployment and management through the whole software development lifecycle. The goal of DevOps toolchain is to help the software development fits the DevOps principles [8][52][5]. Within DevOps toolchain, each tool in the toolchain related to a specific activity in DevOps, for example, Jenkins works as automating tasks.

According to [5], DORA state of DevOps reports [51][53][54] and our previous definition of the DevOps, we summarise the essential component of a DevOps toolchain as below.

Project Management & Planning

Planning software development project, track the tickets and the issues, communication between and within the teams. The project management tools help to implement the DevOps culture, which enhances collaboration and knowledge sharing.

Configuration Management

Configuration management provides a central platform to manage the configuration across the assets. Such a tool allows the team defines the desired state of the assets in a configure file. Then the tool automates the configuration process, which reaching the assets to the defined status. In the cloud environment, a common practice of configuration management is through infrastructure as code, which is define the cloud infrastructure, services configuration and deployment orchestration as configuration file [55].

Continuous Integration

Continuous integration (in short: CI) is the top practice for improving the Deployment Frequency [53]. It is one of the most important parts of DevOps toolchain. As we introduced at 2.2.2, CI allows the developers to integrate their work more frequently to the production products, and it shortens the time to the market of the product. The automatic testing and code analysis integrated into the CI continuously maintain the quality of the product. CI tools also automated the most parts of the software development pipeline, In conclusion, CI helps the system fulfil the DevOps definition (2.3) by speed up the delivery by automation, maintain the quality by continuous quality assurance. The location of the CI server is flexible, depends on the scenario, it could be either on-premises (on a development machine or a local server) or deployed on the cloud. Nowadays, some vendors provide CI as a service. In this case, the CI server is hosted, managed by the vendor and be provided to the user as an online service. As we introduced in 2.2.2, CI brings together all automation tools, and automate the DevOps workflow, which connects

multiple automated processes. CI is the "confluence point" of the most DevOps tools, and thus the core of the whole DevOps toolchain.

Version Control

Version control is an important part of DevOps toolchain. It is a system that could record and track the changes in a set of files over time. Version control simplifies the collaboration between team members. Furthermore, allow the simultaneous development of the different parts of a software system. According to [56] and [53], version control is the top practice when it comes to improving the multiple metrics in DevOps. Version control becomes the indicator of the software system performance [56] infrastructure as code. An important DevOps practise we mentioned at 2.3.3 also relies on version control.

The version control is composed of a repository and the checkout. The repository is a database which records all history versions of the files. The checkout is the local copy of all the files. The user could edit the files in the checkout, then commit the change to the repository. Depends on the location of the repository, there are three types of version control systems [57].

- **Local Version Control Systems:** The repository located locally on the development machine where the user keeps the checkout. However, the repository is stored in a separated version database that keeps all changes of files.
- **Centralized Version Control Systems:** The repository located in a centralised server, while there are multiple checkouts on multiple development machines. This allows multiple developers to work together under a version control system. However, such setup is not fault-tolerated because to the VCS server is centralised.
- **Distributed Version Control Systems:** This is a type of VCS system that leverage the peer-to-peer approach. Most modern VSC system, such as Git, is using such approaches. The file history is not only kept in the server, but also in each development machine where has the checkout. Once the server dies, the history record will not be lost, and the development machine that retains the file history record will copy the file back after the server is up and running again.

Monitoring

The monitoring system is one of 4 basic elements of DevOps, as we mentioned at 2.3.3. In the DevOps toolchain, the monitoring system detects the failure in the whole system and helps the software team find the problems earlier. The team could also collect performance-related matrices with the help of the monitoring system, which could be used for optimising the application. Besides, the monitoring system can also help in the business aspect by collect the KPI-related metrics

form the user's behaviour. Monitoring system combines the data measured from the system and then visualises these data on the dashboard. The visualisation helps people which is not in the operation team understand the data.

Test Automation

The test automation tool could verify the code before it being built. Such tools usually come either an independent tool or a plugin that embedded within IDE, build server and continuous delivery pipeline. The integration of testing with other tools such as continuous integration pipeline makes it easy for the organisation to implement the quality gate in the software development [50]. The test automation tool runs on the local development machines after each build, after committing to the repository, and before the deployment to production. This policy makes sure that the testing and quality control goes through the whole software life cycle. We will introduce when runs automation testing and what kind of testing will be run in section 4.2.5.

2.4 Serverless Computing

In this section, we focus on the concepts of Serverless Computing. We will have more discussion regarding the new cloud service based on Serverless Computing in the next chapter.

Serverless Computing (in short: Serverless) is a cloud execution model in which the cloud provider manages the server and resources allocation. The popularity of serverless is precipitated by the development of microservices and container technologies [58]. A survey by the Cloud Native Computing Foundation (CNCF) showed that in 2019, 41% of respondents used serverless technology in production, compared with 32% in 2018[14]. The report of this survey also shows that serverless architectures and cloud functions are being used by 3.3 million developers [14] in 2019.

In traditional cloud computing services, users rent a fixed number of cloud servers from the cloud provider, and then the cloud provider charges users based on the lease period and server type (pay-as-you-go model). In serverless computing services, developers only pay based on the execution time of the program. Another difference between serverless computing and traditional computing method is that, in serverless computing, users doesn't need to care about the physical machine that runs the application. In addition, in serverless, the environment that runs the application will be destroyed shortly after the application terminates. However, the task is still running on a physical cloud server that is fully managed by the cloud provider. This means that when serverless is used, the user leaves all server provisioning and management tasks to the cloud provider [59].

2.4.1 History

In the early days of cloud computing, the consideration behind cloud computing design was that developers only needed to transfer their deployment environment from a local server to a server on the cloud. Therefore, cloud virtual machines (for example, Amazon Web Service EC2) is the main form of providing cloud services. After Amazon Web Service started offering the service with the virtual machine, Google entered this field for competing with AWS, but in another direction. In 2008, Google released Google App Engine (GAE) ¹[60]. The platform allows developers to run their code without managing the cloud virtual machine. This makes Google the first in the main cloud providers to allow the developer to run code on its cloud without provisioning and to manage the cloud servers. However, the GAE only allows the developer to run the python code that is programmed with Google's framework, rather than running arbitrary Python code. Amazon Web Service (AWS) introduces AWS Lambda in 2014, make Amazon the first public cloud provider that provides serverless computing platform[61]. Since then the serverless computing starts its rapid commercial development. Following AWS, other providers also introduced their serverless computing platforms. Only in a single year (2016), Google ², Microsoft ³, and IBM ⁴ released their serverless computing platform respectively. In the beginning, the serverless computing offering of vendors is limited to function as a service (FaaS), and the company only use the serverless computing in some supportive components like scheduled tasks. Nowadays, the serverless is expanding its application scope together with the extension of serverless offering in cloud vendors. For example, AWS provides a serverless platform ⁵ with the different component for a modern application, as well as tools and services for DevOps. These components are enough for a software team builds microservices architecture backend service for web applications, with DevOps toolchain that also builds in AWS. In a word, the serverless cloud service cloud now covers the entire development life cycle.

2.4.2 Characteristics

we conducted research related to the main characteristics of serverless computing, and we summarise our finding in following four main characteristics based on materials [62][63][59] we read.

Event Driven

Event-Driven means the serverless applications is usually triggered and start running due to an event. There are different kinds of event that could act as a trigger.

¹<https://cloud.google.com/appengine>

²<https://cloud.google.com/functions>

³<https://azure.microsoft.com/en-us/overview/serverless-computing/>

⁴<https://www.ibm.com/cloud/functions>

⁵<https://aws.amazon.com/serverless/>

The first one is the HTTP request. When an HTTP request reaches the server, the serverless application could be triggered to read the context of this request, execute the code, return the HTTP response to the frontend. This kind of pattern matched the nature of web application which allows the developer easily build serverless API for web/mobile applications on top of serverless cloud functions. The serverless application could also be triggered by changes in the database and object storage. This allows the serverless computing to be used as a background task such as data processing. A good example is the serverless computing use case of Thomson Reuters in their social media data analysis project[64]. Thomson Reuters uses AWS Lambda to host a serverless application that triggers when new data is stored. The application processes the data real-time, extracts the hashtag trend data and stores it in Amazon DynamoDB, a database solution by AWS, which is also serverless.

Managed Resources Allocation

Managed resource allocation means that developers only need to deploy code without leaving operational tasks to the cloud. As we mentioned before The developer does not need provisioning or managing any server besides, the developer is not required to install any software or runtime [65] when deploying his/her application.

The cloud provider manages the scaling of the infrastructure which the developers are running their code. This also reflects the managed resource allocation of serverless. In traditional virtual machines, although some cloud providers (such as AWS and Azure) support automatic scaling; however, the scaling strategy must be defined by the user. Moreover, the user needs to set up the cloud infrastructure (such as Auto Scaling Groups and Elastic Load Balancing in AWS) for using autoscaling. In contrast, in serverless computing, the cloud provider will handle everything related to automatic scaling. Furthermore, together with other operational tasks, the availability and security issues of the underlying infrastructure are being taken care of by the cloud provider as well.

Pay-per-use

Pay-per-use is the significant characteristic of serverless computing from non-technical perspective. The traditional cloud server using pay-as-you-go mode. The billing is done based on the type of VM and the rental time of this VM. For the user, this is not economically flexible, because the user must pay the same price for an idle VM as when it is fully loaded. On the contrast, in serverless computing, the users do not need to pay the idle time; they only pay for the time that the application is running. In many scenarios, such payment mode could lower the cost. According to [59], serverless computing could be 6x cheaper than VM when doing on-fly video encoding, with 60x performance. An organisation could save up to 4x-10x when moves application to serverless [59][66].

Extensive Application Scenarios

Serverless computing has a wide range of applications. A common application is to deploy the runtime in a serverless environment. However, as mentioned in section 2.4.1, serverless computing is now not limited to deploy cloud functions but used in all the components that could be used when building modern applications. For example, besides serverless functions (AWS Lambda), the serverless offering in AWS also includes the serverless database, container runtime services, data analysis and Kubernetes cluster. Google cloud also advocates "full-stack serverless" [67]. Like AWS, Google Cloud also provides various serverless solutions ranging from computing and DevOps storage to AI and data analysis. Azure's serverless products also cover a wide range of backend components, including computing, storage, artificial intelligence, monitoring and analysis [68].

2.4.3 Limitations

Serverless computing is not the perfect solution. In some aspects, it still has its limitations.

Performance

This is mainly the problem within the computing task that runs serverless. In the current serverless products of cloud providers, the computing power of serverless computing is limited. For example, in the virtual machine service (AWS EC2) provided by AWS, users can choose virtual machines with up to 96 CPUs and 192 GB RAM. In the serverless AWS computing engine, the maximum RAM size allocated is only 3008MB [69], and the maximum number of vCPUs is not specified in the document. By making it unsuitable for heavy tasks, it limits the application scenarios of serverless computing to the development team. In some cases (such as ML model training), the limitations of hardware selection are also mimicking performance. Research experiments [59] at the University of California, Berkeley show that because AWS Lambda does not support GPU computing, makes it 21 times slower than EC2 instances using GPU [70] when training deep learning models. In this case, longer execution times could make serverless servers more expensive, and the research also shows serverless has poor performance in MapReduce and linear algebra computing. In conclusion, as a development team, selecting serverless computing means limited hardware option and poor performance with a high cost in some scenarios.

Cold Start

The cold start is also a disadvantage of serverless. In when running a function on serverless cloud service, the functions are being served by container [71]. As long as the functions keep being triggered, the container which hosting the functions will stay active. The cold start means the trigger event happens when the function is not

being triggered for a too long time, so the cloud provider has already deactivated the container. In such a situation, the cloud has to deploy the code again and spin up a new container. This will significantly add overhead to the total execution time. Thus, if the development team needs to run a short task frequently, but not so frequently to keep the cloud function "warm", serverless is not the best option. This is because the cold start time could take even longer than the actual runtime, which will lower the performance.

Fortunately, for AWS Lambda, some plugins exist to solve this problem. The common practices of these plugins are to use CloudWatch to ping the function periodically. However, for other serverless services (such as AWS Fargate), there is no way to significantly shorten the cold start time.

Communication Pattern

The communication pattern between serverless services is limited: In current serverless computing offering from cloud providers, there is a lack of peer-to-peer networking between different running serverless instances [70]. This means some heavy lifting inter-communication such as streaming content to another function [72] cannot be done efficiently. For example, in AWS Lambda, replacement of peer-to-peer networking between executing cloud functions is through slow cloud storage [70]. While the communication between virtual machines is through the network interface, which is much faster than cloud storage. Such limitation could further affect the performance of the distributed system that hosted by serverless since the distributed algorithm largely depends on the communication between nodes.

Another limitation is the communication pattern of serverless leads to more inter-instances communication. A good example is the MapReduce [59]. While in VM part of shuffle and aggregation operation could be done within a VM instance, but between different tasks, such operations in Lambda must require inter-instance communication, since each task is on an independent instance. This problem largely increases the need for network communication. The experiments from UC Berkeley shows that during MapReduce operation, the serverless functions cost 15% more than VM [59].

Chapter 3

Overview of Current Serverless Cloud Services Offering in AWS

In this chapter, we will introduce the serverless cloud services in Amazon Web Services (AWS) that we will use in the DevOps toolchain and the experiments.

In Section 3.1, we introduce the AWS Elastic Container Services, the container orchestration services and AWS Fargate, which allow user runs containers with ECS in a serverless way. In Section 3.2, we introduce the AWS Lambda, the serverless cloud function service. We introduce AWS CloudWatch, the monitoring service of AWS in Section 3.3. Lastly, we introduced AWS DevOps tools in Section 3.4.

3.1 AWS Elastic Container Services with Fargate

In our toolchain, we make use of Elastic Container Services with Fargate to run the Jenkins build agent. In this section, we introduce ECS and Fargate, and how do they combine to host the serverless Docker container.

3.1.1 AWS Elastic Container Service(ECS)

Amazon Elastic Container Service is a managed container orchestration service that runs Docker containers. AWS fully manages the ECS service, which means that AWS will be responsible for some operational tasks, such as automatically scaling the running container. To introduce how the container runs on ECS, we first introduce a few concepts.

Task: Task means a container instance that runs in the ECS cluster. A task is defined by task definition, which is a JSON file that contains the following information: container definition, network, hardware configuration and launch type. The task is the instantiation of a task definition [73]. The ECS task scheduler is responsible for putting the task to the cluster.

Service: A service is an abstraction of a set of tasks that include a specified number of tasks runs simultaneously.

Launch Type: Launch type defines on which infrastructure the task will run. Currently, there are two options, EC2 (virtual machine) and Fargate (serverless). The EC2 launch type refers to running the container (task) in a group of EC2 virtual machines. This launch type requires the user to manually create and managing EC2 VMs. The Fargate launch type means run containers in AWS Fargate, and Fargate is a serverless container service in AWS. This launch type does not require user provisioning and managing the infrastructure that runs the containers. Instead, AWS takes over these tasks. Our first experiment in Chapter 5 will be related to a comparison between these two launch types.

The workflow for running a container in ECS is as follows: The first step is to have the task definition to define the specification of the Docker container that is going to run. In the next step, a task defined by this task definition is created.

In the DevOps toolchain, ECS can be used to host Docker-based build agents in the continuous delivery pipeline. Docker-based build agent means running certain stages in the pipeline distributed in Docker containers. ECS supports the use of APIs to perform the two steps we mentioned in the last paragraph, which makes it easy for DevOps tools to deploy Docker-based build agents to ECS cluster. Major continuous delivery tools support the Docker-based build agent. We will thoroughly introduce the Docker-based build agent in Section 4.2.

3.1.2 AWS Fargate

AWS Fargate is a serverless container service by AWS, and as we mentioned above, one of the launch types of ECS. It removes the need for provision, manages the server from the user's side. Fargate also follows the payment mode of serverless computing, which is paid for the runtime of each running container.

We notice that different from other serverless computing services in AWS, for example, AWS Lambda, Fargate cannot be used independently. To use Fargate, the user needs to select it as "Launch type" in Elastic Container Service, which means the container runs under this task definition will run in Fargate. We call this method of using Fargate as "Elastic Container Service with Fargate" in the following chapters. Another way of using Fargate is to run pods in Fargate when deploying Kubernetes cluster to AWS Elastic Kubernetes Service (EKS). Kubernetes is an open-source software which is used for orchestrating and managing container in a cluster environment. Elastic Kubernetes Services helps to runs a Kubernetes cluster on AWS infrastructure, and AWS takes care of the operational task such as monitoring and maintenance.

3.2 AWS Lambda

AWS Lambda is AWS's first serverless service. It was launched in November 2014. In AWS Lambda, users can upload codes called "Lambda functions" to AWS Lambda. AWS Lambda runs the code in its own hosting infrastructure. "Managed" refers to AWS performing all management tasks of back-end services, including server and OS maintenance, server configuration, and scaling. In addition to server-related tasks, AWS will also be responsible for security, monitoring, and logging.

AWS Lambda is event-driven, which means that when an incoming event triggers the function, the deployed lambda function will start running. We introduced the characteristics and applications of event-driven in Section 2.4.2. In addition, AWS allows user to associate Lambda functions with other AWS services. This means that changes in AWS services can be used to trigger our Lambda functions. The combination of event-driven characteristics and association with AWS services allows user to extend the functionality of AWS services. We will introduce how to use this combination in Chapter 4.

3.3 AWS CloudWatch

As we mentioned in Chapter 2, monitoring is one of the DevOps practises. AWS CloudWatch is a monitoring and observability service [74]. It is providing an out-of-box monitoring solution for both infrastructure and deployed applications. CloudWatch could also help on resource utilization and gives a uniform platform to monitoring operational health of the infrastructure in both AWS and on-premises. In addition, CloudWatch gives complete visibility to AWS infrastructure status, because CloudWatch natively integrates with over 70 AWS cloud services [74].

The core function of CloudWatch is to collect metrics and logs of all running AWS services under the current user, display these data in real-time and save the data for further analysis. CloudWatch supports monitoring all services running in serverless and server-based AWS. In addition, it can be used to monitor on-premises services. Monitoring in CloudWatch follows the following workflow:

1. **Collect:** CloudWatch gathers the log from services in AWS. In addition, it also gathers metrics include CPU/RAM utilization, network I/O, e.g..
2. **Monitor:** CloudWatch visualizes application and infrastructure logs and metrics on the dashboard. Users can check the status from the dashboard and can also set CloudWatch alarms.
3. **Act:** CloudWatch continuously monitors the status of AWS services. When certain metrics reach the value set in the CloudWatch alarm, the alarm will trigger the action set by the user. A common use case is to set the alarm about CPU usage and use that alarm to trigger auto-scaling. Alarm actions may also trigger Lambda functions.

4. **Analyze:** CloudWatch can save logs and analyze them later. The analysis includes customizable indicators, contributor insights and logs analysis.

We will introduce how the CloudWatch is being used in our toolchain in Chapter 4.

3.4 AWS Developer Tools

AWS provides a set of cloud-based tools which helps the user to build an integrated DevOps toolchain. These tools include the following four tools.

3.4.1 CodeBuild

CodeBuild is a fully managed build server in AWS. CodeBuild mainly takes care of the automated build and automated testing within configuration delivery. Same with all serverless services, CodeBuild frees the software team from building and managing build servers.

Although as a managed service, still, CodeBuild provides the user with a configurable build environment. Users are allowed to select the hardware configuration of the build machine. CodeBuild provides several out-of-box build environments which include build dependencies for the project in different programming languages. For example, a Java environment including JDK and Gradle; PIP and Python for Python development; Android build environment, etc.¹ Users can also use a self-defined build environment in compatible with their requirement. Furthermore, CodeBuild provides good integration with popular tools. For example, CodeBuild can be integrated into a continuous pipeline in Jenkins by acting as a Jenkins build agent. This could be achieved through Jenkins' plugin "AWS CodeBuild"² developed by AWS CodeBuild engineering team.

3.4.2 CodeDeploy

CodeDeploy is for automating the application deployment to both AWS services and on-premise services. Besides the basic functionality as automated deployment, CodeDeploy also minimized the downtime by using advanced deployment strategies (blue/green deployment and rolling update) and continuous health checking. CodeDeploy also allows users to continuously monitor the running status of deployed applications.

3.4.3 CodePipeline

CodePipeline is for modelling the workflow within the continuous delivery pipeline with both graphic interface and code. The user could use different DevOps tools

¹The full list can be found at <https://docs.aws.amazon.com/codebuild/latest/userguide/build-env-ref-available.html>

²<https://plugins.jenkins.io/aws-codebuild/>

from AWS or third party in each stage of the CodePipeline. In the other way around, CodePipeline could connect different DevOps tools we mentioned above into an integrated continuous delivery pipeline. These tools include AWS DevOps tools such as CodeCommit, CodeDeploy, CodeBuild; Third-party tools such as GitHub, Jenkins, XebiaLabs etc. ³

3.4.4 CodeStar

CodeStar is a uniform platform that joins AWS DevOps tools as an integrated DevOps toolchain. We mentioned CodePipeline above which integrated different tools to an integrated continuous delivery pipeline, while CodeStar brings the integration to a further step. We will introduce what tools does CodeStar include in Section 4.3.5 where we present our implementation of integrated DevOps toolchain in AWS.

³See <https://aws.amazon.com/codepipeline/product-integrations/>

Chapter 4

Design of DevOps Toolchains

In this chapter, we introduce the design and implementation of both toolchains (integrated and non-integrated) and explain how we come to this implementation. We will also compare these two types of toolchains within the scope of functionality and ease of implementation. Note that for the experiments that are answering our two research questions in the next chapter, we implement two different continuous delivery pipelines design with two sets of tools respectively, one with tradition non-integrated tool while another one with the serverless integrated DevOps tools from AWS.

In Section 4.1, we present a case software project that uses DevOps toolchains in the experiments. In Section 4.2, we introduce the design and implementation of our non-integrated DevOps toolchain. Section 4.3 is related to the integrated toolchain, and Section 4.4 is a comparison between integrated and non-integrated toolchain. Lastly, in Section 4.5, we talk about the challenges we met during the implementation.

4.1 Case Project

We first develop the case project. The case project is an example software project which will be used to test our implementation and run the experiments in which we simulate the DevOps development process of the case project on our DevOps toolchain. The case project is a REST API service. We use API services as case projects because such services are an essential part of modern software projects with microservices architecture. Although the type of case project has no impact on our DevOps toolchain at the architectural level, the build dependencies and software configuration in the toolchain may be affected. Thus we need to have an introduction to the case project.

We choose Java as the programming language used for the case project because it's popularity and versatility. Java is one of the most common languages used in commercial software development. According to the TIOBE index of programming language [75], Java is the most popular or second most popular programming

language in the world since the mid-1990s. Besides commercial software development, the Java programming language is also widely used in open-source software development. The report [76] from GitHub shows that Java ranked third in 2019 and second before 2018. One of the main applications of Java in Web development. Currently, 7 out of every 10 [77] most popular websites in the world use Java as the web development language (server-side). Furthermore, Java has good versatility, which means that it is suitable for almost all types of applications. For example, web applications, desktop applications, and besides, Java is the main development language for Android applications.

Java programming language has a whole ecosystem that can be used to improve software development practices for adopting DevOps. These tools include: build, code analysis, testing frameworks, artifact management, build automation & dependency management et. These tools could be easily integrated and act as part of the DevOps toolchain.

In term of developing REST API with Java, the Spring is the most popular framework and has been used in many major Internet companies, including Google, Microsoft and Amazon [78]. So, we choose Spring as the framework to build our application. To develop our Spring application, we use Spring Boot¹. Spring Boot is a project under Spring, which, according to its documentation, is to allow the developer to create Spring application with the minimal effort [79], by simplifying the configuration of Spring framework.

```
Method: GET
Endpoint: /packages
Success Response:
Code: 200
Content:
[
    {
        name : (Package name)
        description : (Package description)
        dependencies : (Dependencies)
    }
]
Error Response:
Code: 500
Content: { msg: Server Error! }
```

Figure 4.1: RESTful API Interface of Case Project

The case project is a simple REST API (Figure 4.1) which returns the info of all

¹<https://spring.io/projects/spring-boot>

installed software packages in the host machine in JSON format when the frontend sends an HTTP GET request to the backend.

4.2 Design of Non-integrated DevOps Toolchain

In this section, we introduced the design of a non-integrated DevOps toolchain. We first introduce the considerations when choosing the tools to build the toolchain. When introducing each component of the toolchain, we will also introduce how the component uses serverless computing.

4.2.1 Architecture

The toolchain implementation is based on the DevOps elements we presented in Chapter 2. Figure 4.2 shows the architecture of our DevOps toolchain. Here we are only presenting architecture on a more general level. The detailed architecture of each component will be introduced in the following sections, both text and graph. From Figure 4.2, we can see except version control, and the whole environment is running in Amazon Web Services. Due to the limitation of space, the internal architecture of certain components is not shown in the graph. Instead, we will show them in the following sections.

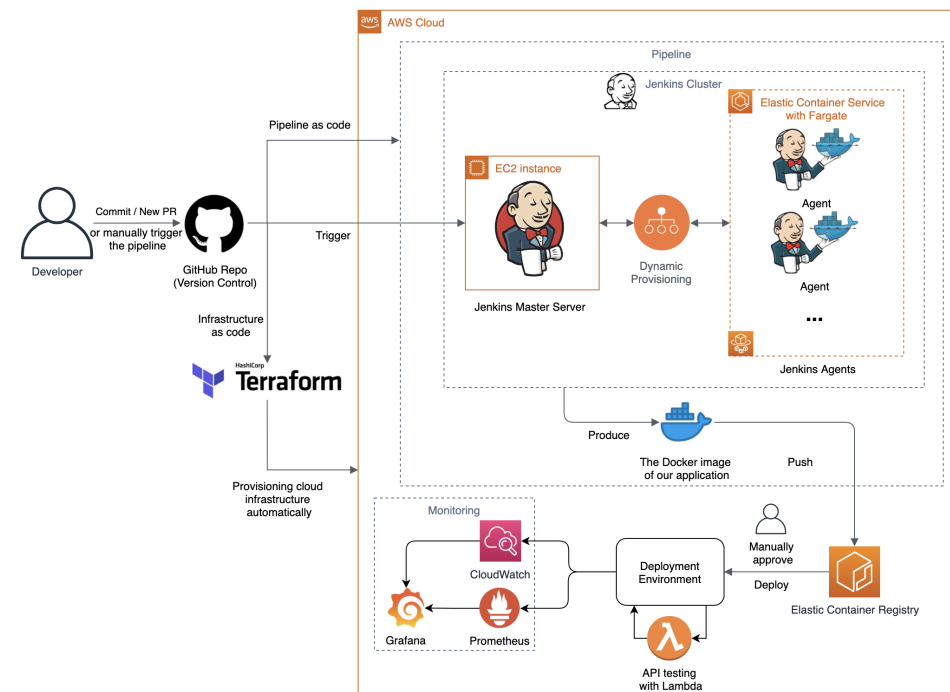


Figure 4.2: Architecture diagram of the non-integrated DevOps toolchain

When the developer pushes a new commit to the repository in GitHub ², GitHub will send HTTP POST request that contains the necessary information to the Jenkins master node. Jenkins master, which triggered by the HTTP request, will create a new job for this project according to the information that the HTTP request contains. The job will first pull the latest code from the git repository, then runs the docker containers with required build environment and build the project. In the end, a docker image for running the project will be created and be pushed to the container registry of AWS. The project will be deployed in the last step.

4.2.2 Introduction to Tools Used in the Implementation

One of the essential steps to build the non-integrated toolchain is to select the proper tool for each component. In this section, we describe our consideration when we select tools.

Continuous Delivery Pipeline

The most popular server-based tools for build continuous delivery pipeline are Jenkins³, Drone⁴, GoCD⁵ and Circle CI⁶. Table 4.1 shows a comparison between these tools. As we can see from the table, Jenkins is the most popular option for CI/CD. Jenkins has wide application in the commercial use case, and the high popularity in the open-source community as well. Although compared with the other three newer tools, Jenkins is more focuses on the "Build" step within the continuous delivery pipeline since it was originally a build automation server. But, the open-source nature of Jenkins gives it a much wider selection of the plugin, which means Jenkins can be used for almost all steps in a continuous delivery pipeline.

	Jenkins	Drone	Circle CI	GoCD
Open Source	Yes	Yes	No	Yes
GitHub stars	15.7k	21.2k	-	5.7k
Github contributors	614	258	-	116
Plugin extensions	Over 1500 ⁷	93 ⁸	110 ⁹	88 ¹⁰
Price of self-hosted solution	Free	Free	\$35 user/month	Free
Number of companies use it in the tech stack ¹¹	2634	82	1368	42

Table 4.1: Comparison of continuous delivery tools

²<https://github.com/>

³<https://www.jenkins.io/>

⁴<https://drone.io/>

⁵<https://www.gocd.org/>

⁶<https://circleci.com/>

Created by Kohsuke Kawaguchi in 2001 as "Hudson", Jenkins is an open-source automation server written with Java, which helps to automate different parts in the software development life cycle. It is suitable for a team of all sizes and varies of languages and technologies [80]. Furthermore, Jenkins attracts software teams with its ease of use and high extensibility [80] and thousands of plugins. Since Jenkins has an active open source community, more plugins are being created and maintained. These plugins can help Jenkins keep up with the rapidly evolving DevOps practices and help Jenkins integrate with emerging tools and cloud services. extensibility makes Jenkins still the most popular tool in the DevOps toolchain, even if it is an aged software created when the term "DevOps" appeared.

Our continuous delivery pipeline is developed with Pipeline plugin¹² in Jenkins. Pipeline plugin allows us to define a continuous delivery pipeline as code in Jenkinsfile. In the pipeline, a conceptually distinct subset of tasks within the continuous delivery pipeline [81] is defined as a "stage"¹³ and each task within a step is called "step". Each pipeline is binding with a "project". An execution runtime of a project/pipeline is called "build", and the machine (virtual machine, container, etc.) for running the build is called "agent".

Build & Test Automation Tool

Within this project, we use Gradle¹⁴ as the build tool. Gradle is a powerful build tool that was originally designed for JVM-based languages, but now it also supports other programming languages such as C++ and Python. Like Jenkins, Gradle also has a dynamic ecosystem with a large number of plugins. This makes it possible to use different types of tools (such as unit testing and code analysis) in a single pipeline in Gradle. Moreover, Gradle makes dependency management easy, and dependencies can be easily added to the project by editing the project's Gradle configuration file. In addition, Gradle supports the configuration as code. This allows developers to define all build configurations of a software project in one file.

For unit testing within the build stage, we are using JUnit¹⁵ as the tool for testing. JUnit is a test framework which allows us to specify unit test by within @Test annotation within the Java Code. The JUnit executes the unit test amount three approaches of test automation we mentioned at 2.2.1. The "unit" means the smallest component of software, usually a method in practice [82]. A unit test is a testing approach which checks if each part of the software works properly. This is

⁷<https://plugins.jenkins.io/>

⁸According to GitHub search result

⁹<https://circleci.com/integrations/>

¹⁰<https://www.gocd.org/plugins/>

¹¹based on data from StackShare

¹²<https://www.jenkins.io/doc/book/pipeline/>

¹³For example, "Build", "Test", "Deploy" step in a continuous delivery pipeline.

¹⁴<https://gradle.org/>

¹⁵<https://junit.org/>

an early test conducted in the DevOps life cycle, and it will be conducted during the build. Therefore, we embed JUnit into the Gradle build process. Using JUnit for unit testing can improve the confidence of developers when developing software [82]. Frequent and early unit testing enables developers to find errors earlier. The earlier errors are found, the easier it is to fix them. Furthermore, the unit test helps the software team to be confident that the new change that passes the test will not break the existing product. This could make the team be more dare to make the change, do faster releases and be agile.

For code analysis, we use SonarScanner for Gradle¹⁶. SonarScanner provided an easy way to embed SonarQube code analysis into the Gradle build process. This means in our implementation, and we use the tool to do static analysis on our Java code to build our project with Gradle. Static analysis means the code analysis tools will analysis our source code but without build and execute the code. The analysis objects include software structure, security, code quality, etc. After the analysis is done, a report generated by SonarQube could give a software development team complete overview of the code quality issues. SonarQube also quantifies the code quality by gives SQALE rating and technical debt ratio to the code. The quantitative code quality metrics make it easy for the team to enforce the quality policy. A common way to enforce the quality policy with metrics is to have a quality gate in the pipeline, which decides the project is "good enough to deliver" when metrics reach a certain threshold. The code analysis, on the one hand, automates the code review process, and enforce code quality policy, on the other hand, helps to maintain code quality in the very early of the DevOps pipeline. Thus, code analysis reflects the philosophy of DevOps: automation, fast software delivery with high quality.

Deployment and Jenkins Agents

We will widely use Docker¹⁷ in our pipeline. Docker is an open-source software which could pack, deliver and run the software as a container with OS-level visualization. Docker is developed by Docker, Inc. and published at PyCone in California, the USA in 2013. In March 2013, Docker became an open-source software [83]. Docker has high compatibility that allows the user to run Docker in any major operating systems(macOS, Linux and Windows) with both X86 and ARM CPU architectures.

To software teams, Docker eases the environment configuration task by replacing VM with lightweight containers. A container is a separate unit that includes the application and all its dependencies which allow application runs in the same way regardless of the host environment [84]. A container is the running instance of a Docker image that defined by Dockerfile. Compared with VM, which emulate the hardware and wraps up the whole operation system, containers only includes the dependency needed to run an application. Multiple containers could run on

¹⁶<https://docs.sonarqube.org/latest/analysis/scan/sonarscanner-for-gradle/>

¹⁷<https://www.docker.com/>

tops of one Linux instance. This means the container is lightweight and cost much fewer resources. With the same hardware, the team could run 4x-6x numbers of application in the container than with VM [85].

Docker fits well with DevOps since it allows the software team to quickly create local development environments that simulate different production environments in a single machine. For example, by creating multiple containers with different operating systems on the build machine, the development team can build C/C++ applications for different OS (Windows, macOS) only on the build machine with Linux system. In addition, Docker allows the team to deploy to the cloud more easily [85]. Furthermore, Docker is designed so it could easily corporate with many DevOps tools [86]. In short, Docker can easily meet the needs of DevOps: it can easily run multiple development environments on the same host, is easy to deploy to the cloud, is very recourse-saving compared with VM, and fits well with DevOps tools.

Docker allows us to specify all system dependencies in a single file (Dockerfile). Dockerfile defined the built environment as code. This allows us to manage the build environment with the version control system, and to test and do quality analysis to the environment defined by Dockerfile. In short, the Dockerfile could improve the automation within the system and help to apply DevOps practices on the build environments.

There will be two main use cases of Docker in our toolchain. Firstly, we run the build stage within the container. This means the pipeline will execute specific steps inside ephemeral Docker containers [87]. It is easier to manage build dependencies in the Docker container. Besides, the container-based agent requires less effort to maintain.

In our case, to build the case application, the host machine needs to have JVM installed. However, we want to make our pipeline not only suitable for Java application but also easily be used to build an application in other programming languages. Docker solves this problem by provides excellent isolation from the host machine. Thus, we can configure the built environment (operating system version, dependencies) runs within a Docker container without actually install anything on the host machine by merely editing the Dockerfile.

Second, we use Docker to Dockerize our application which creates a Docker image of our application. The first reason for using Docker is that Docker reduces the operational effort. With Docker, there is no need to pre-install any Java environment in the environment to run our application. This is because all environment is already being packed in our Docker image. The second reason, Docker improved compatibility of the case project because Docker ensures that the application packed container can run in the same behaviour no matter what host it is running on. The last reason, all major cloud computing providers support Docker. We can easily run the container on the cloud services. This means that our Dockerized applications can be easily cloud-native and can be deployed in a multi-cloud environment. For example, in AWS, there are Elastic Container Services (introduced in Chapter 3) specifically for container orchestration.

Monitoring

We are using CloudWatch for monitoring the status of the cloud infrastructure. CloudWatch is a tool provided by AWS, which we already introduced in Chapter 3. However, CloudWatch cannot give us insight about the status of Java Spring-Boot application. More generally, the metrics within the application framework. In Spring Boot, these metrics are HTTP statistics, CPU load and JVM statistics. Prometheus¹⁸, together with Grafana, could fill this gap.

Prometheus is an open-source monitoring and alerting solution initially built by SoundCloud in 2012 [88]. Prometheus is used for reading numeric metrics that are recorded in time series. In Spring Boot, there are some plugins, for example, Micrometer¹⁹ exist which could export all the Spring Boot specific metrics to Prometheus in time series. Thus Prometheus has a perfect fit with our case project. In addition to collecting metrics, we could also set the alarm within Prometheus, which could alert the user when some metrics are not in the normal range. we can also query the metrics from the past. Although Prometheus supports simple graph which shows the metrics' change with time, it is not user friendly enough. Thus, we introduce Grafana to better visualizes the data collected by Prometheus.

Grafana is an observation platform in which it's core feature is in the visualization. Users could define dashboards according to their needs with JSON files. Another main feature if Grafana is that it could gather data from the different platform into one dashboard. In our case, with Grafana, we can display the metrics from Prometheus and CloudWatch in a single page. This could give us a great overview of both Spring Boot application and cloud infrastructure at the same time.

4.2.3 Infrastructure as Code (IaC)

Infrastructure as code the common practices to implement configuration Management in the cloud-based environment. And Configuration management is one of the components of the DevOps toolchain that we mentioned in Chapter 2.

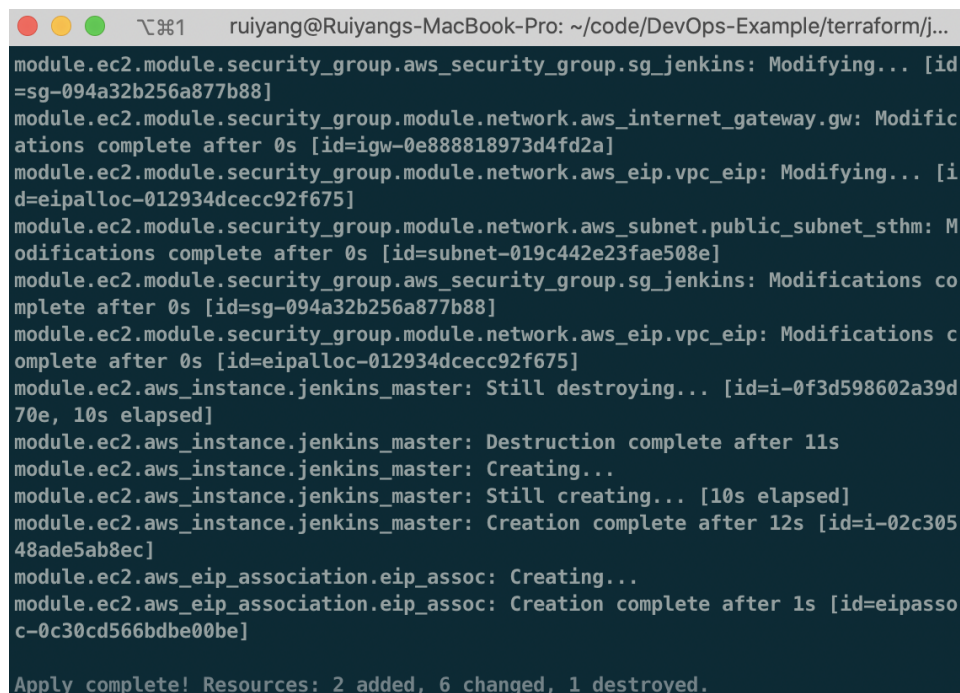
Terraform²⁰ is one of the most popular tools to manage cloud Infrastructure with infrastructure as code practice. Infrastructure as code enables the software team to version control and test the infrastructure to track its change. Furthermore, it enables the team to apply DevOps practice, such as test automation on the infrastructure.

Terraform is not limited to defined infrastructure as code; it also helps to do the cloud infrastructure orchestration. Compared with tools like Ansible, which the code defines each step of the automated process, Terraform code only define the final status of the infrastructure. This makes the code more understandable. Also, Terraform is portable, means it support infrastructure orchestration of all major cloud providers

¹⁸<https://prometheus.io/>

¹⁹<https://micrometer.io/>

²⁰<https://www.terraform.io/>

A terminal window with a dark background and light text. The window title bar shows three colored circles (red, yellow, green) and the text 'ruiyang@Ruiyangs-MacBook-Pro: ~/code/DevOps-Example/terraform/j...'. The terminal output shows Terraform CLI progress for various AWS resources. It includes status updates for security groups, internet gateways, EIPs, subnets, and EC2 instances. The final line indicates the application is complete with 2 resources added, 6 changed, and 1 destroyed.

```
ruiyang@Ruiyangs-MacBook-Pro: ~/code/DevOps-Example/terraform/j...
module.ec2.module.security_group.aws_security_group.sg_jenkins: Modifying... [id=sg-094a32b256a877b88]
module.ec2.module.security_group.module.network.aws_internet_gateway.gw: Modifications complete after 0s [id=igw-0e888818973d4fd2a]
module.ec2.module.security_group.module.network.aws_eip.vpc_eip: Modifying... [id=eipalloc-012934dcecc92f675]
module.ec2.module.security_group.module.network.aws_subnet.public_subnet_sthm: Modifications complete after 0s [id=subnet-019c442e23fae508e]
module.ec2.module.security_group.aws_security_group.sg_jenkins: Modifications complete after 0s [id=sg-094a32b256a877b88]
module.ec2.module.security_group.module.network.aws_eip.vpc_eip: Modifications complete after 0s [id=eipalloc-012934dcecc92f675]
module.ec2.aws_instance.jenkins_master: Still destroying... [id=i-0f3d598602a39d70e, 10s elapsed]
module.ec2.aws_instance.jenkins_master: Destruction complete after 11s
module.ec2.aws_instance.jenkins_master: Creating...
module.ec2.aws_instance.jenkins_master: Still creating... [10s elapsed]
module.ec2.aws_instance.jenkins_master: Creation complete after 12s [id=i-02c30548ade5ab8ec]
module.ec2.aws_eip_association.eip_assoc: Creating...
module.ec2.aws_eip_association.eip_assoc: Creation complete after 1s [id=eipassoc-0c30cd566bdbbe00be]
Apply complete! Resources: 2 added, 6 changed, 1 destroyed.
```

Figure 4.3: Creating a cloud environment with Terraform CLI

In our implementation, we define our cloud infrastructure and all AWS resources, including EC2 virtual machine, ECS cluster, security groups and network Infrastructures in a series of configuration files. Then we create the cloud environment by simply using CLI interfaces. Figure 4.3 shows the creation of the cloud environment with Terraform.

4.2.4 Version Control

Version Control System (VCS) is the process that record the changes in source code set over time [89], and versioning the history of these files. Version control can not only apply to the software source code, but also the cloud infrastructure, build configuration, Docker images, continuous integrating pipelines that we define as code in our project. VSC is suitable for track the development progress and manages the goal within a software development team [90]. Among all software for version control, Git is the most popular one nowadays. Git is a distributed version control software created by Linus Torvalds. Git is based on the command-line interface (CLI), which allow the user to execute commands in different environments, for example, within continuous integration pipeline, within Gradle build or within a Docker container.

The survey [91] from Synopsys shows that in 2019, 71% of the project today is using Git as it is versioning system while SVN that ranks in second only be

used in 25% of the projects. We use Git as the version control system since it is used by most of the software development teams nowadays. We use GitHub for hosting the case project. Github is the biggest platform in the world that hosting a version-controlled software project for free using Git. It provides interfaces with different DevOps related tools which makes it easy to be integrated into all kinds of DevOps toolchains. The setup of Git includes the client, which is the development machine, and the remote server, which is the server that hosting the file history on the network. A branch means a diverge from main codebase which allows the developer to work on it without touch the code in the main codebase [92]. When the developer adds new code to on the feature branch, the code in the main branch (main codebase) might also be changed by other developers. And when the developer what to submit the change to the main branch, he/she has to uses merge. Merge is to combine two branches, by combining two sets of commits in two branches to a unified history [93].

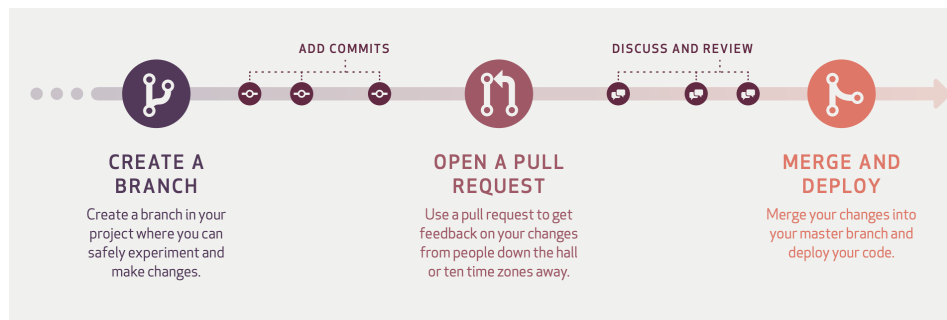


Figure 4.4: GitHub Workflow [94]

The Git flow [95] proposed in 2010 is a successful workflow for working with Git. Git flow has already been widely used and has been approved by the software industries. However, the git-flow is centred with "release" which makes it not suitable for a DevOps team which could deploy to production everyday[94]. To better cope with the frequent release nature of DevOps, the Github flow – a simplified version of Git flow is proposed by GitHub. Github flow is light and branch-based workflow that the team could follow for collaborating on a software project. Compared with git-flow, GitHub flow Therefore, we choose GitHub flow [96] the basis of our workflow in the project. The simplified version of this GitHub flow is shown as in Figure ??

Several general principles followed by us when adapting GitHub flow, we refer to principals in [96] to design our workflow.

- When working on the new feature, make a new branch for this feature. The name of this branch should be descriptive, which reflect the content of this feature. Commit the new code related to this feature to the feature branch. And push from this feature branch to the branch with the same name on the remote server (github.com) when necessary.

- Open a pull request²¹ when the feature is ready to merge, or when developer feel that he/she need help or comments from other team means on this feature. Others also do the code review in the pull request. Pull request allow others in a software team to inspect and comment on the change will be made before the developer merges his change to another branch.
- When the code is reviewed and is good to be merged, the developer should merge the code to the master.
- After the code of this feature is in the master, the code will and should be immediately deployed. There should not be any rollback in the master branch. If there are any issues within the newly merged code, a new commit or a new branch should be made to fix the issue rather than rollback on the master.
- Master branch is always deployable. This means when deploying the continuous delivery pipelines in our toolchain, only the master branch can be deployed to production. Moreover, there should not have any code which is not good to be deployed in the master branch.

Note that in our Git workflow, there are several time points that we need to run the continuous delivery pipeline within the toolchain. The continuous delivery pipeline will also vary with the time point within the version control workflow. We will introduce this in detail on Section 4.2.5.

4.2.5 Continuous Delivery Pipeline

Figure 4.5 shows the six Jenkins stages in our pipeline. The bottom part of this Figure shows the task distribution between the master node and agent nodes. The master node is an EC2 virtual machine while agents run on Fargate instances within an ECS cluster. In section 2.3.1 We mentioned that the development of the monitoring system should be in parallel with the main software project, thus in our pipeline, the build and deploy of the monitoring system is in parallel with the case project.

As we can see from the Figure, when the master node starts a job, it will create a Docker container in AWS Fargate as the agent. The agent will pull codes from VCS, build the code, and then send the build artifacts back to the master node. After this, the container will be terminated. The master node will continue the rest steps.

Build Agents Build agent is an independent computation unit (VM or Docker container) that could exchange data with the Jenkins master node and run a certain part of the pipeline. To implement a Jenkins build cluster, we need first to implement build agents. We discussed why we use Docker-based agent in our Jenkins

²¹<https://docs.github.com/en/github/collaborating-with-issues-and-pull-requests/about-pull-requests>




Figure 4.5: The Stages and Distributed Build in Our Pipeline

build a cluster on 4.2.2 and we decide to it in our implementation. The first step of our implementation is to develop our own Docker image²² of the Jenkins agent. We use the "jenkins/inbound-agent"²³ as the base image, this allows our Jenkins agent to establish an inbound connection to the Jenkins master with TCP. The next step is to set up the built environment within the agent. We add shell script for auto-install all build dependencies of our case project when we build this Docker image. In the last step, we build the Docker image for build agent and push it to our registry²⁴ in DockerHub.

We also discussed how Fargate allows us to run container serverless. To make use of Serverless offering of AWS, we let Docker-based Jenkins agents run on AWS Fargate to cut the operational effort and automate the scaling of Jenkins cluster. To implement this, we use Jenkins plugin "Amazon Elastic Container Service (ECS) / Fargate"²⁵, which is the Jenkins plugin allow us to host Jenkins agent in Fargate.

Considerations in Designing the Workflow of Distributed Pipeline The considerations behind to our design are that the first two steps take most of the time in our pipeline and according to Figure 4.6 runs more frequently than other steps, which the reason will be discussed in next section "Workflow in Production". The running time will be further extended when building a larger project. These two stages will be the bottleneck of the pipeline if we have it on the master mode. So we need to offload these steps to Jenkins agents for better performance.

²²The Docker image we developed could be found at <https://hub.docker.com/r/dry1995/jnlp>

²³<https://hub.docker.com/r/jenkins/inbound-agent/>

²⁴<https://hub.docker.com/u/dry1995>

²⁵<https://plugins.jenkins.io/amazon-ecs/>

The second reason is: As we mentioned in the introduction to Docker on 4.2.2, the built environment inside the Jenkins agent running in a Docker container is easier to change. When the team wants to build the release for different OS (Which happens in C/C++ development) or wants to have a different build environment for various projects, Docker could help to eliminate tasks such as configuration and installation different environment. With Docker, they can just modify the Dockerfile that defines the Docker image of the Jenkins agents. However, we cannot put the stage of builds the Docker image in Jenkins agents. This is because AWS Fargate does not allow Docker container runs in privileged mode, which means we cannot use Docker within the container that runs in Fargate. This is one significant limitation of Fargate. Instead, we had to move the step back to the master node. Fortunately, in our case project, Docker build only takes a short time (average 1s). Therefore, this will not slow down the entire pipeline.

We also notice that the Deploy stage also takes a long time. Still, we do not have it in the distributed build because: first, it is on the end of a pipeline so it will not block the further steps, second, the pipeline runs the stage less frequently than first two stages as shown in Figure 4.6. Thus there will be less possibility that there are many jobs runs at "Deploy" stage in parallel.

Workflow for Continuous Delivery Figure 4.6 shows our proposed workflow of a project that goes through the continuous delivery pipeline. We can see when the event on the feature branch triggers the pipeline, and it only runs through the first two stages. This is because according to the practices of continuous integration mentioned by us in 2.2.2 and by Martin Fowler in [27], a developer should merge (the "integration" in continuous integration) his/her work into main branch couple times per day. Therefore the code with this new feature runs through the whole pipeline at least several times a day. This already ensures the code could frequently be tested and deployed into the test environment.

The developer only commits to the feature branch. The pipeline runs first two stages after a developer pushes local commits to Git. It first pulls the newly pushed code, and then build. In the build stage, the code first is analyzed, then we do unit testing to make sure the code could pass the test cases defined by the developer during development. In the end, the code will be built into Java ARchive file (.jar). The purpose of putting the code analysis step before the build is that the code analysis will check for warnings, errors, and code quality so that we can ensure that the code is runnable without syntax errors and pass the quality gate before putting it in the build. This build will not run the code exists error or not passing the quality gate set by the team. This measure can reduce the cost by lowering pipeline runtime if code is not runnable or not quality acceptable.

If no error returns after finishing all the above steps, the developer can open a pull request view the code change and ready to merge the code to the dev branch. Before the merge, the pull request needs to pass the code review by another developer. Code review is to make sure that the automated tests do not miss any bugs. After

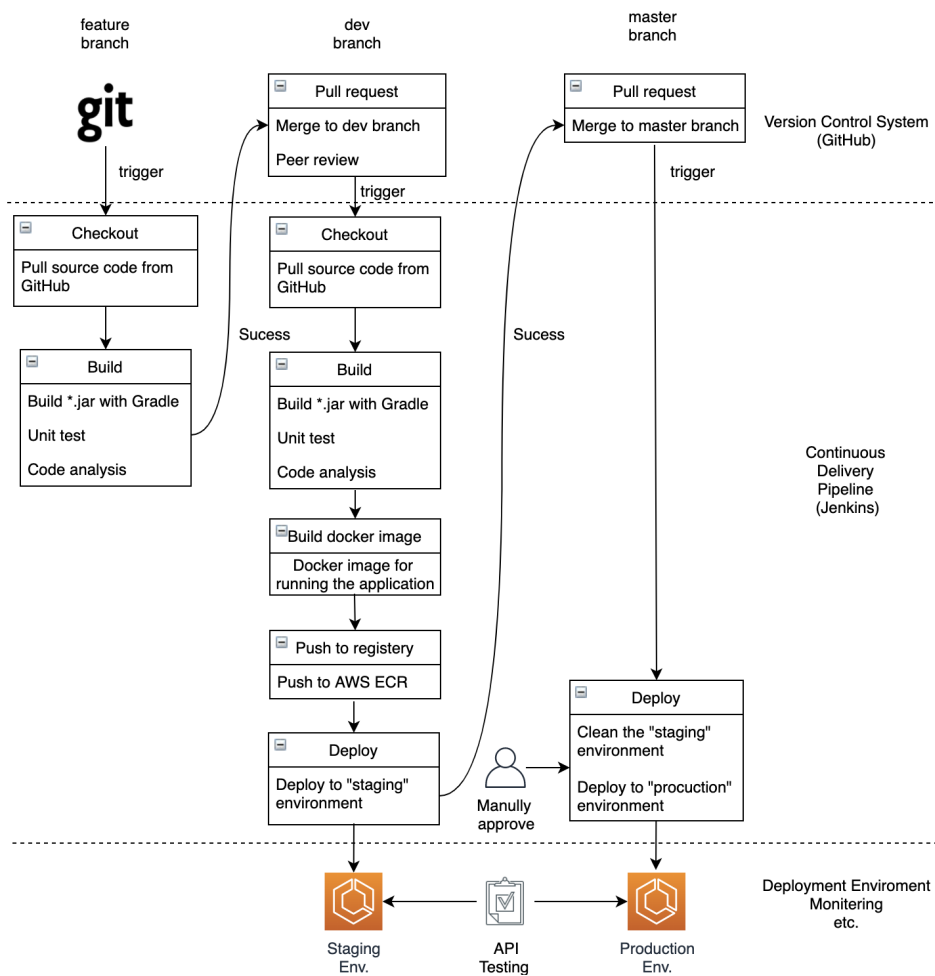


Figure 4.6: The Workflow of Continuous Delivery Pipeline in Our DevOps Tool-chain

the code review passed, the reviewer or the developer him/herself merge the code to the dev branch.

After the code merged to the dev branch, the pipelines run again, this time it runs the whole pipeline. First, the pipeline executes the first two stages as in the feature branch. Now we have the Java ARchive file. The Java ARchive is an executable package of our Spring Boot application. Next step is to Dockerizing our application which generates the Docker image our application. Then we push the image to the Amazon Elastic Container Registry AWS (AWS ECR) for further use.

The next step of the pipeline is deployment, the pipeline pull image in ECR that we pushed in the last stage, and then deploy it to the deployment environment in ECS with AWS CLI. The deployment strategy we are using is the rolling update. In the rolling update, we are gradually replacing instances in our deployment environment with the newer version of code.


```

Testing the url http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/
Testing the url http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages
Get result of url: http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/
1 out of 2 endpoints tested, 1 succeed
Get result of url: http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages
2 out of 2 endpoints tested, 2 succeed
All endpoints are tested, result:
{
  "http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/": "Succeed, response body (first 300
characters): Hello world",
  "http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages": "Succeed, response body (first
300 characters): [{\"name\": \"libws-commons-util-java\", \"description\": \"Common utilities from the
Apache Web Services Project\", \"dependencies\": [\"\"], \"reverseDependencies\": []}, {\"name\": \"python-pkg-
resources\", \"description\": \"Package Discovery and Resource Access using
pkg_resources\", \"dependencies\": [\"python (\\u003e= 2.6)\", \"python (\\u003c
}

```

Figure 4.7: API Test Result (Jenkins log output)

In the dev branch, the pipeline will deploy the application to the staging environment. The deployment to staging environment should be automated. This is because the staging environment is only for testing and only visible within the team. In the staging environment, we will conduct API testing (last stage shows in Figure 4.5) for test if our deployed API works and if it works as expected. The test is being done by trigger a Lambda function. The Lambda function sends a test HTTP request to the deployed endpoints, and verify if the HTTP response is correct. Figure 4.7 shows the test result of our case project. From the figure we can see if the test case is passed, and part of the HTTP response for manual inspection. If the deployed function passes the test, this shows the deployment works as expected and ready for the deployment. The developer could now open a pull request, merge code to master branch. The pipeline will deploy the application to the production environment, which is visible to the customers.

4.2.6 Deployment Environment

We create a simple deployment environment with AWS Elastic Container Service and Elastic Load Balancer. Same with Jenkins agents, we use Fargate to host our containerized case project.

AWS Fargate allow us to run our containerized application without having to manage servers, makes it easier for us to build a functionality complete DevOps toolchain implementation. We chose ECS instead of EKS (Elastic Kubernetes Service) because ECS is free, and EKS charges an extra fee for the running time of the cluster. Compared with EKS, ECS also provides better integration with other AWS services, such as AWS DevOps toolchain and AWS CloudWatch monitoring.

Figure 4.8 shows our deployment architecture. The deployment region in Stockholm (EU-north-1). The Fargate instances in ECS cluster are automated scaled according to the number of incoming requests. To improve the availability of the product, we deploy the case project into two different availability zones within the region. When one availability zone is down, the load balancer can route the request to ensure that the request can still reach the healthy availability zone. Besides




Figure 4.8: Deployment Environment

the availability improvement, the load balancer also distributes incoming requests across Fargate instances which maximizes the resources utilization rate within our ECS cluster.

4.2.7 Monitoring

Different from test automation which usually integrated with the continuous delivery pipeline, the monitoring is independent of the pipeline. Usually, monitoring does not act as one step within the continuous delivery pipeline but as an independent component.

In Chapter 3, we introduced AWS CloudWatch as one of the serverless services in AWS. In our toolchain, we will use it as a tool for monitoring the cloud infrastructure. With Cloudwatch, we can get the realtime log and quantitative metrics from our deployed container in the ECS. Figure 4.9 shows the monitoring dashboards in CloudWatch and Grafana.

In the last paragraph of Section 4.2.2, we mentioned that the CloudWatch is not capable of monitoring what happens inside the application framework, in our case, Spring boot. Thus, we use Prometheus to gather the metrics within our Spring Boot application. Prometheus gather the realtime metrics with the help of Micrometer plugin of Spring boot. Micrometer exports all numeric metrics in realtime; these metrics include JVM statistics, CPU/RAM utilization, number of request to each API endpoint and the user's self-defined metrics within Java code. Prometheus records these metrics every second, and save the metrics in time series to Database. In addition to Prometheus, we use Grafana to read the metrics from Prometheus and

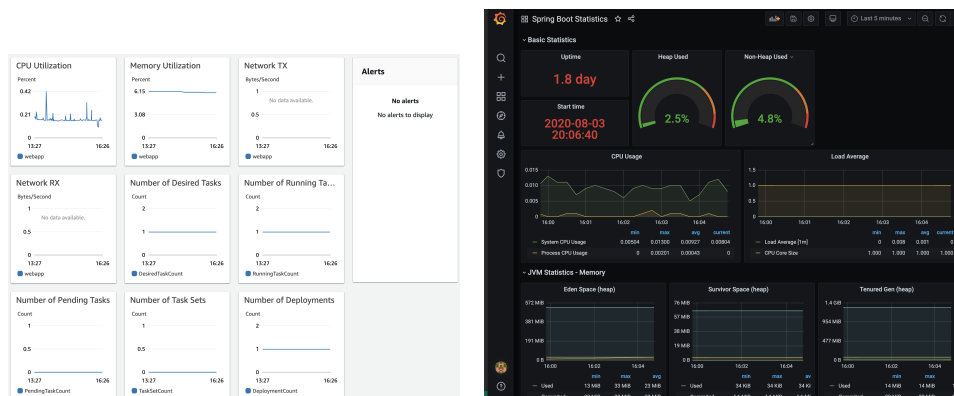


Figure 4.9: The Monitoring Dashboard of CloudWatch (Left) and Grafana (Right)

display the metrics on the dashboard.

Another service we introduced in Chapter 3 is AWS lambda and discussed how could it be used in our DevOps toolchain in which monitoring is one of the use cases. In our monitoring system, AWS lambda is used as an extension for CloudWatch, and we use it for two cases.

Auto-Scaling the ECS Cluster with Custom Alarm in Cloudwatch As we mentioned in 4.8 Deployment Environment, The deployment could be auto-scaled by defining the auto-scaling policy within the ECS cluster. Such auto-scaling in practice is: When the watched resources utilization is above/below a certain threshold, an alarm in Cloudwatch will be triggered. The alarm will further trigger the scaling event if the scaling policy was being set before.

According to Luca Tiozzo's article [97], with Cloudwatch alarm based scaling, the developer team has to use two groups of alarm watching RAM and CPU, respectively. When the ECS cluster lack of CPU resource but not lack of RAM resource, the CPU alarm is triggered, and the ECS scaled up. Now the ECS cluster has enough CPU resource, but the problem is, it may have too much RAM resource so that it triggers the RAM alarm to scale-in. So the cluster will scale in again. This will cause the cluster to keep scaling up and back without finding a suitable size.

A good practice solves the problem is to use a single group of alarm that only triggered by single metrics [97]. The developer team can set an AWS lambda function that read different metrics and then aggregate it to a single custom metric. The software team determines the threshold and metrics according to the deployed project. Once the aggregated metric reach the threshold, the lambda function triggers an alarm that can trigger the scaling of ECS cluster.

Custom Project-Specific Metrics The second application scenario is related to the first one. The Cloudwatch has support on resource utilization metrics. How-

ever, some metrics are project-specific and not related to resources utilization and performance. For example, the number of successful payment has been made in a payment service. In such a case, Lambda could fill the gap within the scope of CloudWatch. The team could set up a Lambda function which gets the number by monitoring the log with PutMetricData provided by CloudWatch. This Lambda can further forward the metrics to metrics analysis and visualization platform, for example, Grafana ²⁶ to give the management team an overview of the KPIs.

4.3 Design of Integrated Serverless DevOps Toolchain

AWS provides a set of serverless DevOps tools which could help us build a completely serverless DevOps toolchain. We introduced these tools in Chapter 3. In the section, we introduce the design of Serverless toolchain based on DevOps tools of AWS. Part of the toolchain is the same as the non-integrated DevOps toolchain in the previous section: some such as CloudWatch are already AWS serverless tools, some such as GitHub cannot be replaced with AWS tools, and some are tools embedded in the build pipeline, such as Gradle. Therefore we will not introduce these components but will focus on how do we make use AWS DevOps toolchain. Figure 4.10 shows the general workflow of a project delivered by our integrated DevOps toolchain.




Figure 4.10: Integrated Serverless DevOps Toolchain

²⁶<https://grafana.com/grafana/>

4.3.1 Continuous Delivery Pipeline with AWS CodePipeline

The workflow of our continuous delivery pipeline is the same as the pipeline in 2.2.2. Instead of Jenkins, which is server-based, we build the pipeline with AWS CodePipeline. Figure 4.10 shows the activity within the CodePipeline in a single graph. Different from Jenkins who can do the whole continuous delivery process solely with the help of plugins, the CodePipeline just provides a platform which the development team can configure a workflow with AWS DevOps tools or other third-party tools.

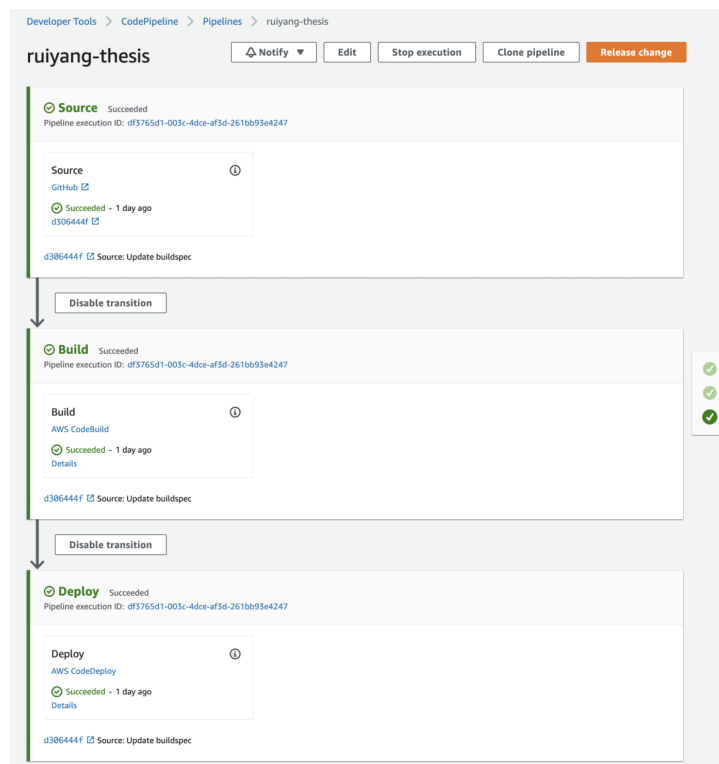


Figure 4.11: Our Workflow in CodePipeline

We use CodeCommit to replace GitHub as the version control system, CodeCommit is a fully-managed source control services [98] from AWS, which hosts Git repositories. Same with GitHub, it helps used managed the server which hosting Git repositories and eliminates the need to managed and scale the infrastructures. It also supports the pull request and code review, which is necessary for work under Github flow.

In the next step, we are using AWS CodeBuild, which we introduced in Chapter 3. AWS CodeBuild does the same procedure as in Jenkins pipeline. It does code analysis, unit test and builds the Java application with Gradle, build the Docker image of the application and push to the ECR. The CodeDeploy deploys our application to ECS with Blue-and-green deployment strategy.

The implementation of continuous delivery in CodePipeline is straightforward compared with Jenkins. In Jenkins, without the help of the plugin, the workflow of a continuous delivery pipeline can only be defined by groovy code, while CodePipeline natively provides a graphical user interface for continuous delivery pipeline modelling.

4.3.2 Source Control with AWS CodeCommit

We use CodeCommit to replace the GitHub that we used in the non-integrated toolchain based on the following reasons.

First, within CodePipeline, it is faster to clone a project from AWS CodeCommit than from GitHub. Our test shows that it takes CodePipeline on average 6 seconds to clone the case project from GitHub. However, the average clone time from CodeCommit is only 3 seconds.

Second, CodeCommit has better integration with AWS. CodeCommit supports manage each user's access to the Git repositories with AWS Identity and Access Management (IAM). IAM is a centralized tool to manage each user's permission to all different AWS services. It is clear that IAM cannot manage the user's access to GitHub but can manage user's access to CodeCommit repositories. We think it is always good that the team could manage the access to all component within the DevOps toolchain. Thus use CodeCommit instead, Github can ensure centralized access management within the DevOps toolchain.

Third, it is easier to use CodeCommit within the CodePipeline. CodeCommit could be added into AWS CodePipeline by select the repositories that the project located, while GitHub requires an additional manual login process.

4.3.3 Build and Test with AWS CodeBuild

Same with the design of our Jenkins pipeline, AWS CodeBuild also executes the build within Docker container. The image of the Docker container provided by AWS already contains environment for the build of different programming languages. It also includes the Java environment and Gradle which needed by our case project. Therefore we could save time in setting up the pipeline since we do not need to define the Docker image for the build by ourself.

As we mentioned in 4.3.1, the process within CodeBuild is the same as we have in Jenkins before the stage "Deploy". We will not describe the process again here. Same with Jenkins, the workflow of CodeBuild is defined in a YAML configuration file. The only difference in the build workflow is that CodeBuild stores the build artifacts to S3. The build artifacts stored in S3 are configuration files, which define the deployment configuration in CodeDeploy. This is because CodeDeploy requires the deployment configuration from the step before it to run automatically.

4.3.4 Blue/Green Deployment with AWS CodeDeploy

One of the advantages of AWS DevOps tools is good integration with other AWS services. During the design and implementation of our toolchain, this advantage shows in the deployment to ECS with CodeDeploy.

In Jenkins, there is a lack of specific plugin that helps us deploy the project into ECS or EKS. Thus we have to deploy our project to ECS with AWS command-line interface(CLI). The problem with AWS CLI is that it only supports the most basic deployment strategy, which is rolling update deployment. The rolling deployment strategy is to replace the old code running on the instances with new code gradually, instance by instance. Such a difference shows better integration between CodeDeploy and AWS infrastructure, which allows us using more advanced deployment strategies.

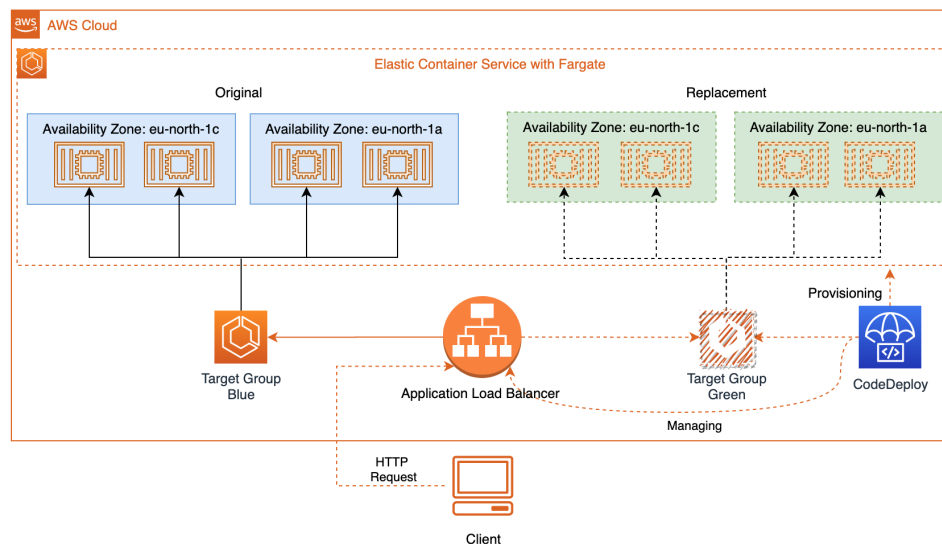


Figure 4.12: Blue and Green Deployment for Our Deployment Solution

In real-life production, the team would like to make sure the deployment is reliable with minimized downtime. Thus the safety is highly valued in deployment strategy. In answering this need, a strategy called blue/green deployment, which is now widely used in the industry. AWS CodeDeploy natively supports the blue/green strategy. A blue/green deployment is a deployment strategy that requires two sets of totally identical deployment environment that runs the new and original version of code respectively, while the load balancer is gradually routing more incoming requests to the environment that runs the newer version of code. Blue/green deployment could minimize the downtime to nearly zero[99].

Figure 4.12 shows the visualisation of blue/green deployment. It also shows our design on how to implement blue/green deployment with CodeDeploy (shown before in Figure 4.8). CodeDeploy controls routing policy within the load balancer. When new deployment comes, CodeDeploy does the following steps:

- Provisioning new identical deployment environment (replacement environment) and deploy a newer version of code on it. In ECS, the deployment is called "task set".
- Control the load balancer, rerouting incoming traffic gradually to replacement ECS task set. The rerouting rule is configurable according to the need. For example, the default rule is 10% per minutes. This means there will be 10% more requests from the client will be routed to the replacement ECS task set every minute until all the traffic has been routed to the replacement task set. We do not rerouting all traffic at once to ensure the service will not fully down if the new deployed task set not works properly. This minimizes the downtime of our deployment.
- Wait for a certain time (depends on the deployed project) after rerouting is done. During the rerouting, the load balancer keeps doing the health check to the new deployment by sending a request to the health check API endpoint. CodeDeploy read the health status from the load balancer. If the replacement tasks set is un-healthy, CodeDeploy does rollback by rerouting incoming traffic back to origin tasks set.
- If the new deployment is still healthy after waiting time, CodeDeploy terminates the running origin tasks set, and this means the old deployment environment is removed. Now the whole deploy process is done.

When the error happens with a newer version of code, with blue/green deployment, we can immediately roll back to the last version by switching the rerouting to Blue [100]. While in the rolling update we are using in Jenkins pipeline, we have to redeploy the previous version. This difference reflects the better failed-safe of blue/green deployment. Under the same circumstance, the rollback with the rolling update is nevertheless taking a too long time, since we have to replace the already deployed code to the previous version. This could cause longer downtime of the server. However, because an identical environment must be run, blue/green deployments also bring more costs.

The better integration of CodeDeploy with the rest of AWS also brings benefit in the monitoring of the deployed solution. Aside from the existing monitoring with CloudWatch, CodeDeploy also provides us with a dashboard to show the deploy progress and the traffic rerouting process. Figure 4.13 shows the dashboard of CodeDeploy that shows the status of our case project during the deployment.

In comparison, with our non-integrated toolchain, we can not easily do the blue/green deployment. A possible solution is to set up an AWS Lambda function which will be triggered after the deployment. The invoked Lambda function control the load balancer to gradually redirect the traffic from the previous deployment to the new deployment. At the same time, the second lambda function continuously read the health status of the new deployment from the load balancer, and trigger the rollback if the new deployment is unhealthy. Besides all these, a

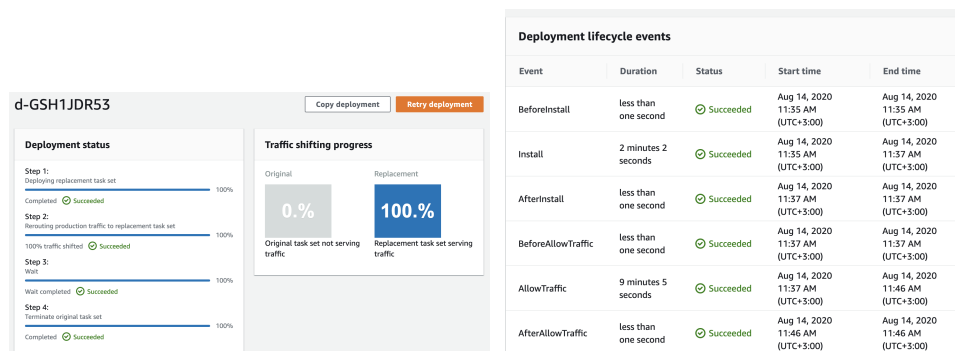


Figure 4.13: Part of the CodeDeploy Dashboard, Deployment Status (Left) and Deployment Timeline (Right)

dashboard is needed for the developer to monitoring the rerouteing status. On the contrary, in CodeDeploy, all the above tasks are being taken cared for. There is fully automated routing control on the load balancer, a dashboard that can monitor the realtime traffic percentage to blue and green instance, as well as the warning once the new deployment is failed.

4.3.5 Integration of AWS DevOps Tools using CodeStar

In 4.3.1, we introduce how do we integrate different stages in the continuous delivery pipeline (CD pipeline) with CodePipeline. However, an integrated continuous delivery pipeline is not called integrated toolchain. The DevOps toolchain is centred with the continuous delivery pipelines, but the pipeline is not all of the toolchains. To integrate other AWS tools with the CD pipeline and get the integrated toolchain as a single application, we use CodeStar.

Figure 4.14 show the user interface of the CodeStar dashboard. We can see, besides the CodePipeline, the CodeStar also integrate tools like monitoring, project management, and version control and then compose our integrated toolchain. In conclusion, the AWS integrated toolchain includes the following tools:

- **CodeStar:** Integrate below tools into a single toolchain project management functionality. The project management could be extended by integrating with third-party tools, i.e. Jira.
- **CodePipeline:** Modelling Continuous Delivery Workflow, integrate tools that used within the CD pipeline.
 - **CodeCommit:** Version control. Git repository which store the source code. Supports Team collaboration through branching, pull request and code review.
 - **CodeBuild:** Build code of the software project.
 - **CodeDeploy:** Deployment and monitoring.

- **CloudWatch:** Monitoring deployed solution.

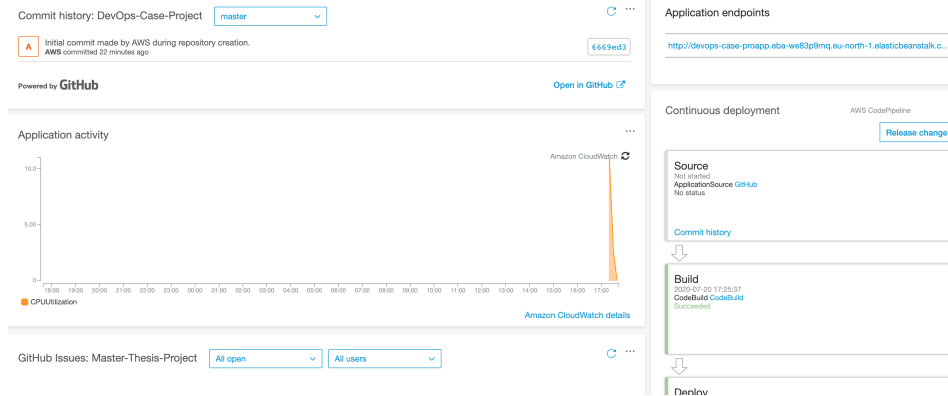


Figure 4.14: CodeStar Dashboard

4.4 Comparison between Integrated and Non-integrated Toolchain

In this section, we discuss the difference between these two kinds of toolchains. The scope of comparison will be limited within the scope of functionality and ease of implementation. Furthermore, it is only based on our experiences with the tools used in our implementation. We summarize the difference between these two toolchains as in Table 4.2. We will do more comparison related to the performance and cost in Chapter 5.

4.4.1 Implementation and Cloud Deployment

The cloud-based integrated toolchain is delivered as a hosted solution in a serverless model. However, we noticed that the non-integrated toolchain could also be completely serverless if we are using hosted tools for all the components.

For example, in our solution, we only have a continuous integration pipeline which is on-promised and need to be deployed to the VM manually. If we replace Jenkins with some other hosted tools, for example, Travis CI²⁷ we can actually build a fully hosted but non-integrated DevOps toolchain. But, for the following reasons, it is not a satisfactory solution; thus, a non-integrated toolchain usually has some on-promised modules that need operational effort and cloud knowledge.

- The hosted tools, especially tools for continuous integration pipeline, are all closed-source commercial solution. This means there is not an open-source community like in Jenkins. The extensibility therefore limited, for

²⁷<https://travis-ci.org/>

	Non-Integrated Toolchain	Integrated Toolchain
Open-source	Open-source solution existed	No, usually hosted commercial solution
Delivery method	Each part is a stand-alone tool either hosted or on-promised, depends on the tools selection	As a single cloud hosted software
Implementing time	Long	Short
Operational effort	High	Low
Visibility on status	Depends on tools, a well-integrated toolchain could gives good overview on the whole toolchain.	Easy to see the status as a whole without additional implementation effort, low visibility on under-laying server since it's hosted solution
extensibility and tool selection freedom	Free to select tools for each part of the toolchain.	Limited integration with third-party tools

Table 4.2: Comparison of DevOps Toolchains

example, AWS DevOps tools can usually integrate with the certain tools that partner with AWS. Besides, the commercial user always needs to pay for these hosted tools.

- Hosted tools run in the vendor's server, and it requests user log in to use this tool which brings extra integration difficulty. This means for two hosted toolchain to integrate, it not only need to do the integration in the data transfer but also needs to connect their account system, for example, with OAuth. This extra inconvenience makes most hosted DevOps tools only do the integration support with other most popular tools which largely limited the extensibility. AWS natively only support the tools belongs to the DevOps Partner solutions²⁸ integrate to the AWS DevOps tools. Is not easy for the user to develop a tool and use it directly without making an agreement with AWS.

Therefore, a non-integrated toolchain usually has some on-promised module in

²⁸<https://aws.amazon.com/devops/partner-solutions/>

real-life use. In our deployment process, we find it requires a lot of work to put an on-promised tool to cloud, especially if the developer is not familiar with the cloud platform that deploys these tools. For example, when setting up the Jenkins cluster that is only one component of the DevOps toolchain, we need to do the following steps:

1. Create a cloud virtual machine (EC2 instances) for hosting the Jenkins master.
2. Setup IAM role for Jenkins master VM, make sure it has access to other AWS recourse that needed during the build.
3. Setup security group and networking for the VM makes sure it can be accessed from the internet but only accessible within the company's IP range, and only port needed is opened to the public.
4. Install Jenkins in the VM. Research what plugin is needed and install necessary plugins.
5. An tedious setup process for setup Jenkins cluster that supports the distributed build. This includes setup ECS cluster for a build agent. Although Terraform makes the provision of cloud resources easier, still, prior knowledge for AWS is needed. The experiences in AWS also need to correctly configure the ECS cluster that maximizing the performance of build agents.
6. Develop a Docker image for the container that runs Jenkins agents.
7. Setup integration with other tools in toolchains by finding correct plugins and configure these plugins.

Only after these steps, we can start using Jenkins within the DevOps toolchain. In comparison, the core feature of the integrated toolchain in AWS is an out-of-the-box feature which means there is no previous cloud knowledge needed, and there is no deployment and environment configuration required before we use it. We are free from all the steps we mentioned above.

In conclusion, the integrated toolchain could save time

4.4.2 Extensibility and flexibility

The integrated toolchain is a hosted platform that runs by a vendor. Similar to we mentioned in 4.4.1, we find all the currently integrated toolchain are all commercial and closed-source and not friendly to third-party plugins. So the integration of their third-party tools is usually only limited to popular tools. For example, AWS DevOps tools only support 21 tools within it is "DevOps Partner Solutions"²⁹.

²⁹<https://aws.amazon.com/devops/partner-solutions/>

Nevertheless, different from the single hosted tools we mentioned in 4.1.1, a hosted integrated DevOps toolchain mostly has everything needed for DevOps lifecycle, so it does not need to be able to integrate with third-party tools. Still, the limitation in third party tool support might make the software team facing trouble when they want to use certain tools which are not very supported.

A non-integrated toolchain allows the software team to pick any tools for each component, as long as those tools can be integrated. The tools in the toolchains could also be open-source, which allow the software team to modify the tools according to their need. For example, develop a plugin for Jenkins that allows the integration of internal company tools with Jenkins.

In conclusion, in terms of extensibility and flexibility, non-integrated toolchains are better than integrated toolchain.

4.4.3 Integration Between Tools

As we mentioned in 4.4.1, sometimes it is hard for tools within a non-integrated toolchain to integrate, especially between the hosted tools. During our implementation, we also realized that, first, it requires some configuration work for tools to be able to work together. Secondly, sometimes the integrating could be buggy is the configuration was done properly. For example, in our toolchain, the ECS build agent cannot connect to the Jenkins master, because the networking within the ECS cluster was not correctly set. This means the software teams need further maintaining of the toolchain. In integrated toolchain, the toolchains are delivered as a single cloud-based software, and each part naturally coped with each other, which makes integration much more straightforward. The better integrating between each component also makes it easier to monitor the toolchain as a whole.

4.4.4 Visibility

In 4.4.3, we mentioned that the integrated toolchain is easier to be monitored as a whole, however, when comes to every single component, in our implementation, we find out that integrated toolchain is lack of visibility. We met difficulties with integrate/non-integrated toolchain respectively, and the experience in solving these two problems shows how visibility could affect the ease when it comes to finding where the problem is.

The first problem was with the non-integrated toolchain, which we have full visibility to the underlying virtual machine. Was that Jenkins master was having difficulty in the provision and connect to the agents. Since Jenkins is a web service deployed in our EC2 virtual machine, we solve the problem easily by reading the Error message within the Jenkins log file.

The second problem we met was within AWS CodeDeploy, which we have no visibility to the machine it runs on. The CodeDeploy failed to deploy the case project to the ECS cluster. We could not find the reason at that time since we cannot find the log of CodeDeploy anywhere since it is not shown in the web interface, and

we have no access to the underlying cloud infrastructure to find the log file either. In the end, we have to reread the documentation and find our load balancer was not properly set so that CodeDeploy could do the health check. Thus the CodeDeploy cannot provision the new deployment environment for us.

The lack of visibility is a problem with all hosted serverless services since the users do not have the visibility to the infrastructure behind the service.

4.5 Challenges in Implementation and Design of DevOps toolchains

In this section, we discuss the challenge that we met during the implementation.

4.5.1 Challenge I: The Enormous and Unregulated Jenkins Plugins System

Jenkins has more than 1600 plugins which brings the amazing software extensibility, which is one of the main advantages of Jenkins. However, there are two problems with Jenkins' plugins; First, there are usually more than one plugins that have the same functionality, for example, there are at least five different plugins related to running Docker container as Jenkins agent. Second, most of the plugin is developed by the open-source community, compared with a commercial product, the open-source project is lack of support such as case support to the specific problem. Besides, the open-source project could end up without further support once the developer decides to discontinue the project.

During our implementation, we find out there are two plugins that support run Jenkins agent in ECS cluster. However, we find only one work after we tried both plugins. Besides, the documentation of Jenkins plugins sometimes is abysmal. For example, the documentation of the plugin that we use for Jenkins agent is too brief to tell us how to use the plugin, and it is not even mentioned the security setting needed in Jenkins master node that allows agents to connect the master node.

As a result of the above three factors, for a developer who does not has previous experience in build DevOps toolchain, we end up spending a very long time in selecting and configuring tools and trying to solve the problem which nether mentioned in the documentation and on the internet.

4.5.2 Challenge II: Fargate Does not Supports Container runs in Privileged Mode

As mentioned in the title, this is a limitation in AWS Fargate to prevent containers from gaining access to critical resources on the host. As a result, we cannot use Docker within a Docker container that runs on Fargate. This makes it impossible for us to distribute the "Build docker image" and "Push to ECR" stages to agents. Instead, we have to run them on the master. Luckily, these two stages take a short

time (15s in total), so this limitation will not slow down the pipeline too much when multiple builds run in parallel.

A possible solution solving this problem is to run these two stages in AWS CodeBuild, AWS CodeBuild has support to Jenkins, which allow us to run certain Jenkins stages in CodeBuild. Moreover, CodeBuild supports fully parallel execution as in Fargate.

4.5.3 Challenge III: Slow Starting Time for Agents in AWS Fargate

On average it takes around 60s from sending a Jenkins job to agent, to running the job in an agent. To our case project that takes 90 seconds to go through the whole pipeline, this is a relatively long time. During this 60s, Jenkins master node sends task definition³⁰ to ECS, provision a Fargate instance within ECS, then pull the image we developed for Jenkins agent, start the agent container within Fargate and then connect to the Jenkins master. This challenge is due to the nature of the serverless computing that we discussed in Chapter 2, and we believe there is not an economical way to solve the challenge with the current setup.

4.5.4 Challenge IV: No Enough Visibility in AWS DevOps tools

As we mentioned in 4.4.4, the lack of visibility of the underlying processes (especially in CodeDeploy) has brought some obstacles to our debugging pipeline. When we tried to deploy the case project to ECS via blue/green deployment, there was a problem: CodeDeploy got stuck in creating a replacement service. We know there was something wrong within either the configuration of ECS, load balancer, health-check or security/network setting. However, there is no log output of the underlying deployment process in CodeDeploy. Finally, we had to check all the possible causes of the problem one by one and found that it was a failed health check due to the configuration error in the load balancer. Such a way of detecting issues with the deployment was very time-consuming.

³⁰Define specification of a container runs in ECS.

Chapter 5

Performance Comparison and Evaluation

In this chapter, we describe our experiments regarding two research questions we proposed in Chapter 1. The experiments based on the DevOps toolchains we implemented in Chapter 4.

In Section 5.1, we will examine how does the serverless compute engine for containers (Amazon ECS on AWS Fargate) could influence the performance of non-integrated toolchains. Thus, in the experiment, we implement the solution with a different type of cloud environment (with/without serverless) as a comparison group. In Section 5.2, we focus on answering research question 2, in which we will compare the performance of continuous delivery pipeline composed of fully-managed serverless DevOps tools in AWS with our Jenkins-based pipeline that runs on the virtual machine.

5.1 Experiment 1: Experiment on Serverless Container Services

The Docker agent has already been supported by many CI/CD tools and we introduced in Chapter 4. The serverless container services in AWS (AWS Fargate) provides the possibility to ease the infrastructure management task for the Docker build agents. This experiment is a controlled experiment which examines whether serverless container service could improve the continuous delivery pipeline from various perspectives.

5.1.1 Test Task and System Description

In this experiment, we run the continuous delivery process of a Spring Boot web application with our DevOps toolchain. From the experiments, we could verify our assumption in Chapter 3 and better-answering RQ 1.

As we described in Chapter 4, the continuous delivery pipeline includes the following steps:

1. *Checkout*: Pull the most recent change from GitHub repository
2. *Build*: Build the application with Gradle, with automating testing with JUnit integrated into Gradle.
3. *Build the docker image*: Build the docker image of our Spring Boot application.
4. *Push to Container Registry*: Push the docker image from the last step to the AWS elastic cloud registry (ECR) for further deployment.

In these four steps, the step "Build", and "Checkout" is being done in parallel within the ECS cluster. As we mentioned in CH4, when the new job started in the Jenkins master server, Jenkins will provision a new container instance within the ECS cluster. The container is managed directly by AWS, so we don't need to create and manage the virtual machine that runs the container. We use this setup in our initial implementation as the control group.

In the experimental group, we replace AWS Fargate with traditional VM, which is EC2 in the Amazon Web Services. The parallelisation pattern remains the same; this means as in the control group, only the first two steps are being run distributively in the Jenkins nodes. The EC2 instances belong to an auto-scaling group that will scale up when CPU Utilisation rate reach 70%. The initial size for an auto-scaling group is 1.

Figure 5.1 shows the architecture of 2 groups in this experiment. The experimental group on the left is a Jenkins server with the traditional virtual machine as workers node that hosting the container agent. The architecture of the control group on the right has agent nodes dynamically provisioned as serverless containers hosed by AWS Fargate.

Hardware

The hardware of the instance that runs Jenkins agents is the independent variable that exposed to the change in the experiment.

The experiments are conducted on Amazon Web Services (AWS). The hardware of Jenkins master node in both experiment groups is the same, which is EC2 instance of type t3.medium with two virtual CPU, 4 GB RAM and 30 GB disk. The EC2 instances as worker node are type t3.small, with two virtual CPU and 2GB RAM. Each EC2 instance can run one container at the same time.

In the control group, which is the implementation we presented in CH4, the Jenkins agents run on AWS ECS powered by AWS Fargate. The virtual hardware resources that are allocated to each serverless container is two virtual CPU, and 2 GB of RAM. The identical hardware setup between to groups makes sure that each container shares the same hardware resources as in another group, so the hardware will not affect the result.

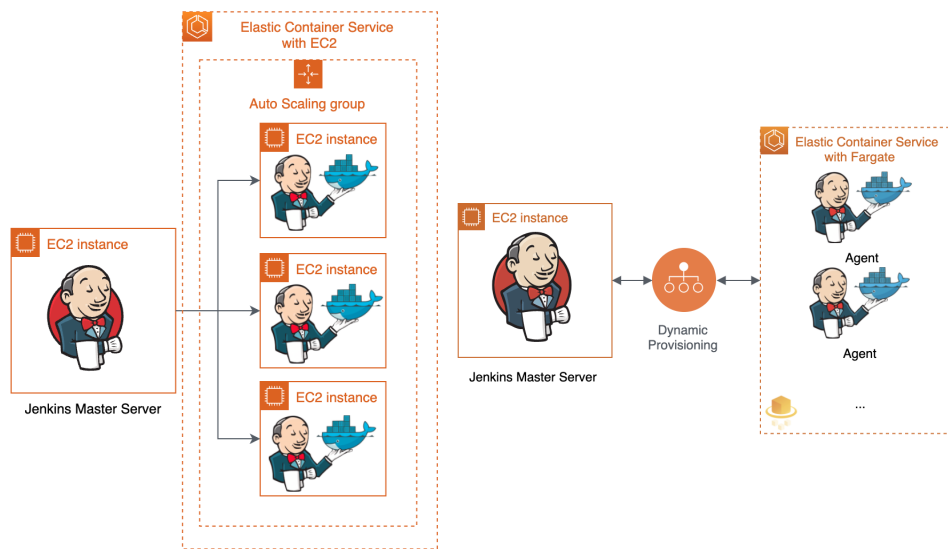


Figure 5.1: Architecture diagram of the test Jenkins cluster with agents running in traditional virtual machines (left) and on ECS with AWS Fargate (right)

Software

We maintain the same software setup in each group. The version of the software is all the latest version as for February 2020. The operating System for EC2 instance that runs Jenkins master node is Ubuntu Server 18.04. The version of Jenkins that runs on the server is 2.222.3. For connect ECS and Fargate which works as the Jenkins agents, we use Jenkins plugin "Amazon Elastic Container Service (ECS) / Fargate", version 1.34. The container in Fargate/EC2 for running the Checkout and Build steps is from our developed docker image, which can be seen at ¹. The docker image includes essential dependencies that will be used to build the Spring Boot application. It's the base image include a program which allows container connects Jenkins master as an agent. The "Build" step in our pipeline uses Gradle (version: 6.2.1) as the build tool for the application, with OpenJDK 1.8.0.252 as Java virtual machine (JVM). This step also includes the automated testing and code analysis, as plugins of Gradle.

To shows how does the two setups performance within the teams with different sizes, we run by run the different number of tasks parallel through the pipeline. This scenario simulates the different team size and shows the scalability when it comes to the need for task parallelisation in larger organisations. It also imitates the DevOps process of a microservices software project, which could have multiple jobs for different service runs in parallel.

¹<https://hub.docker.com/t/dry1995/jnlp>

5.1.2 Performance Properties and Evaluation

We run the pipeline through 2 different setups, and we will get the result of the following properties:

- *Runtime* describes the total time for finishing all the jobs. If the jobs run in parallel, the runtime is from the start of jobs until the end of the last finished job.
- *Cost Structure* describes the daily cost of 2 setups under the same workload, within the same period.
- *Resource Utilisation* describes the average CPU/RAM usage for each instance during a single run of the pipeline. The purpose of this comparison is to show the performance of the same application in a different environment (EC2 and Fargate).

5.1.3 Result and Evaluation

Here shows the result of this experiment. We also evaluate our experiment result by analysing the factors that lead to the results.

Runtime

We first compare the runtime of these two setups. Except for test the runtime of single job runs with two setups respectively, we also test the runtime of each pipeline setup under different number jobs executed in parallel. Figure 5.2 shows the test result.




Figure 5.2: Runtime of Pipeline with Different Jenkins Agents Under Different Number of Jobs Runs in Parallel

The test result shows that when it comes to the execution of a single task. The traditional VM has a faster delivery speed over serverless solution (AWS Fargate). However, with the number of jobs that run in parallel increases, the total runtime on the traditional VM decrease. On the contract, on the serverless solution(Fargate), the runtime remains almost the same.

We analyse the reason behind this result, and we found out that the longer runtime with the single job on Fargate is because the longer starting time of Jenkins agent. In EC2, the Jenkins will simply provision a Docker container within EC2 VM, and connect to the Jenkins master node. However, in Fargate, the Jenkins can only connect to the agent once AWS finishes the initialisation of underlay infrastructure that runs the serverless container, which takes a significantly longer time. The slow provisioning shows one of the limitations of serverless computing (cold-start) that we mentioned in 2.4.3.




Figure 5.3: Execution Mode of Pipelines with Different Jenkins Agents

When it comes to parallel job execution, the performance of Fargate is significantly better. The performance difference is because, in Fargate, each container runs on an independent instance on AWS's infrastructure. Therefore, AWS provisions one Fargate instance for each running job. The independence between Fargate instance ensures the agent will not compete for the resource. On the other hand, when we run multiple agents in our EC2 instance, due to the limitation of resource, part of running jobs has to wait until the resource on EC2 instance available until they can start the execution. To further investigate the reason for the result, we

observe the parallel execution mode when it runs three jobs in parallel. Figure 5.3 shows the execution modes. We find that the easily scalable character (mentioned in 2.4.3) helps the serverless suites better with the parallel task. The long wait time is the reason that makes the total runtime in EC2 much longer.

We also notice that in Figure 5.2, when the parallel task reaches 10, the EC2 runtime becomes shorter. The shorter runtime of EC2 is because we set auto-scaling for our EC2 instance. So in the later part of our experiment, the EC2 scaled from 1 to 2 and then to 3. Nevertheless, even with 3 EC2 instances, only three jobs are allowed to run in parallel, while in Fargate is easy to have ten jobs runs in the complete parallel method. This because the scaling of EC2 VMs is much slower because it is heavier to create a VM than create a new Fargate instance. The other reason is the auto-scaling of EC2 VM is triggered by reaching certain resource utilisation threshold, but in Fargate is based on task number(Jenkins agent number). Even we set the scaling policy of EC2 to a more aggressive pattern, the AWS still more "hesitate" to create new instances compared within the Fargate. Therefore, in the real-life software development, when the number of jobs surges in Jenkins cluster, the Fargate could react faster in terms of scaling. And when the job drops, Fargate also scale-in faster, which saves cost.

Resource Utilization

We compared the average resource utilisation within containers that runs Jenkins agent in two setups. The data from the AWS CloudWatch (Figure 5.4) shows that the resource utilisation rate is similar in these two setups. The similarity is because the "run in the same way regardless of the host environment" [84] feature of Docker container that we mentioned in 4.2.2.



Figure 5.4: The comparison of Resource Utilization

Cost Analysis

The container orchestration service ECS itself is free of charge. We only pay for the resources we are using, which is Fargate or EC2 virtual machine.

In AWS Fargate we pay only by resource we are using and the runtime. The price for Fargate service in EU(Stockholm) is \$0.04165 per GB RAM per hour plus \$0.0049 per vCPU per hour². Thus, in our experiment setup (2vCPU, 2 GB RAM), the price should be \$0.0931 per running agent for one hour's runtime. In

²<https://aws.amazon.com/fargate/pricing/>

AWS EC2 our cost depends on the type of VM we are using. The type of VM we use for running Jenkins agents is t3.small that costs \$0.0216 per hour ³. The Fargate-based Jenkins agent is more expensive than EC2-based agent.

However, the pay-per-use characteristic of serverless service makes Fargate more competitive in terms of cost. The instance only up and run when Jenkins master distributes the job. Once the job finished, the AWS will terminate Fargate instance immediately when it finds no container is running on it. However, EC2 not so flexible due to the resource-utilisation-based scaling policy. The EC2 instance is not scaled in immediately after job finished, and the user has to pay for the instance runtime before it gets terminated – even there is no job running in EC2 instance any more. So depends on the frequency and length of the build task, the user could have runtime per hour when using Fargate instance.

5.1.4 Conclusion

In this experiment, we compare the serverless (AWS Fargate) and non-serverless solution (EC2) for hosting distributed continuous delivery pipeline. The experiment shows that the performance of the serverless solution is worse when it comes to single job execution, which is because of the cold-start issue with serverless computing. However, in terms of the parallel job executing, the serverless solution has better performance over traditional VM.

An better parallel execution performance is meaningful for apply DevOps practice in the microservices development. A microservices could have hundred of services. In the philosophy of microservices, the release of each service should be independently [101]. To ensure this, one practices it to use one pipeline for each service. However managing a hundred pipelines is not an easy task, and better practice is to have multiple services share a single pipeline [102]. Therefore, when multiple services share the same pipeline, this improvement allows services to be released independently without waiting for each other.

In term of cost, the serverless solution is more expensive per hour. However, the difference could be offset by the less runtime. Moreover, resource utilisation is similar in both solutions. Table 5.1 show a comparison between two experimental groups.

5.2 Experiment 2: Experiment on Integrated DevOps Toolchain

For solving the RQ2, we compare the implementation of our design of two different the DevOps toolchain – the non-integrated Jenkins based toolchain and the AWS DevOps toolchain implementation with AWS DevOps tooling.

³<https://aws.amazon.com/ec2/pricing/on-demand/>

	EC2 Agents	Fargate Agents
Build Time	98s for a single task. Largely increases when run multiple tasks in parallel	155s for a single task. Remains unchanged when multiple task runs in parallel
Cost	Lower per hour price (\$0.0216)	Higher per hour price (\$0.0931), however the billable hour is shorter, the total price could be lower.

Table 5.1: Conclusion of Experiment 1

5.2.1 Test Task and System Description

This second experiment is similar to the first experiment, and we deliver our case project through two DevOps toolchains. The first toolchain is our Jenkins-based toolchain in Section 4.2, with agents runs on AWS Fargate. The second toolchain is the toolchain with AWS DevOps tooling that we described in Section 4.3.

During the experiment, we have set up as follows:

Hardware AWS DevOps tools allow us to select the hardware configuration for underlying computing resources. The configuration we chose is two virtual CPU and 3 GB RAM. The Hardware configuration for Jenkins build agent is two virtual CPU and 4 GB RAM. We cannot set the RAM to 3 GB since AWS Fargate is not allowing ⁴ RAM to below 4GM when using two virtual CPU. However, we still allowed to set an additional software RAM limit to the build agent runs in Fargate agent. Thus, we limit the RAM that a running build agent to 3 GB, which ensure both setups have identical hardware configuration.

Software The version of Jenkins that runs on the server is 2.222.3. We use Jenkins plugin” Amazon Elastic Container Service (ECS) / Fargate”, version 1.34 for connecting ECS and Fargate which works as the Jenkins agents. The Jenkins cluster has the same configuration as the last experiment. The built environment in both setups is the same, with Gradle (version:6.2.1) and JVM version is OpenJDK 1.8.0.252.

5.2.2 Performance Properties and Evaluation Criteria

We run the pipeline on these two different setups, and we will get the result of the following properties:

⁴<https://docs.aws.amazon.com/AmazonECS/latest/developerguide/task-cpu-memory-error.html>

- *Runtime* describes the total time for finishing all the jobs. If the jobs run in parallel, the runtime is from the start of jobs until the end of the last finished job. AWS monitors its parallel execution ability in the document⁵⁶ of both CodePipeline and CodeBuild, thus we will also validate this ability by analysing the parallel executing pattern of AWS integrated toolchain.
- *Cost Structure* describes the daily cost of 2 setups under the same workload, within the same period.

The resource utilisation comparison is not available in the experiment since we cannot get the resources utilisation in underlying hardware resource when running AWS DevOps tools.

5.2.3 Quantitative Experiment Result and Evaluation

This section shows the result of our experiment. We also give an evaluation and analysis of the reason behind the result.

Runtime

As in Experiment 1, we compared the running time of each toolchain under different loads⁷. We only compare the runtime without the final deploy stage. This is because, first, although we can still use the AWS toolchain to deploy to ECS in the same way as the non-integrated toolchain: by using the AWS CLI in CodeBuild to deploy. However, by doing so, we skipped CodeDeploy and instead deployed our project using CodeBuild. This is not a natural practice in production because there is another tool dedicated to deployment (CodeDeploy). Secondly, the deployment mode used in CodeDeploy is very different in terms of the speed of the deployment method used in Jenkins. Figure 5.5 shows the result of this experiment. The runtime of integrated toolchain increases with the number of the task runs in parallel, while the non-integrated toolchain remains stable regardless of the number of parallel tasks. We will analyse the reason later in this section.

From Figure 5.6, we can see that during the execution of a single job, two toolchains have similar runtime. The similar runtime is because the build stage, which takes most of the runtime, is running in a similar environment in both toolchains. Both execution environment is within a Docker container runs with the same hardware in AWS. Although, the Docker image for the container is different in two toolchains, and we are not sure what if the actual hardware is the same since we have no visibility to the hardware in both toolchains. However, still, the performance is not so different in both toolchains.

We further observe the time allocation between stages within the runtime in both toolchains. The figure shows our observation. To get the runtime in Figure 5.6, we

⁵<https://aws.amazon.com/codepipeline/features/>

⁶<https://aws.amazon.com/codebuild/features/>

⁷Workload refers to the different number of jobs running in parallel in the two pipelines



Figure 5.5: Runtime of the pipeline with on Different Toolchain

run the pipeline in each toolchain for five times and get the average runtime of each stage. The first difference shows in Figure 5.6 is that the Git Checkout stage in AWS integrated toolchain does not run on the container build agent. Instead, it is running in the unknown environment fully managed by AWS.




Figure 5.6: Observation of Runtime of Each Stage on Two Toolchains

We also compared used GitHub and use CodeCommit for the source control in the AWS integrated toolchain. We observed that it takes 10 seconds to do the git checkout when using third-party tools (GitHub) as the source control system. During these 10 seconds, the actual Git checkout only takes around 5 seconds, while the AWS toolchain does not do anything in the first 5 seconds. We presume that AWS provisions environment for runs the Git checkout stage in this 5 second, so this stage is also running in a serverless environment within AWS. However, we are not allowed to configure anything in this environment. In contract, it only takes 3 seconds to check out new code from CodeCommit and no waiting time before the checkout. This speed difference could because of the CodeCommit is the part

of the AWS services. Thus it has faster data transfer between CodePipeline, or because of no environment provision needed before checkout.

The second difference is that in AWS integrated toolchain, the provisioning of the build environment is much faster. We presume this is due to AWS manages both build agent and this toolchain, and AWS is optimising the provisioning process of the build environment. Furthermore, the cloud instance that runs build environment in CodeBuild might already be started (used for the build task of other users) before CodePipeline sends our build job to CodeBuild. On the other hand, instances (Fargate) that run Jenkins agents is a cold start, which means it is not running before we send build job there. Nevertheless, we can also notice that the build time in the AWS integrated toolchain takes a longer time. It is hard for us to find the reason since the hardware configuration used by AWS is unknown except the size of RAM and the number of virtual CPU.

Also, we observe that with the workload goes up, the runtime of AWS integrated toolchain increases with it. We already answered why the runtime of the pipeline in our non-integrated toolchain with agent runs in AWS Fargate does not change over time in Section 5.1.3. To explain this, we also check the runtime of each stage when runs several jobs in parallel. We notice that when it has multiple jobs runs in parallel, essential if the parallel jobs' number goes over 5. AWS will start limiting the resource allocation, which limits the number of jobs that runs at the same time. As a result, part of our running jobs has to enter the "Queued" status before entering the "Provisioning" stage. In 10 jobs we run in parallel for the experiment, the time spent for queuing for resource various from 1 second (get resource allocated directly) to 130 seconds. Among these ten jobs, four jobs were put into the queue before resource are available to them. Figure 5.7 shows the distribution of queued time among jobs. This validates the claim from AWS about parallel execution ability, which we mentioned in Section 5.2.2. However, the parallel pipeline execution is not fully in parallel but with some limitations on available resources.

The runtime of the non-integrated toolchain is also slightly going up because of the runtime of the build and push container to ECR increases. The reason is obvious, we have these two stages runs on the master node, so the increased data transfer between agent and master and limited computing resource on the Jenkins master increase the runtime. However, as we observed in Experiment 1, the build stage is always fully paralleled, and the runtime of this stage remains unchanged despite the increasing workload. The queuing time for resource in each job is negligible.

Cost Analysis

As a serverless tool, AWS DevOps tooling charges us according to the time and type of hardware configuration (in CodeBuild) we are using. Each pipeline in CodePipeline cost \$1 per month, which is a negligible amount of \$0.0014 per hour, price of with the build agent type general1.small (Linux instance with 3GB




Figure 5.7: Observation of Queued Time in AWS Integrated Toolchain During The runtime Experiment, When 10 Jobs Run in Parallel

Memory and two vCPU) that we used in the experiment is \$0.3 per hour. The CodeDeploy is free of charge under the condition that we are not deploying to on-premises servers. The cost of S3 which store artifacts (1 MB in our case) is negligible according to AWS⁸. In conclusion, the price should be \$0.30014 per running job for one hour's runtime.

In our non-integrated toolchain, our experiment setup (2 vCPU, 4 GB RAM) is different from what we have in 5.1. The price should be \$0.1029 per running agent for one hour's runtime. While the cost of EC2 instance that hosts the Jenkins master node costs \$0.0432 per hour, this price will remain constant when multiple jobs run in parallel. In general, the cost of the non-integrated toolchain is \$0.1452 per hour when only one agent is running. From the calculation, we can see that the AWS integrated toolchain is more expensive under similar performance.

5.2.4 Conclusion

By analysing and explain the results, we could see that two toolchains have similar performance when it runs our case project. However, by further observe the runtime in each stage, we notice that the AWS integrated toolchain is faster in provisioning the build environment, but, slower in running the build and testing task itself. Such runtime distribution means it might be suitable for light but more frequent build task. The software project in real-life software development is much larger, and as a result, the build time will take a larger proportion in the total runtime. Therefore, the AWS integrated toolchain may become slower since it's building is slower than Jenkins non-integrated toolchain. From the parallel exe-

⁸<https://aws.amazon.com/getting-started/projects/set-up-ci-cd-pipeline/services-costs>

	Jenkins-centered Non-integrated Toolchain	AWS Integrated Toolchain
Build Time	Slighter faster(123s), better performance under parallel build because Fargate instances does not competing for resources with each other.	125s for a single task. Increased with the number of the parallel executing task goes up, could be because of the queueing for the resources within CodeBuild.
Cost	Lower per hour price, \$0.1452 per running agent (includes the cost of master node).	Higher per hour price, \$0.3 per hour per running CodeBuild agent. However no constant additional cost as the master node in Jenkins, might be cheaper in a long run.

Table 5.2: Result of Experiment 2

cution part, we conclude that with our AWS integrated toolchain, the AWS CodeBuild, which takes most of the runtime, do have a limitation on the resource that we can use at the same time. Thus, before the system provisioning any resource, some tasks have to wait for the resource allocation in the "Queued" status. The "Queued" status largely prolongs the total runtime.

In general, the AWS integrated toolchain is more expensive and slower than our Jenkins-centred non-integrated toolchain, note that this does not mean the non-integrated is better than integrated toolchain. As we discussed in the last chapter, the AWS integrated toolchain is easier and faster to implement, is more stable and with better customer support. Table 5.2 summarised the result of this experiment.

Chapter 6

Conclusions and Future Work

The changes in software development and businesses have affected teams to aim for more efficient practices and tooling, which are the key factors of DevOps. Thus in this master thesis project, we focus on the tools aspect within the DevOps, more specifically, is the DevOps toolchain that helps the software team to implement DevOps practises. On the other hand, the development of cloud computing technologies, especially serverless computing, has brought many changes in the development and deployment of DevOps toolchains. Serverless computing services in the cloud could either act as the runner of automated build and automated testing, monitoring tools or even replace the whole DevOps toolchain with a hosted integrated DevOps toolchain. With the motivation to combine the DevOps toolchain and serverless computing, this master thesis project explores the possibility that serverless computing could improve the DevOps toolchain in term of performance and functionality. This is done by, first, implemented a traditional non-integrated toolchain and deploy to AWS, and find what improvement could AWS serverless computing services bring to this toolchain for answering RQ1. Second, implement an integrated toolchain with AWS serverless DevOps tools, and compare with the traditional toolchains, which answer to our RQ2.

In this chapter, we first make conclusions by summarising our findings in Chapter 3 and Chapter 4. The conclusion is presented to answer our research questions in Section 6.1. In Section 6.2, we discuss possible future works.

6.1 Conclusions

In this section, we summarise our main findings during the implementation and evaluation as the answer to two research questions.

RQ1: *How can serverless computing services in Amazon Web Services helps the DevOps toolchain?*

Our main conclusion for RQ1 is that serverless computing services can benefit the DevOps toolchain by improving performance, eliminating server management

tasks, extend functionalities and reducing the cost of idle time. We identified that from following aspects, the AWS serverless computing services could improve the DevOps toolchain.

The first improvement is made by the managed serverless container service (AWS ECS with Fargate), which could work as the build agent which host the build and testing process within the CD pipeline. According to our experiments and analysis, the feature of serverless computing, combined with the beneficial brings by Docker container, could improve the parallel execution performance and save cost. The high performance in parallel pipeline execution could be helpful in several aspects: First, it improves the performance of continuous delivery for the microservices project as we stated in 5.1.4. Second, the serverless container service reduces the effort when setting up the build agents. Lastly, our experiments also show that Pay-per-use model makes Fargate cheap to use for hosting the build agent.

The second improvement is from the serverless function (AWS Lambda) could be used within the monitoring part of the toolchain. The event-driving computing model of serverless functional is perfect for processing with the logs, alarms and realtime performance metrics within monitoring. To a software team that sets up DevOps toolchain, they can use AWS Lambda to extend the monitoring system, such as customizing metrics and sending notification. In addition to these, AWS Lambda can help the software team to set up an AWS-event-driven workflow that triggers Jenkins job when changes happen in AWS services. [An example could enable Jenkins to build starts when new data is uploaded to S3 bucket.] better remove

The third improvement is from the serverless managed tools (AWS CloudWatch e.t.). Those tools could be used directly as a component in the toolchain. For example, A software team can use CloudWatch directly as the monitoring solution in the toolchain. This type of tools frees the development team from developing and maintaining their monitoring solution. Suppose these tools are provided by the cloud provider that deploys the DevOps toolchain. In that case, such as monitoring AWS deployed DevOps toolchain with CloudWatch, the good integration between the tool and the cloud may provide more detailed monitoring to the cloud infrastructure than any self-built third-party tool.

However, there are still limitations when using serverless computing services. One most significant limitation is that the Fargate does not support runs privileged container. This was largely limiting what we operation can we do within the running Docker build agent; for example, we cannot build a Docker image for the project that is being built in the build agent. Moreover, the performance of the build agent runs in the serverless computing service is lower. This is due to the cold-start problem of serverless computing. Also, the average time of build stage of the same software project is longer in the serverless build agent (92s) than in EC2 based build agent(65s). This difference shows that the serverless computing engine has lower performance when it runs the automated build and automated testing within a continuous delivery pipeline.

RQ2: *How does the newly emerged integrated toolchain in Amazon Web Services compare with the traditional non-integrated toolchain?*

Compared with non-integrated toolchain, the integrated toolchain stands out with lower development difficulty, better integration with other AWS services, less time and effort to develop and better customer support. However the lower performance and higher cost makes it not always a better solution to the software development team.

The first advantage for integrated toolchain is it is an out-of-box solution that needs much less time and effort to develop and setup. In contrast, building the non-integrated toolchain needs many works, namely, tool selection, cloud infrastructure design and management, manual configuration and even software development. Besides, an integrated toolchain provides a better sight on the DevOps toolchain as a whole. This is easier to achieve because the integrated toolchain is delivered as a single platform. In our experiment on performance, we find that the integrated toolchain is faster in provisioning the computing resource for continuous delivery. We believe due to the computing resources are also managed by AWS, it easier for AWS to optimize the provisioning process.

However, we also find under the same hardware configuration, and it takes the integrated toolchain longer time to run the automated build and automated testing. The gap of the runtime between integrated and non-integrated toolchain increased with the number of parallel executing jobs. Our research shows this is because the AWS is imitating the maximum hardware resources that we could use. The restriction shows the first disadvantage of the integrated toolchain – the development team does not have full control on the underlying hardware. A similar problem with integrated toolchain is that the team also has low visibility on the status of the underlying hardware, which brought some challenges in our implementation. Moreover, our analysis shows that the cost of the integrated toolchain is higher in the non-integrated toolchain. The high cost is because, first, not possible to manly use open-source solution, second, the AWS charges much more on computing resource when using it from it's DevOps toolchain. Furthermore, the integrated toolchains limit the tool selection, which on the one hand reduce the time needed for tool selection, on the other hand, it limiting the use of some special tools that are not supported by the toolchain.

Although the AWS integrated toolchain has a higher cost per running hour, the higher price also means better services and less time and effort to develop and maintain the toolchain. The team also need to consider if they have specific tools they need to use that cannot be integrated with AWS toolchain. In the performance aspect, the delivery speed of AWS toolchain could vary dramatically depends on the service number in the microservices and the project complexity. Generally, it is slower compared with Jenkins build agent with the same hardware consideration. The team need to consider if the speed is a sensitive factor in their business and if the project needs parallel pipeline execution.

6.2 Future Work

Based on our current implementation and findings, we propose following further works could be done.

- Due to the tight schedule in finishing the thesis, our case project is rather simple, and the test cases are rather small. With the increasing popularity of microservices architecture, it is interesting if we extend the case project to a microservices. Such an extension could better simulate the software project in real-life development. We can further include more type of software projects into the case projects, which could strengthen our conclusion by show that the conclusion is applicable to all kinds of the software projects.
- Current our experiment only simulate the DevOps workflow of a simple project. However, the delivery frequency, the size and the development behaviour of a software team could also affect the toolchain's performance. It could be interesting if we could test the toolchains in a real-life software development team to see how does them performs. We can also get a more user-experience related result by interviewing the team members about the experience and difficulties when develop and use the toolchains.

Bibliography

- [1] Kent Beck, Mike Beedle, Arie Van Bennekum, Alistair Cockburn, Ward Cunningham, Martin Fowler, James Grenning, Jim Highsmith, Andrew Hunt, Ron Jeffries, et al. Manifesto for agile software development. 2001.
- [2] Marco Miglierina. Application deployment and management in the cloud. In *2014 16th International Symposium on Symbolic and Numeric Algorithms for Scientific Computing*, pages 422–428. IEEE, 2014.
- [3] Ramtin Jabbari, Nauman bin Ali, Kai Petersen, and Binish Tanveer. What is devops? a systematic mapping study on definitions and practices. In *Proceedings of the Scientific Workshop Proceedings of XP2016*, pages 1–11, 2016.
- [4] Gene Kim, Jez Humble, Patrick Debois, and John Willis. *The DevOps Handbook:: How to Create World-Class Agility, Reliability, and Security in Technology Organizations*. IT Revolution, 2016.
- [5] What is a devops toolchain? – bmc blogs. <https://www.bmc.com/blogs/devops-toolchain/>. (Accessed on 03/13/2020).
- [6] Devops - wikipedia. <https://en.wikipedia.org/wiki/DevOps>. (Accessed on 02/24/2020).
- [7] Christof Ebert, Gorka Gallardo, Josune Hernantes, and Nicolas Serrano. Devops. *Ieee Software*, 33(3):94–100, 2016.
- [8] Devops toolchain - wikipedia. https://en.wikipedia.org/wiki/DevOps_toolchain. (Accessed on 03/11/2020).
- [9] Liming Zhu, Len Bass, and George Champlin-Scharff. Devops and its practices. *IEEE Software*, 33(3):32–34, 2016.
- [10] Serverless computing – amazon web services. <https://aws.amazon.com/serverless/>. (Accessed on 05/25/2020).
- [11] Serverless computing vs. containers how to choose — cloudflare. <https://www.cloudflare.com/learning/serverless/serverless-vs-containers/>. (Accessed on 05/25/2020).

- [12] Devops as a service: Automation in the cloud — sumo logic. <https://www.sumologic.com/insight/devops-as-a-service/>. (Accessed on 05/25/2020).
- [13] Gartner says worldwide iaas public cloud services market grew 31.3% in 2018, 07 2019. (Accessed on 05/21/2020).
- [14] Kim McMahon. The state of cloud native development. *KEY INSIGHTS FOR THE CLOUD NATIVE COMPUTING FOUNDATION STATE OF DEVELOPER NATION Q2 2019*, 05 2020.
- [15] Y. Kim and J. Lin. Serverless data analytics with flint. In *2018 IEEE 11th International Conference on Cloud Computing (CLOUD)*, pages 451–455, 2018.
- [16] Alfonso Pérez, Germán Moltó, Miguel Caballer, and Amanda Calatrava. Serverless computing for container-based architectures. *Future Generation Computer Systems*, 83:50–59, 2018.
- [17] Stefan Nastic, Thomas Rausch, Ognjen Scekcic, Schahram Dustdar, Marjan Gusev, Bojana Koteska, Magdalena Kostoska, Boro Jakimovski, Sasko Ristov, and Radu Prodan. A serverless real-time data analytics platform for edge computing. *IEEE Internet Computing*, 21(4):64–71, 2017.
- [18] Alex Glikson, Stefan Nastic, and Schahram Dustdar. Deviceless edge computing: extending serverless computing to the edge of the network. In *Proceedings of the 10th ACM International Systems and Storage Conference*, pages 1–1, 2017.
- [19] Vitalii Ivanov and Kari Smolander. Implementation of a devops pipeline for serverless applications. In *International Conference on Product-Focused Software Process Improvement*, pages 48–64. Springer, 2018.
- [20] Shashikant Bangera. *DevOps for Serverless Applications: Design, deploy, and monitor your serverless applications using DevOps practices*. Packt Publishing Ltd, 2018.
- [21] Jim Highsmith. What is agile software development? *crosstalk*, 15(10):4–10, 2002.
- [22] Michael A Cusumano and Stanley A Smith. Beyond the waterfall: Software development at microsoft. 1995.
- [23] Kent Beck. Embracing change with extreme programming. *Computer*, 32(10):70–77, 1999.
- [24] Agile software development - wikipedia. https://en.wikipedia.org/wiki/Agile_software_development#Iterative,_incremental_and_evolutionary. (Accessed on 03/18/2020).

- [25] Kent Beck, Mike Beedle, Arie Van Bennekum, Alistair Cockburn, Ward Cunningham, Martin Fowler, James Grenning, Jim Highsmith, Andrew Hunt, Ron Jeffries, et al. Principles behind the agile manifesto. *Agile Alliance*, pages 1–2, 2001.
- [26] Sten Pittet. Continuous integration vs. continuous delivery vs. continuous deployment. *Web-article. Atlassian.* <https://www.atlassian.com/continuous-delivery/ci-vs-ci-vs-cd>. Fetched, 24:2018, 2018.
- [27] Martin Fowler and Matthew Foemmel. Continuous integration, 2006.
- [28] Test automation in a ci/cd pipeline — spritecloud. <https://www.spritecloud.com/test-automation-with-ci-cd-pipeline/>. (Accessed on 03/19/2020).
- [29] Build automation - wikipedia. https://en.wikipedia.org/wiki/Build_automation. (Accessed on 03/19/2020).
- [30] Paul E Ceruzzi, E Paul, William Aspray, et al. *A history of modern computing*. MIT press, 2003.
- [31] M Fowler. Continuous delivery. may 30, 2013, 2013.
- [32] What is continuous delivery? - continuous delivery. <https://continuousdelivery.com/>. (Accessed on 03/23/2020).
- [33] Why agile isn't agile without continuous delivery. <https://www.atlassian.com/continuous-delivery/principles/why-agile-development-needs-continuous-delivery>. (Accessed on 08/10/2020).
- [34] Is continuous delivery replacing agile? <https://www.tempio.io/blog/2016/continuous-delivery-replacing-agile>. (Accessed on 08/10/2020).
- [35] Marko Leppänen, Simo Mäkinen, Max Pagels, Veli-Pekka Eloranta, Juha Itkonen, Mika V Mäntylä, and Tomi Männistö. The highways and country roads to continuous deployment. *Ieee software*, 32(2):64–72, 2015.
- [36] Len Bass, Ingo Weber, and Liming Zhu. *DevOps: A software architect's perspective*. Addison-Wesley Professional, 2015.
- [37] Leonardo Leite, Carla Rocha, Fabio Kon, Dejan Milojicic, and Paulo Meirelles. A survey of devops concepts and challenges. *ACM Computing Surveys (CSUR)*, 52(6):1–35, 2019.
- [38] Devops in a scaling environment - tajawal - medium. <https://medium.com/tech-tajawal/devops-in-a-scaling-environment-9d5416ecb928>. (Accessed on 03/27/2020).

- [39] Lucy Ellen Lwakatare, Pasi Kuvaja, and Markku Oivo. Relationship of devops to agile, lean and continuous deployment. In *International conference on product-focused software process improvement*, pages 399–415. Springer, 2016.
- [40] Ian Buchanan. Agile and devops: Friends or foes. *Atlassian Agile Coach*, 2015.
- [41] Mandi Walls. *Building a DevOps culture*. ” O’Reilly Media, Inc.”, 2013.
- [42] Lucy Ellen Lwakatare, Pasi Kuvaja, and Markku Oivo. Dimensions of devops. In *International conference on agile software development*, pages 212–217. Springer, 2015.
- [43] Dror G Feitelson, Eitan Frachtenberg, and Kent L Beck. Development and deployment at facebook. *IEEE Internet Computing*, 17(4):8–17, 2013.
- [44] There’s no such thing as a ”devops team” - continuous delivery. <https://continuousdelivery.com/2012/10/theres-no-such-thing-as-a-devops-team/>. (Accessed on 03/26/2020).
- [45] Jordan Shropshire, Philip Menard, and Bob Sweeney. Uncertainty, personality, and attitudes toward devops. 2017.
- [46] FMA Erich, Chintan Amrit, and Maya Daneva. A qualitative study of devops usage in practice. *Journal of software: Evolution and Process*, 29(6):e1885, 2017.
- [47] Devopsculture. <https://martinfowler.com/bliki/DevOpsCulture.html>. (Accessed on 03/27/2020).
- [48] Matej Artac, Tadej Borovssak, Elisabetta Di Nitto, Michele Guerriero, and Damian Andrew Tamburri. Devops: introducing infrastructure-as-code. In *2017 IEEE/ACM 39th International Conference on Software Engineering Companion (ICSE-C)*, pages 497–498. IEEE, 2017.
- [49] M HERING, D DeGrandis, and N Forsgren. Measure efficiency effectiveness, and culture to optimize devops transformation. devops enterprise forum, 2015.
- [50] Michael Hüttermann. *DevOps for developers*. Apress, 2012.
- [51] N Forsgren, J Humble, and G Kim. Accelerate: state of devops report: Strategies for a new economy. dora (devops research and assessment) and google cloud, 2018.
- [52] Toolchain - wikipedia. <https://en.wikipedia.org/wiki/Toolchain>. (Accessed on 03/11/2020).

- [53] Nicole Forsgren Velasquez, Gene Kim, Nigel Kersten, and Jez Humble. State of devops report, 2014.
- [54] Nicole Forsgren, Dustin Smith, Jez Humble, and Jessie Frazelle. 2019 accelerate state of devops report. 2019.
- [55] M. Artac, T. Borovssak, E. Di Nitto, M. Guerriero, and D. A. Tamburri. Devops: Introducing infrastructure-as-code. In *2017 IEEE/ACM 39th International Conference on Software Engineering Companion (ICSE-C)*, pages 497–498, 2017.
- [56] Source and version control in devops – bmc blogs. <https://www.bmc.com/blogs/devops-source-version-control/>. (Accessed on 05/09/2020).
- [57] Ben Chacon, Scott; Straub. *Getting Started – About Version Contro*. 2014. (Accessed on 08/11/2020).
- [58] Ioana Baldini, Paul Castro, Kerry Chang, Perry Cheng, Stephen Fink, Vatche Ishakian, Nick Mitchell, Vinod Muthusamy, Rodric Rabbah, Aleksander Slominski, et al. Serverless computing: Current trends and open problems. In *Research Advances in Cloud Computing*, pages 1–20. Springer, 2017.
- [59] Eric Jonas, Johann Schleier-Smith, Vikram Sreekanti, Chia-Che Tsai, Anurag Khandelwal, Qifan Pu, Vaishaal Shankar, Joao Carreira, Karl Krauth, Neeraja Yadwadkar, et al. Cloud programming simplified: A berkeley view on serverless computing. *arXiv preprint arXiv:1902.03383*, 2019.
- [60] Alexander Zahariev. Google app engine. *Helsinki University of Technology*, pages 1–5, 2009.
- [61] Serverless computing - wikipedia. https://en.wikipedia.org/wiki/Serverless_computing. (Accessed on 06/01/2020).
- [62] Paul Castro, Vatche Ishakian, Vinod Muthusamy, and Aleksander Slominski. The rise of serverless computing. *Communications of the ACM*, 62(12):44–54, 2019.
- [63] Hyungro Lee, Kumar Satyam, and Geoffrey Fox. Evaluation of production serverless computing environments. In *2018 IEEE 11th International Conference on Cloud Computing (CLOUD)*, pages 442–450. IEEE, 2018.
- [64] Thomson reuters case study. <https://aws.amazon.com/solutions/case-studies/thomson-reuters/>. (Accessed on 06/01/2020).
- [65] Serverless computing – amazon web services. https://aws.amazon.com/serverless/#Serverless_application_use_cases. (Accessed on 06/02/2020).

- [66] Tim Wagner. Serverlessconf 2018 keynote - debunking serverless myths. <https://www.slideshare.net/TimWagner/serverlessconf-2018-keynote-debunking-serverless-myths>, ServerlessConf 2018 2018. (Accessed on 08/11/2020).
- [67] Serverless computing — google cloud. <https://cloud.google.com/serverless>. (Accessed on 06/02/2020).
- [68] Azure serverless — microsoft azure. <https://azure.microsoft.com/en-us/solutions/serverless/#solutions>. (Accessed on 06/02/2020).
- [69] Aws lambda limits - aws lambda. <https://docs.aws.amazon.com/lambda/latest/dg/gettingstarted-limits.html>. (Accessed on 06/03/2020).
- [70] Joseph M Hellerstein, Jose Faleiro, Joseph E Gonzalez, Johann Schleier-Smith, Vikram Sreekanti, Alexey Tumanov, and Chenggang Wu. Serverless computing: One step forward, two steps back. *arXiv preprint arXiv:1812.03651*, 2018.
- [71] Keeping functions warm - how to fix aws lambda cold start issues. <https://www.serverless.com/blog/keep-your-lambdas-warm/>. (Accessed on 06/03/2020).
- [72] Tim Wagner. Serverless networking is the next step in the evolution of serverless — by tim wagner — a cloud guru, 10 2019. (Accessed on 08/11/2020).
- [73] What is amazon elastic container service? - amazon elastic container service. <https://docs.aws.amazon.com/AmazonECS/latest/developerguide/Welcome.html>. (Accessed on 07/20/2020).
- [74] Amazon Cloud Services. Amazon cloudwatch - application and infrastructure monitoring. <https://aws.amazon.com/cloudwatch/>. (Accessed on 08/12/2020).
- [75] TIOBE. index — tiobe - the software quality company. <https://www.tiobe.com/tiobe-index/>. (Accessed on 06/11/2020).
- [76] GitHub. The state of the octoverse. <https://octoverse.github.com/>, 2019. (Accessed on 06/11/2020).
- [77] Programming languages used in most popular websites - wikipedia. https://en.wikipedia.org/wiki/Programming_languages_used_in_most_popular_websites. (Accessed on 06/11/2020).
- [78] Spring — why spring? <https://spring.io/why-spring>. (Accessed on 06/11/2020).

- [79] Spring boot. <https://spring.io/projects/spring-boot>. (Accessed on 06/11/2020).
- [80] John Ferguson Smart. *Jenkins: The Definitive Guide: Continuous Integration for the Masses*. "O'Reilly Media, Inc.", 2011.
- [81] Jenkins. Pipeline-jenkins user documentation. <https://www.jenkins.io/doc/book/pipeline/>. (Accessed on 06/16/2020).
- [82] Software Testing Fundamentals. Unit testing - software testing fundamentals. <http://softwaretestingfundamentals.com/unit-testing/>. (Accessed on 08/12/2020).
- [83] Abel Avram. Docker: Automated and consistent software deployments. *InfoQ. Retrieved*, pages 08–09, 2013.
- [84] What is a container? — app containerization — docker. <https://www.docker.com/resources/what-container>. (Accessed on 06/22/2020).
- [85] Steven J Vaughan-Nichols. What is docker and why is it so darn popular? <https://www.zdnet.com/article/what-is-docker-andwhy-is-it-so-darn-popular>, 2014. (Accessed on 08/12/2020).
- [86] Ben Lloyd Pearson. Who's using docker? — opensource.com. <https://opensource.com/business/14/7/docker-through-hype>, 07 2014. (Accessed on 08/13/2020).
- [87] Overview — drone. <https://docs.drone.io/runner/docker/overview/>. (Accessed on 06/10/2020).
- [88] Overview — prometheus. <https://prometheus.io/docs/introduction/overview/>. (Accessed on 08/05/2020).
- [89] Git - about version control. <https://git-scm.com/book/en/v2/Getting-Started-About-Version-Control>. (Accessed on 06/12/2020).
- [90] Jon Loeliger and Matthew McCullough. *Version Control with Git: Powerful tools and techniques for collaborative software development*. "O'Reilly Media, Inc.", 2012.
- [91] Compare repositories - open hub. <https://www.openhub.net/repositories/compare>. (Accessed on 06/12/2020).
- [92] Git. Git - branches in a nutshell. <https://git-scm.com/book/en/v2/Git-Branching-Branches-in-a-Nutshell>. (Accessed on 08/13/2020).
- [93] Atlassian. Git merge — atlassian git tutorial. <https://www.atlassian.com/git/tutorials/using-branches/git-merge>. (Accessed on 08/13/2020).

- [94] GitHub Guides. Understanding the github flow, 2013.
- [95] Vincent Driessen. A successful git branching model. *URL* <http://nvie.com/posts/a-successful-git-branching-model>, 2010.
- [96] Scott Chacon. Github flow. 2011, 2011.
- [97] Luca Tiozzo. Aws ecs host auto-scaling with custom cloudwatch metrics and aws lambda. <https://medium.com/thron-tech/aws-ecs-host-auto-scaling-with-custom-cloudwatch-metrics-and-aws-lambda-b9a9f55faf1d>, 04 2019. (Accessed on 07/05/2020).
- [98] Aws codecommit — managed source control service. <https://aws.amazon.com/codecommit/>. (Accessed on 08/14/2020).
- [99] Bo Yang, Anca Sailer, and Ajay Mohindra. Survey and evaluation of blue-green deployment techniques in cloud native environments. In Sami Yangui, Athman Bouguettaya, Xiao Xue, Noura Faci, Walid Gaaloul, Qi Yu, Zhangbing Zhou, Nathalie Hernandez, and Elisa Y. Nakagawa, editors, *Service-Oriented Computing – ICSOC 2019 Workshops*, pages 69–81, Cham, 2020. Springer International Publishing.
- [100] Cloud Foundry Documentation. Using blue-green deployment to reduce downtime and risk — cloud foundry docs. <https://docs.cloudfoundry.org/devguide/deploy-apps/blue-green.html>, 04 2019. (Accessed on 07/07/2020).
- [101] Zhamak Dehghani. How to break a monolith into microservices. *Thought Works,[Online]*. Available: <https://martinfowler.com/articles/breakmonolith-into-microservices.html>. [Accessed 2 January 2019], 2018.
- [102] BILL DOERRFELD. How to scale microservices ci/cd pipelines - devops.com. <https://devops.com/how-to-scale-microservices-ci-cd-pipelines/>, 05 2020. (Accessed on 08/15/2020).