

A Cloud-Based DevOps Toolchain for Efficient Software Development

Ruiyang Ding

Master's Thesis in Computer Science
EIT Digital Master's Programme in Cloud Computing and Services

Supervisor: Prof.dr.ir. D.H.J. Epema
Company Supervisor: Mikko Drocan

31.08.2020

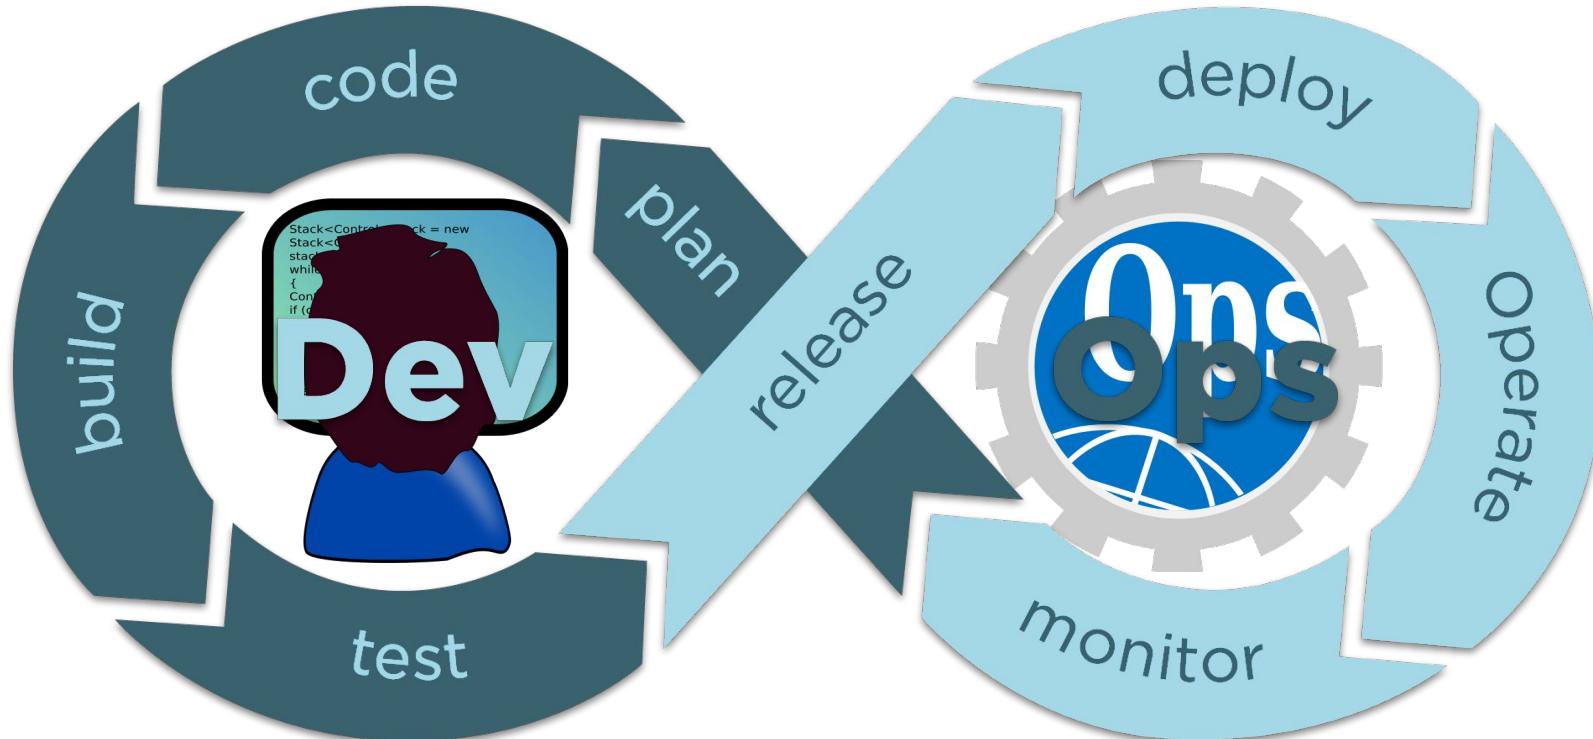


About Eficode



- Leading DevOps company in Europe
- Helping the customer in their DevOps transformation.

Introduction - DevOps



DevOps - Components

Culture



Automation



Organization

Monitoring &
Measurement

DevOps Toolchain

“The integration between tools that specialised in different aspects of the DevOps ecosystem, which support and coordinate the DevOps practices.”

Serverless

“A cloud execution model in which
the cloud provider manages the **server** and **resources allocation**”

Compared with Infrastructure as a service (IaaS):

- Event Driven
IaaS: Keep running
- Managed Resources Allocation
IaaS: Server creation and configurations are needed
- Pay-per-use, only pay for the code execution time
IaaS: Pay for type of VM and the rental time of this VM

Serverless in AWS



Advanced AWS Partner

[

We want to see how AWS serverless platform can helps in building the DevOps toolchains

]

DevOps ❤️ AWS Serverless ?

Computing



Monitoring



DevOps



.....

AWS serverless platform

Problem 1:

There is not enough research on how serverless can be used in DevOps toolchain



DevOps for Serverless Application Development?

1. DevOps for Serverless Applications: Design, deploy, and monitor your serverless applications using DevOps practices (Shashikant Bangera, 2018)
2. Implementation of a DevOps pipeline for serverless applications (Ivanov, Vitalii and Smolander, Kari, 2018)



Use Serverless in DevOps?

1. Several blogs such as “The Serverless Path to DevOps” (Sharjeel Yusuf, Feb 27, 2020)
2. No directly related paper so far

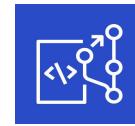
DevOps ❤️ Serverless ?

Problem 2

Now the software teams have two kinds of toolchains to select.

- Non-integrated DevOps Toolchain:
Individual tools which were stand-alone and from different vendors
- Integrated Toolchain:
Delivered as a single platform from a single vendor

AWS
DevOps
Tools



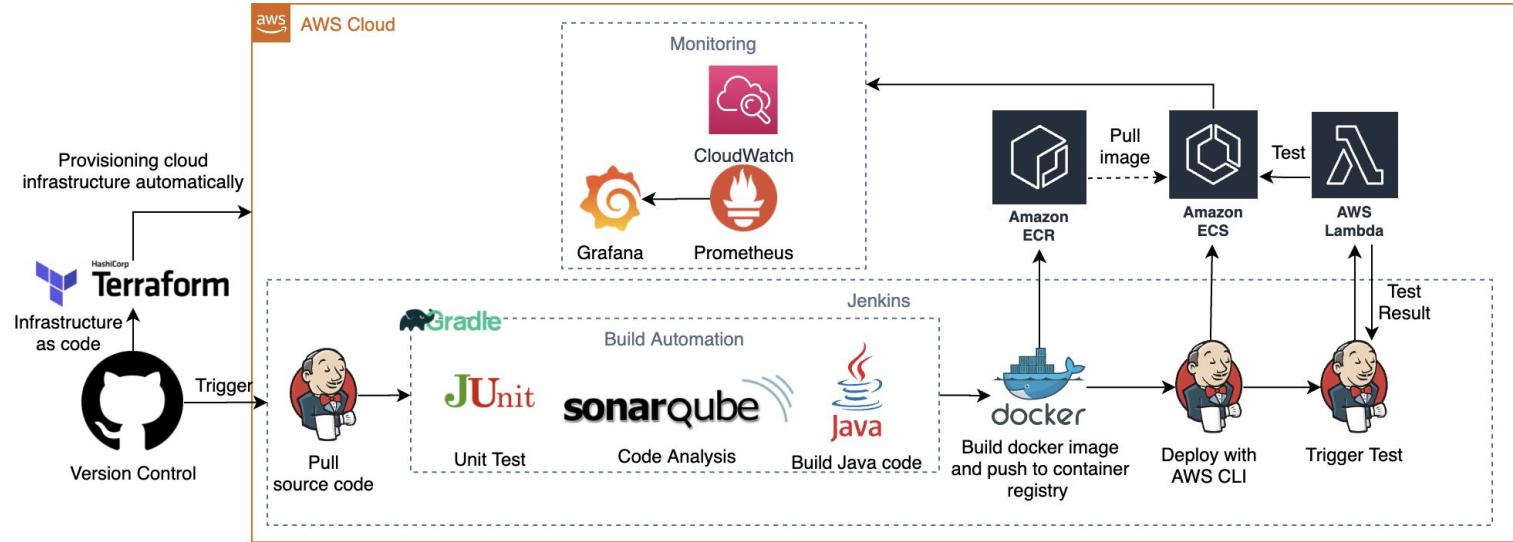
Research Questions

1. How can serverless computing services in Amazon Web Services enhance the DevOps toolchain?
2. How does the integrated toolchain build with AWS DevOps Tools compare with the traditional non-integrated toolchain?

Research Approach

1. Investigate Serverless Offering in AWS
2. Develop Cloud-Based DevOps toolchain
 - a. Non-integrated toolchain + Serverless computing services -> **Experiment I & RQ1**
 - b. Integrated toolchain -> **Experiment II**
3. Experiments
 - a. **Experiment I** : On non-integrated toolchain -> **RQ1**
 - b. **Experiment II** : Compare integrated/non-integrated toolchain -> **RQ2**

Non-integrated toolchain - Overview



Problems in implementation:

- Lot of time and effort for:
 - Setting up cloud infrastructure
 - Configuration of Jenkins
- Lack of documentation, unstable Jenkins plugins
- Previous cloud knowledge needed

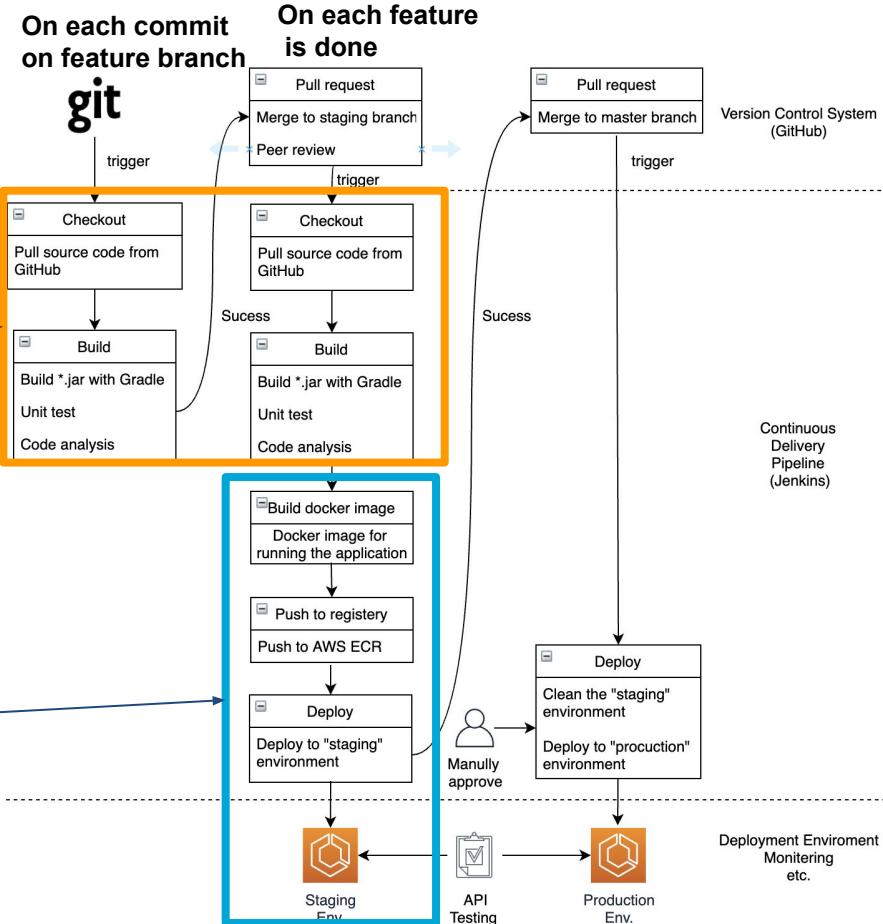
Non-integrated toolchain - Workflow

- Heavy, 80% pipeline runtime.
- More frequent.

Q: How to accelerate?

Why not run on feature branch?

- According to DevOps practice, developer should merge code often (several time a day)
- Already pretty frequent
- No need to run per commit

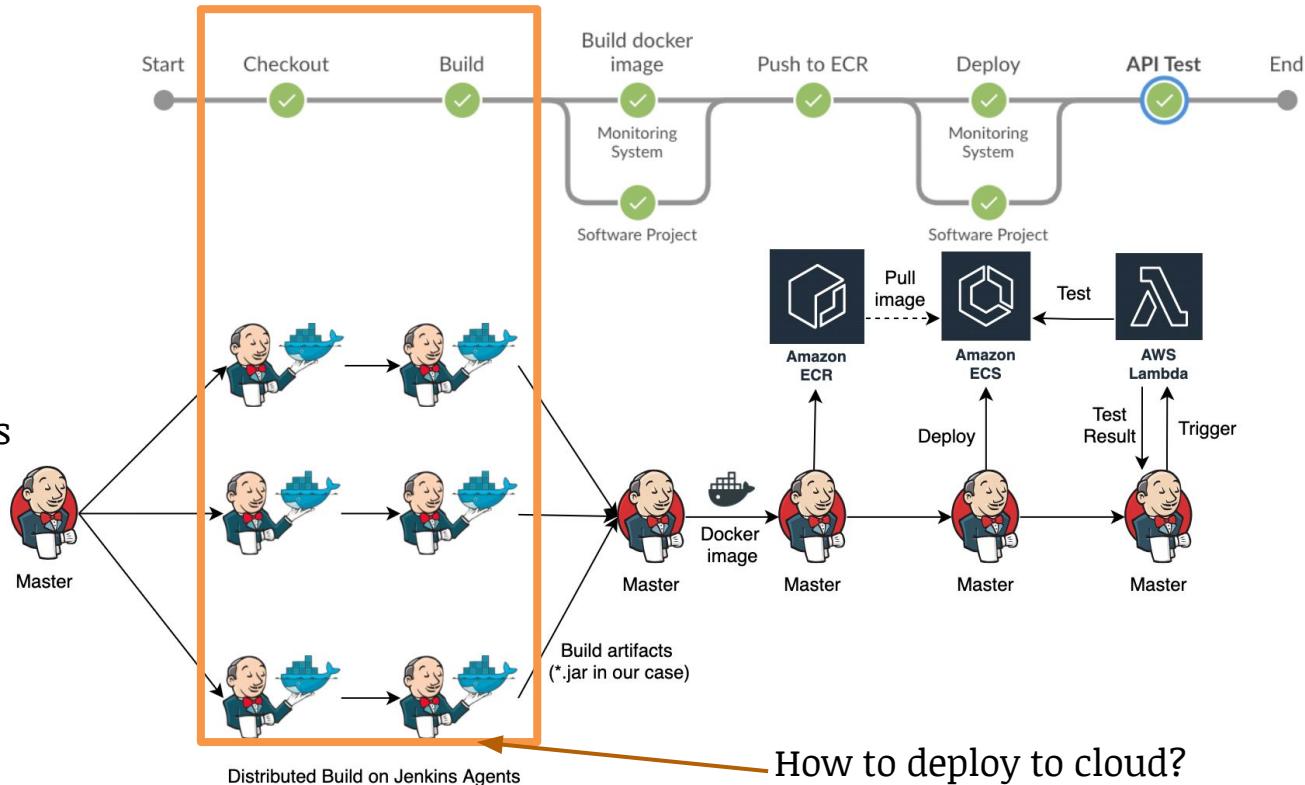


Non-integrated toolchain- Distributed Build

Solution:

Distributed Build with **build agents** run in **Docker containers**

- Parallel the heavy steps
- Why Docker? - Faster start, Easy to set up, Easy to change the build environment

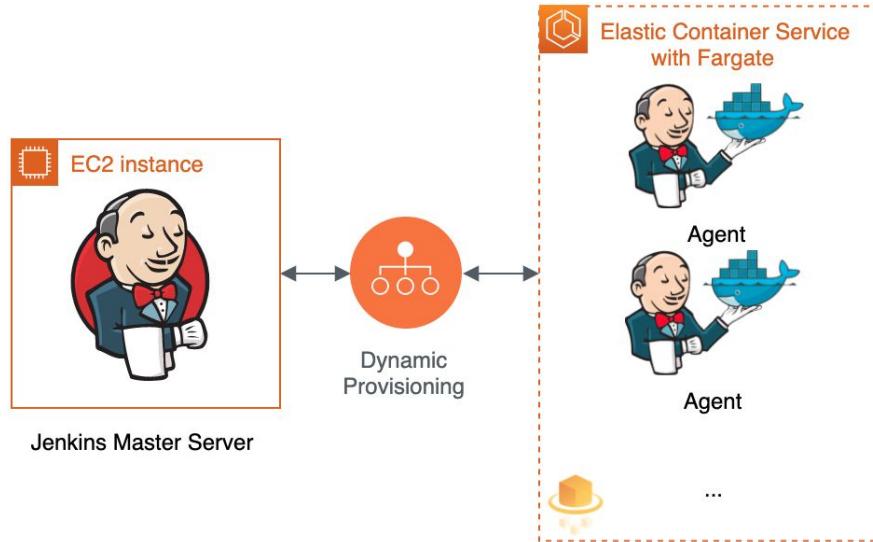


Using AWS serverless services - Fargate

Solution:

Hosting docker-based distributed build agent in **AWS Fargate**

- Master: Runs 24/7
-> Traditional VM
- Jenkins agents:
 - Runs in Docker container
 - on demand, runs only when new job comes
-> Serverless container services



Using AWS serverless services - Lambda

Test deployed API
with **AWS Lambda**

Manually?

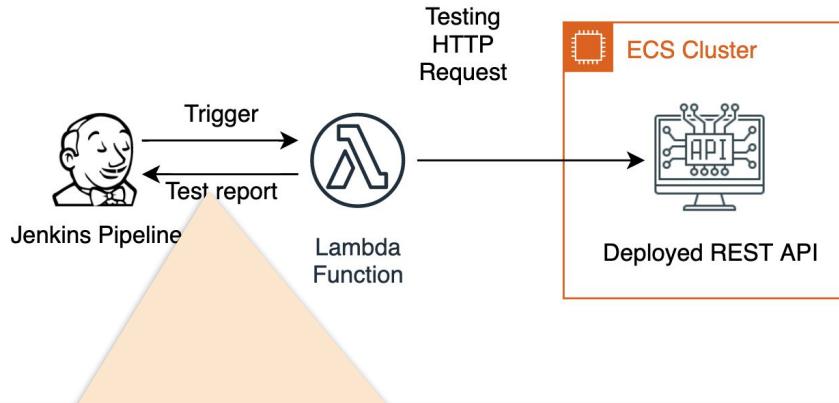
- We need automation

VM?

- Complicated
- Pay for idle time

Lambda

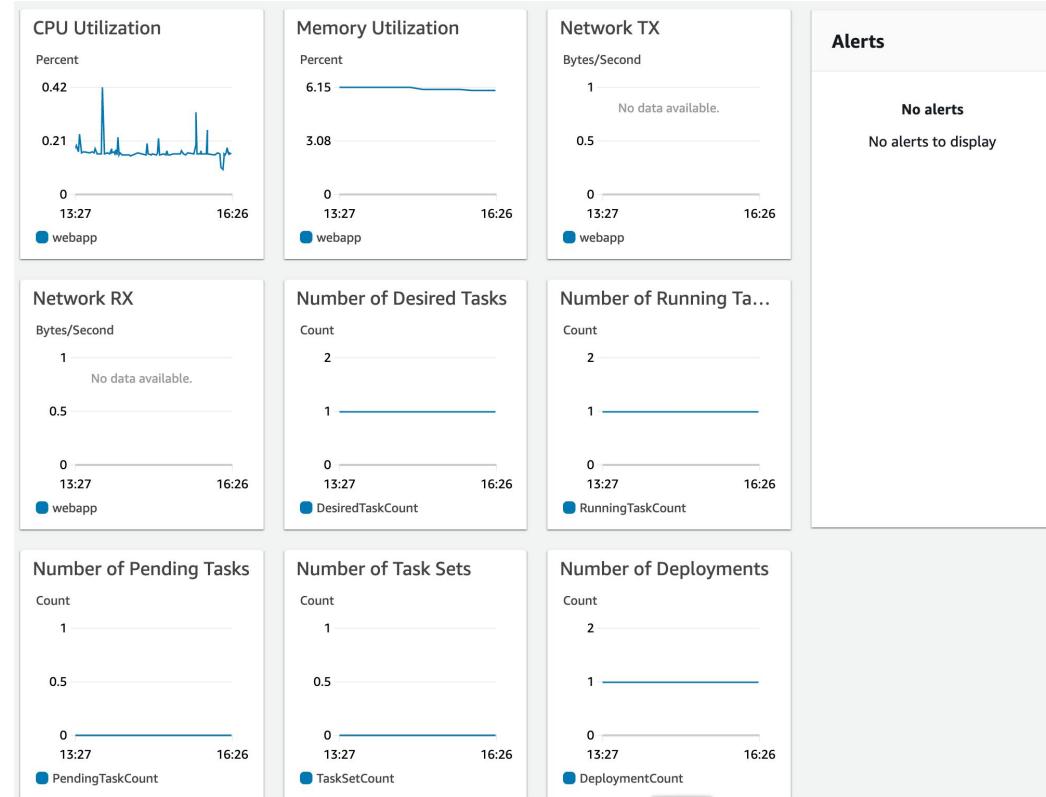
- Event-driving
- Easy to set up
- Pay per use



```
Testing the url http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/  
Testing the url http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages  
Get result of url: http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/  
1 out of 2 endpoints tested, 1 succeed  
Get result of url: http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages  
2 out of 2 endpoints tested, 2 succeed  
All endpoints are tested, result:  
{  
    "http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/": "Succeed, response body (first 300  
characters): Hello world",  
    "http://bluegreen-alb-1565976610.eu-north-1.elb.amazonaws.com/packages": "Succeed, response body (first  
300 characters): [{"name": "libws-commons-util-java", "description": "Common utilities from the  
Apache Web Services Project", "dependencies": ["\""], "reverseDependencies": []}, {"name": "python-pkg-  
resources", "description": "Package Discovery and Resource Access using  
pkg_resources", "dependencies": ["python (\u003e= 2.6)", "python (\u003c"]}  
}
```

Using AWS serverless services - CloudWatch

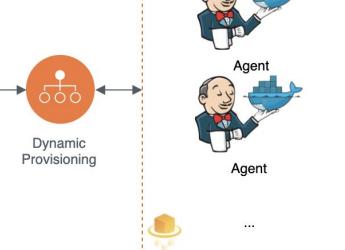
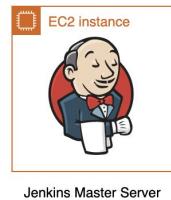
Monitoring ECS cluster with [AWS CloudWatch](#)



Experiment I - Fargate vs EC2

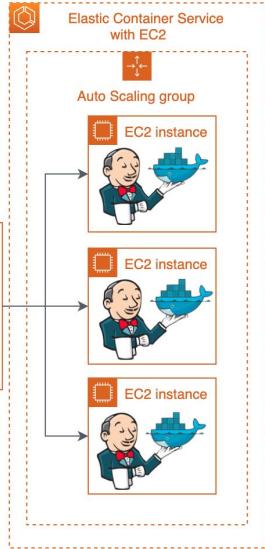


Case project:
A simple REST API web
service developed by me
with Spring Boot



Jenkins agents hosted by
AWS Fargate

VS



Jenkins agents hosted by a
group of Auto-Scaling AWS
EC2 virtual machines

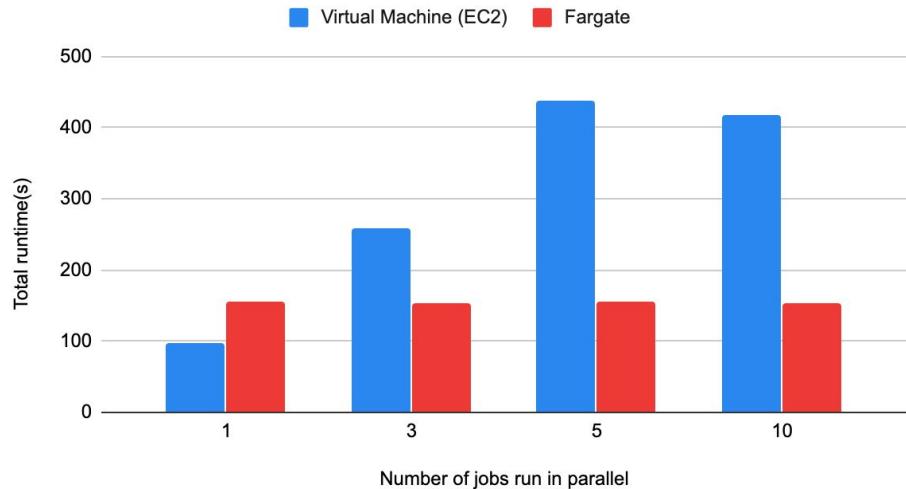


Experiment I - Result, Runtime

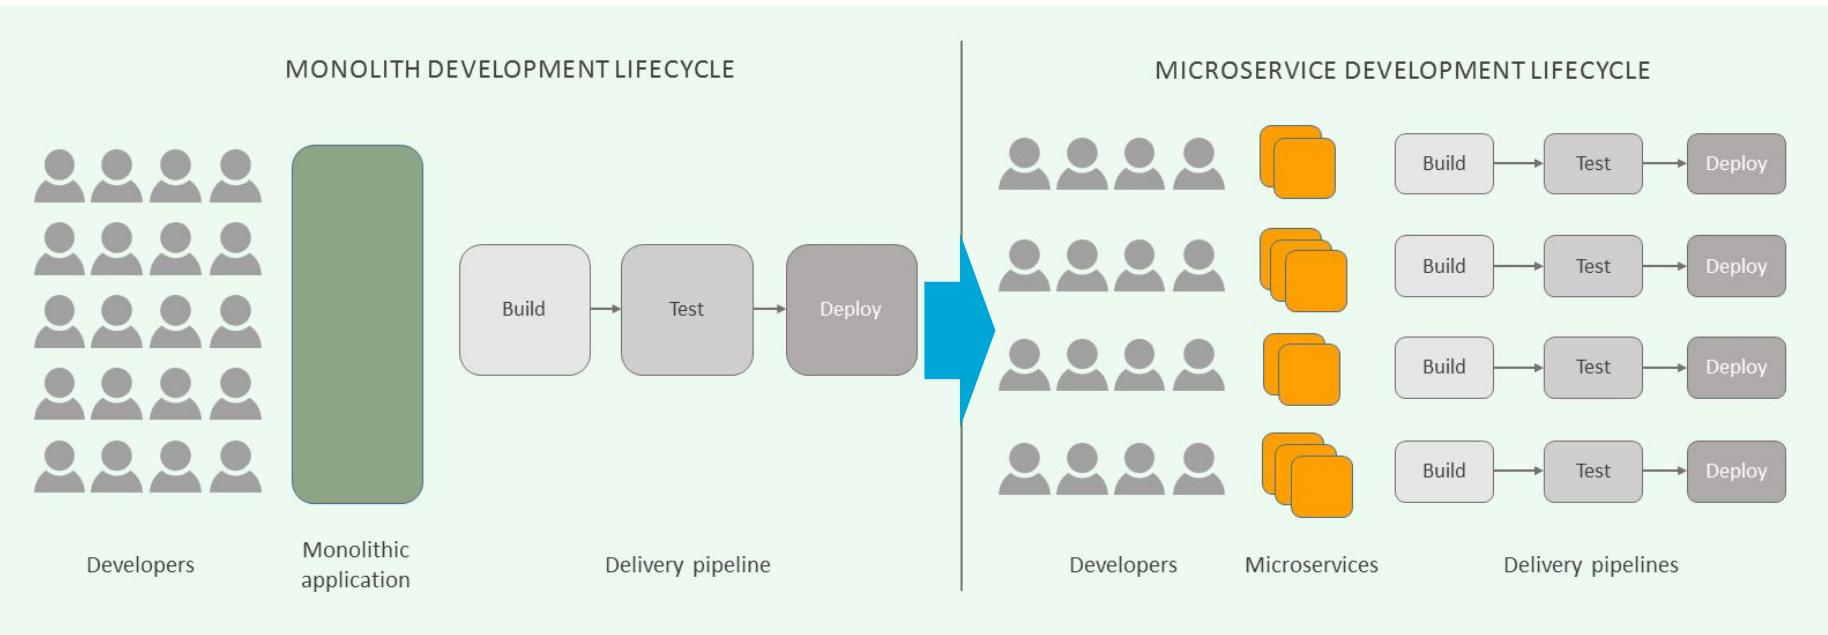
Performance with severless Jenkins agent:

2.8 X

under 5 tasks in parallel



Why Parallel performance is important?



Imaging having 100+ services in this microservices
DevOps Engineer:

- Manage 100+ CI servers (Jenkins server in our case)
- Manage one server, 100 + task runs in parallel

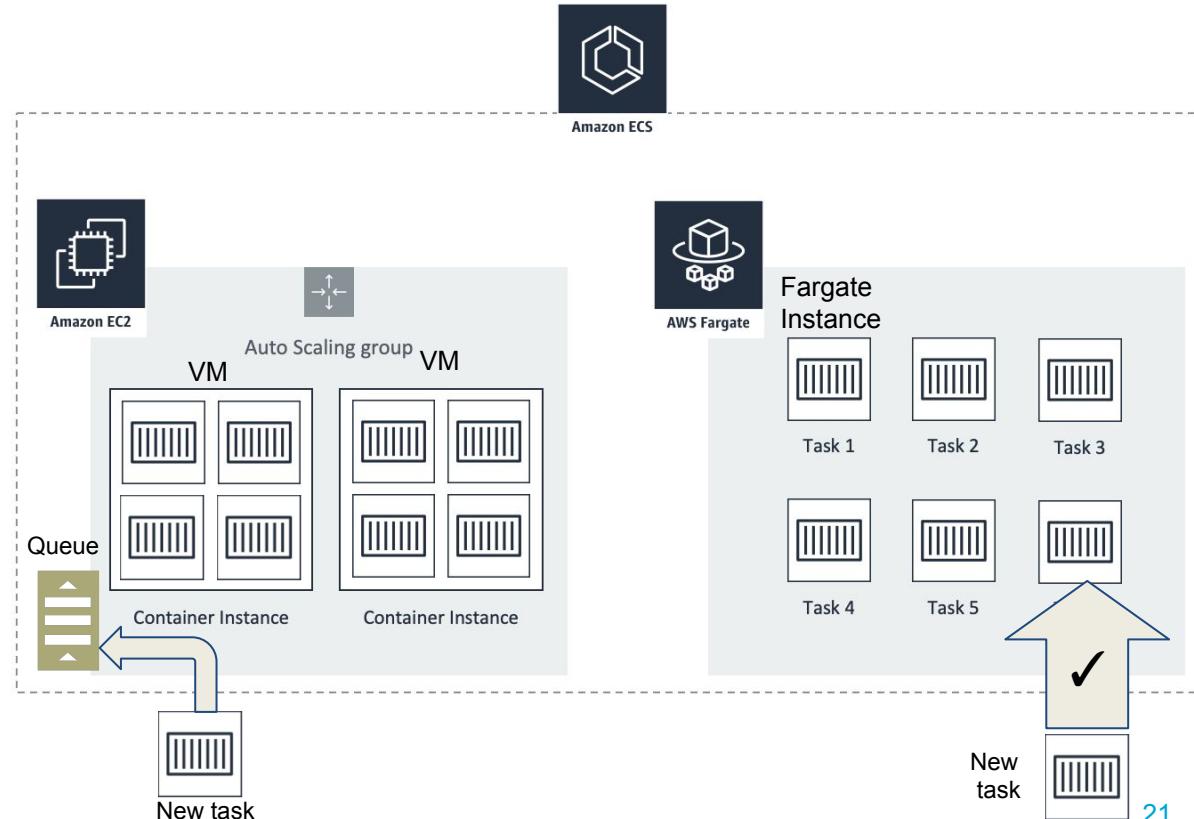
Experiment I - Result Evaluation, Runtime

EC2:

- Scaling is not related to the number of incoming tasks
- Competition for resource between agents on same VM
- Longer time to create new VM

Fargate:

- Unconditionally create a new instance
- Independent and isolated, no competition.
- Faster provision



Experiment I - Result, Cost & Resource Utilization

- Similar CPU/RAM utilization within container
- Fargate has higher cost per running hour
- No cost during Idle time - actual cost might lower

\$0.0216 vs \$0.0931

per hour runtime per agent, 2v CPU, 2 GB RAM

Integrated toolchain - Overview

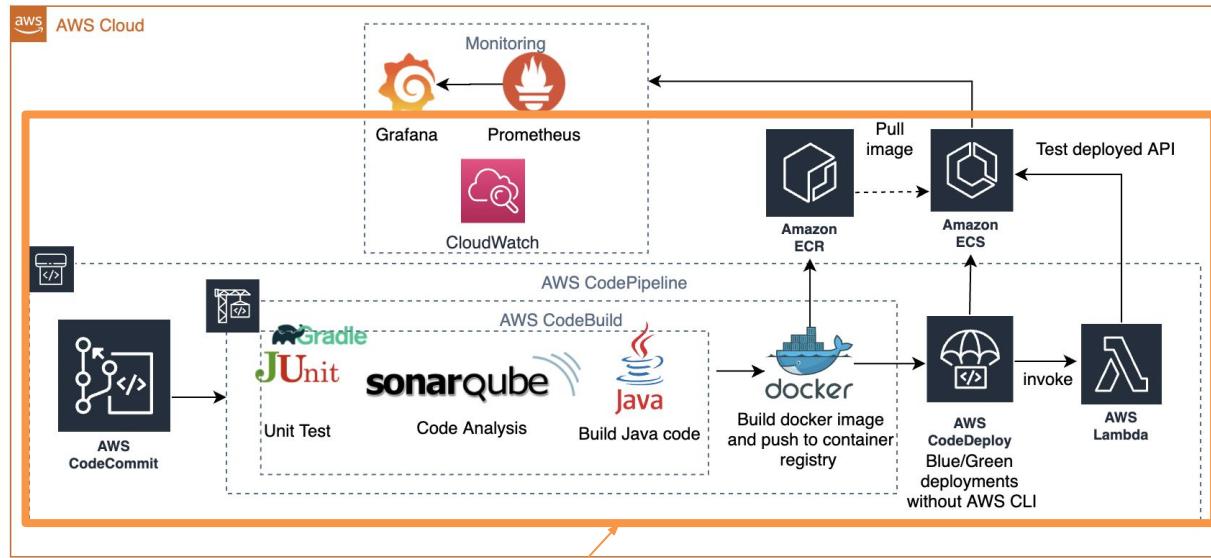
Compare with non-integrated toolchain:

Advantages:

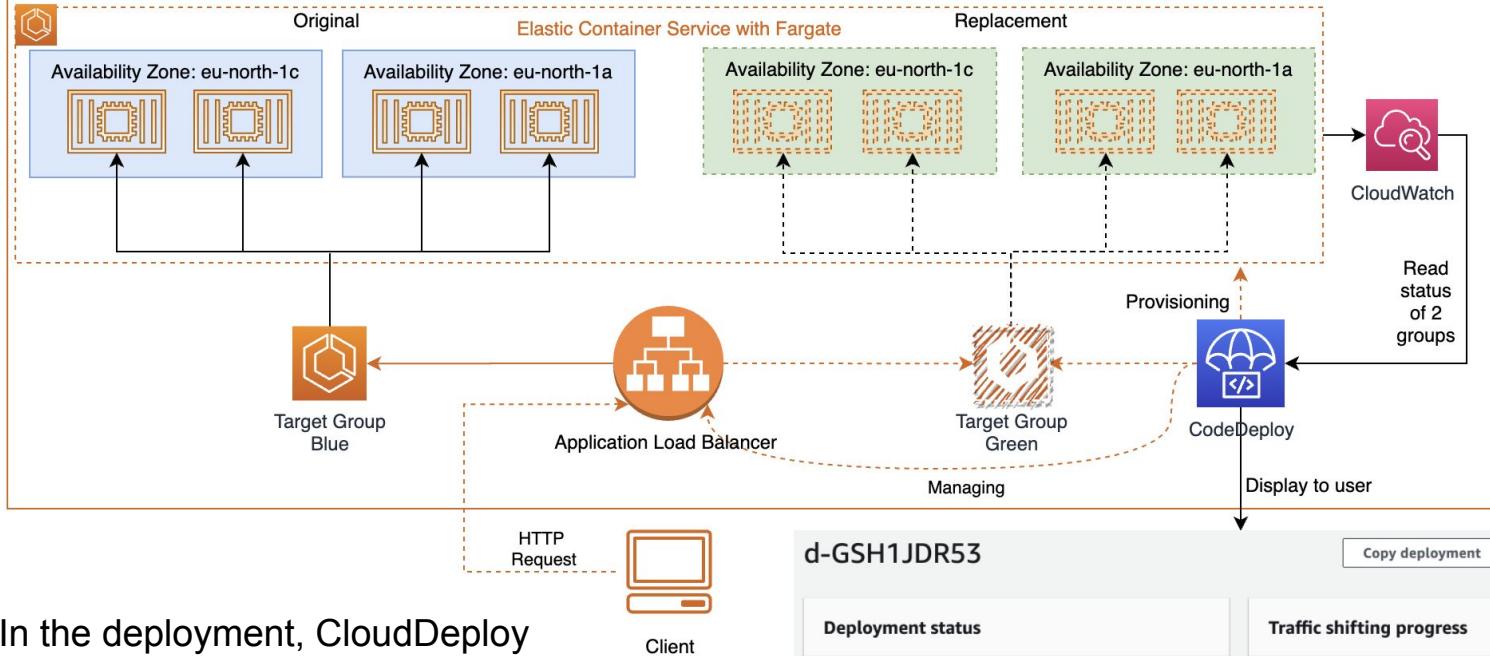
- + Less time and effort to set up
- + Better customer support
- + Better integration with AWS services -> **Blue/Green Deployment with CodeDeploy**

Disadvantages:

- Low visibility to infrastructure -> **Cannot identify error when CodeDeploy failed**
- Low freedom and extensibility -> **Third-party tools limited**

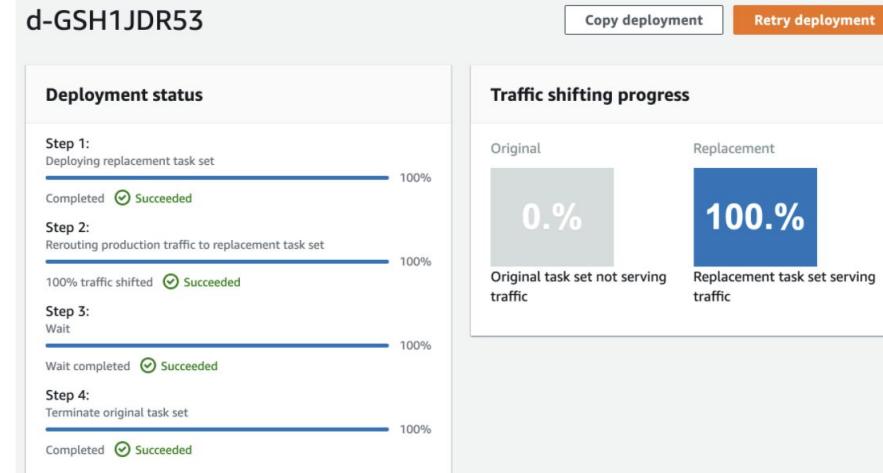


A single platform & serverless



In the deployment, CloudDeploy controls

- ALB
- Deployment Env. (ECS)
- Cloudwatch



Experiment II - Result, Runtime/Cost

Evaluation on integrated/non-integrated DevOps toolchain

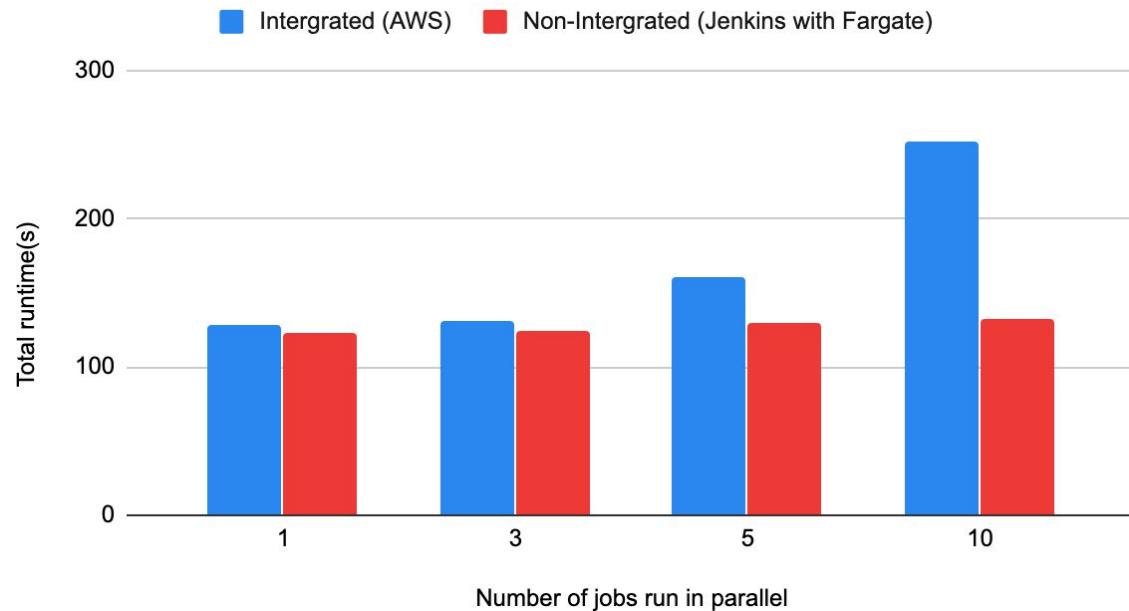
Non-integrated

- Better performance in parallel execution
- Lower cost, \$0.1452/agent/hour

Integrated

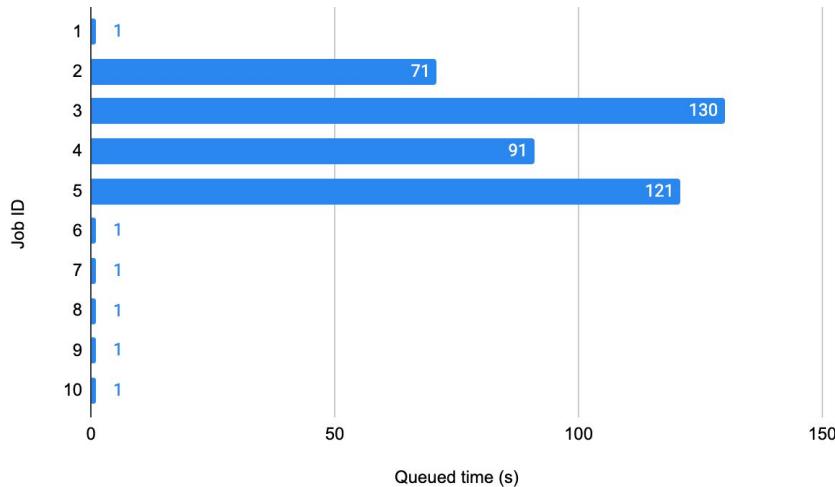
- Lower performance in parallel execution
->Why (next slides)
- Higher per hour cost, \$0.3/agent/hour.

However, no need to pay the master node



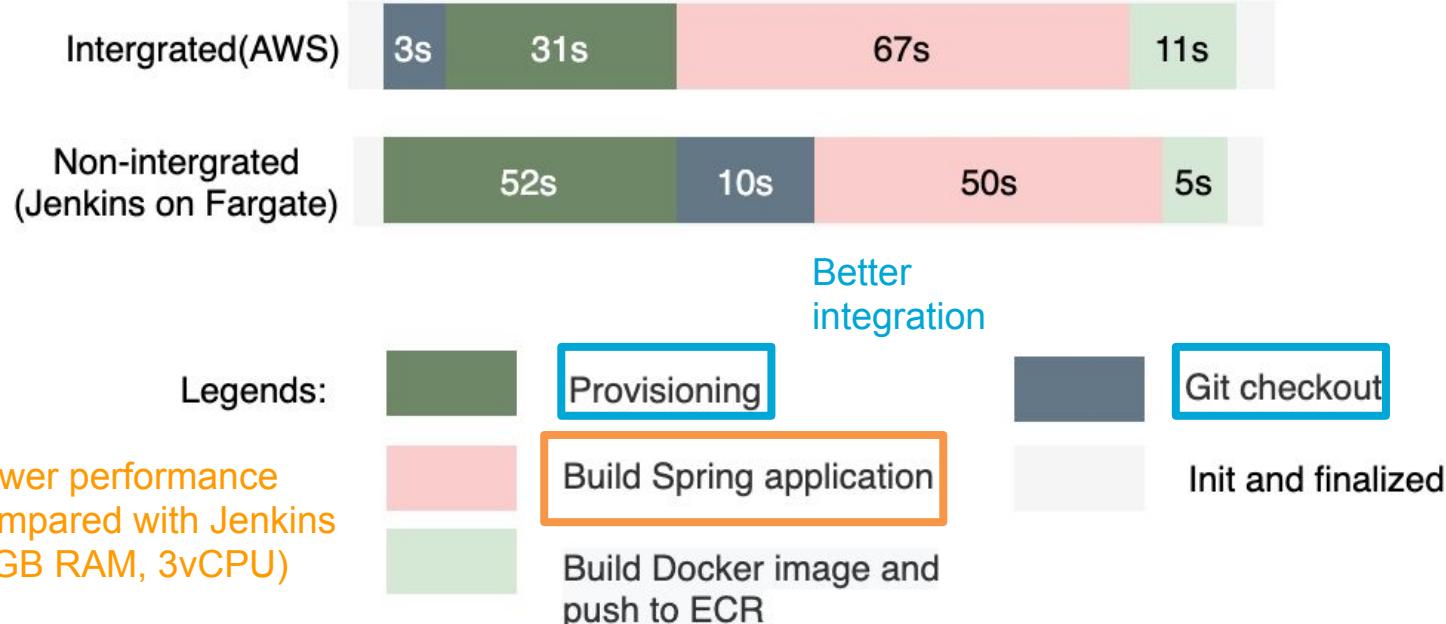
Experiment II - Result analysis, Runtime

- AWS limiting resource during parallel build
- 4 out of 10 jobs have to queue



Phase details				
Name	Status	Context	Duration	
SUBMITTED	✓ Succeeded	-	<1 sec	
QUEUED	✓ Succeeded	-	61 secs	
PROVISIONING	✓ Succeeded	-	17 secs	

Experiment II - Result analysis, Runtime



Decomposition of runtime of each stages in the pipeline

Conclusions

RQ1: How can serverless computing services in Amazon Web Services enhance the DevOps toolchain?

- Three possible services
 - Fargate for hosting distributed agents
 - Better parallel execution performance good for microservice development and larger team.
 - Higher hourly cost, total cost might not higher
 - CloudWatch for monitoring solution
 - Lambda for event driving task (e.t. integration test of deployed API)
 - Less maintenance effort
 - Lower cost
- In conclusion:
Improving performance, eliminating server management tasks, and reducing the cost of idle time.

Conclusions

RQ2: How does the integrated toolchain build with AWS DevOps Tools compare with the traditional non-integrated toolchain?

- Performance: Lower than Jenkins with agent in Fargate
 - Resource limitation -> long queueing time
 - Longer build time under same hardware.
- Cost: Higher per hour price, however no cost for idle time
- Development: Easy integration with other AWS services, less time and effort to develop and better customer support. Low visibility to the underlying infrastructure.

Future Works

- **More evaluation method on toolchains**
i.e. Have two software team in real-life test the toolchains.
- **More complicated case project in experiments**
 - Could be a microservice with different type of services
i.e. Database, REST API, Lambda functions, etc.
 - Strength our conclusion

Thanks for your attention!



Q & A

The following pages are skipped

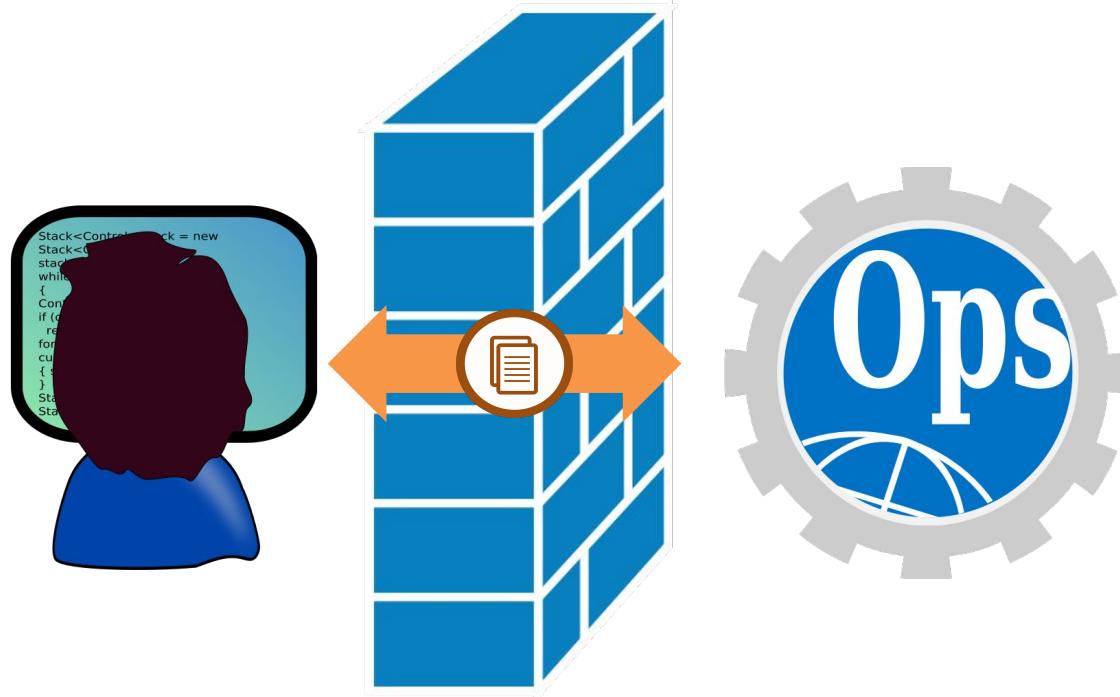
3mins, problem / RQs

10 mins p1

8 mins p2

4 mins conclusion

What is DevOps?



AWS Serverless platform

- AWS Fargate: Run container serverless
- AWS Lambda: Run code serverless, could be triggered by AWS events
- AWS CloudWatch: Monitoring application
- AWS DevOps Tools:
 - CodeCommit: Git version control; cooperation.
 - CodeBuild: Build Automation
 - CodeDeploy: Deploy application to AWS/ on-premise servers, support advanced deploy methodology
 - CodeStar: Connecting above services

Agenda

01
Background

03
Implementation

05
Result

02
Problem
Statement

04
Experiments

06
Conclusion

About Me



EIT Digital Double Degree
Cloud Computing and Services



DevOps Consultant

Eficode 

Helsinki, Finland 

What is DevOps?

A set of practices aims to combine traditionally separated disciplines (DEV + OPS), intended to speed up software delivery, while ensuring high quality.

DevOps - DevOps Toolchain and DevOps Practises

Infrastructure as
Code



Build Automation



Project Management &
Planning



Continuous Integration &
Delivery



Version Control



Monitoring



Test Automation

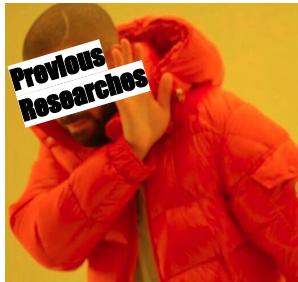


Problem: There is not enough research on how serverless can be used in DevOps toolchain



DevOps for Serverless Application?

1. DevOps for Serverless Applications: Design, deploy, and monitor your serverless applications using DevOps practices (Shashikant Bangera, 2018)
2. Implementation of a DevOps pipeline for serverless applications (Ivanov, Vitalii and Smolander, Kari, 2018)



Use Serverless in DevOps?

1. Several blogs such as “The Serverless Path to DevOps” (Sharjeel Yusuf, Feb 27, 2020)
2. No directly related paper so far

Services in AWS - Elastic Container Service(ECS) with AWS Fargate

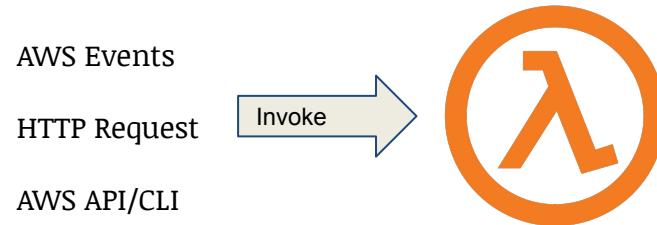


ECS: Container orchestration services.

Could runs Docker containers in:

- EC2 cluster
- Fargate -- Serverless service for Docker containers

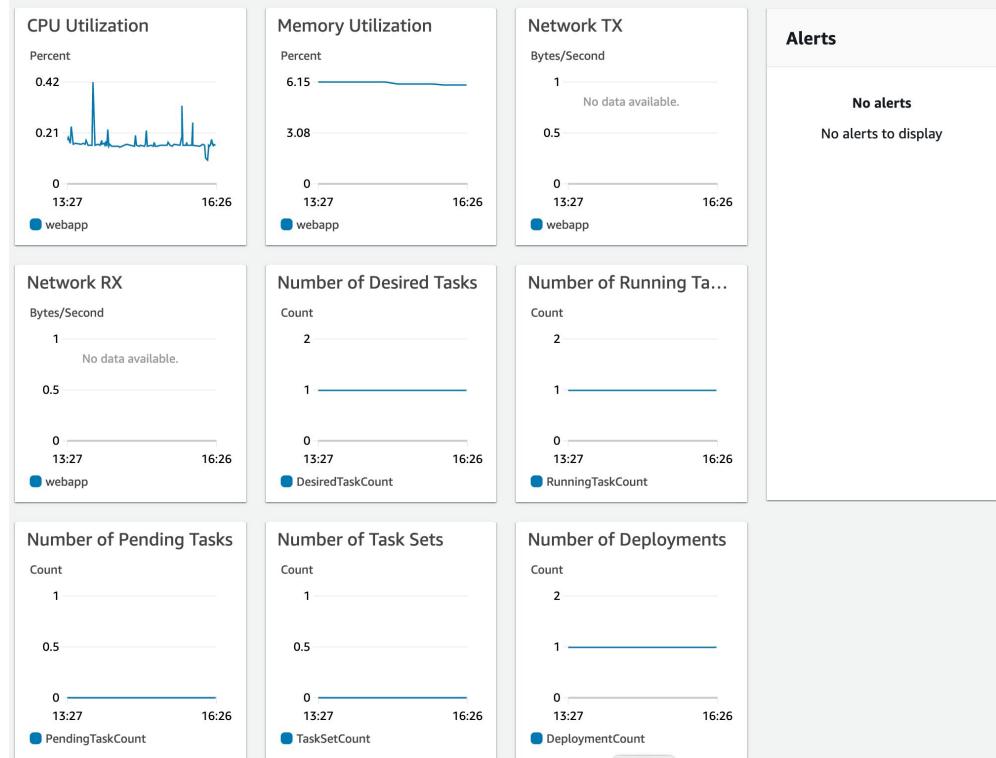
Services in AWS - AWS Lambda



- Run code without thinking about servers
- Event Driving
- Pay-as-you-go

Services in AWS - AWS CloudWatch

- Out-of-the-box
- Monitoring deployed application in AWS
- Logs and metrics



- [Open Blue Ocean](#)
- [Rename](#)
- [Pipeline Syntax](#)
- [GitHub Hook Log](#)

Build History

trend [=](#)

find X

#165 Jul 15, 2020 1:10 PM

#164 Jul 15, 2020 1:10 PM

#163 Jul 15, 2020 1:10 PM

#162 Jul 15, 2020 1:10 PM

#161 Jul 15, 2020 1:10 PM

#160 Jul 15, 2020 1:10 PM

#159 Jul 15, 2020 1:10 PM

#158 Jul 15, 2020 1:10 PM

#157 Jul 15, 2020 1:10 PM

#156 Jul 15, 2020 1:10 PM

#155 Jul 15, 2020 9:24 AM

#154 Jul 15, 2020 7:30 AM

#153 Jul 15, 2020 7:30 AM

#152 Jul 14, 2020 1:39 PM

#151 Jul 14, 2020 1:39 PM

#150 Jul 14, 2020 1:38 PM

#149 Jul 14, 2020 11:21 AM

#148 Jul 14, 2020 11:20 AM

#147 Jul 14, 2020 11:19 AM

Average stage times:
(Average full run time: ~2min)

#165
Jul 15
16:10
No Changes

#164
Jul 15
16:10
No Changes

#163
Jul 15
16:10
No Changes

#162
Jul 15
16:10
No Changes

#161
Jul 15
16:10
No Changes

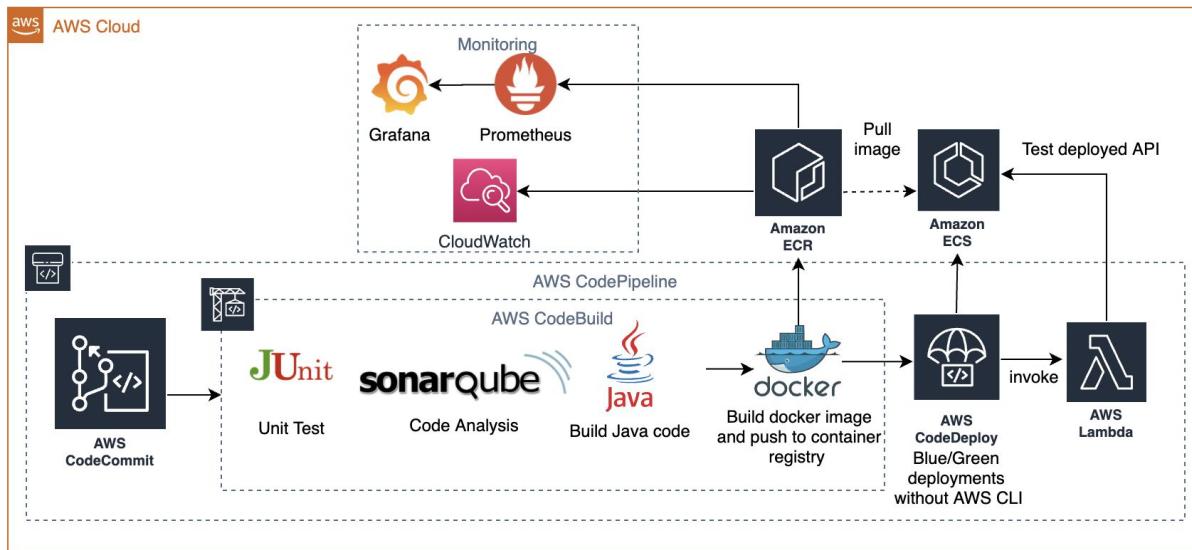
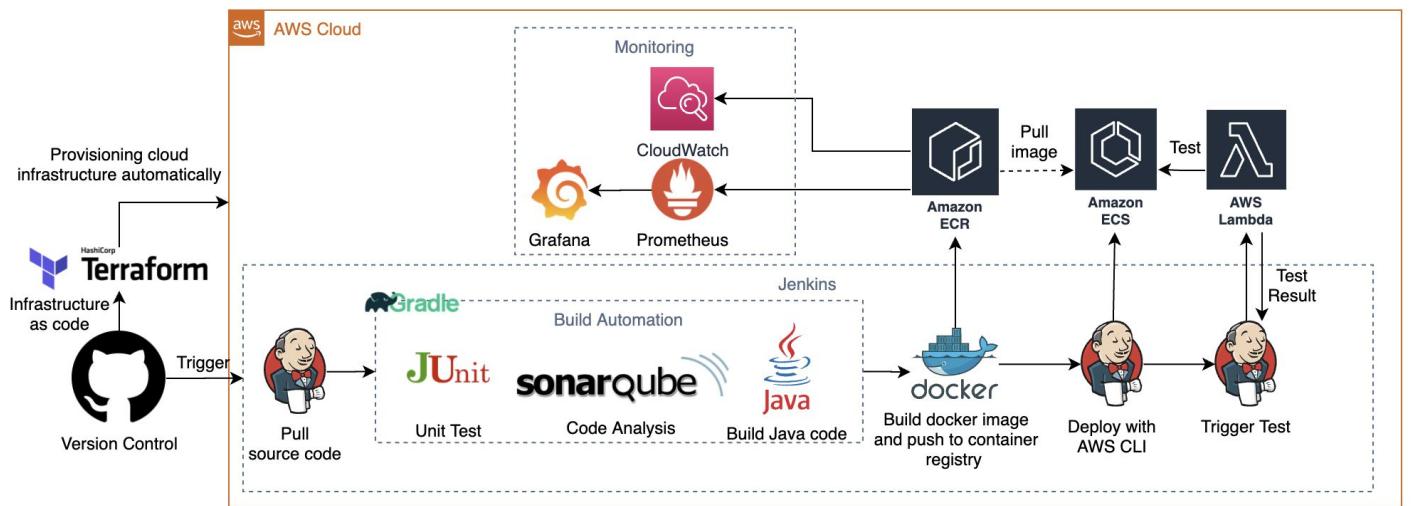
#160
Jul 15
16:10
No Changes

#159
Jul 15
16:10
No Changes

#158
Jul 15
16:10
No Changes





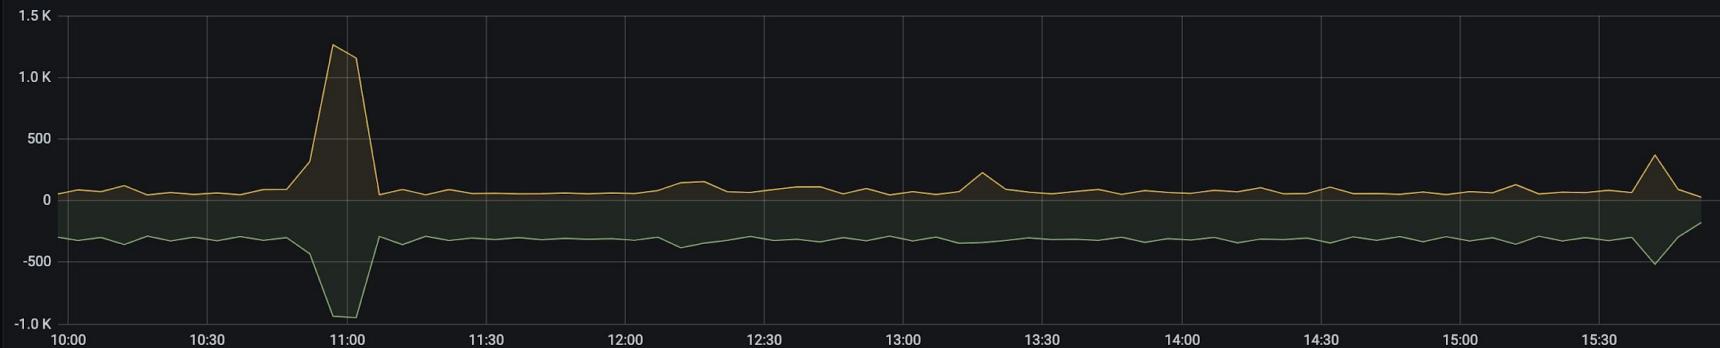




Amazon EC2



Total network traffic outbound (+) and inbound (-) [bytes/sec]



Network details

Inbound network traffic per instance [bytes/sec]



Outbound network traffic per instance [bytes/sec]



Inbound network packets per instance [IOPS]



Outbound network packets per instance [IOPS]

